

Zero Touch Kubernetes

Vollautomatisierte Infrastruktur mit Flatcar Container Linux



FroSCon

Free and Open Source Software Conference

2022-08-20

Hallo!

Ich bin Thilo.



Thilo Fromm
Engineering manager,
Microsoft

Github: [t-lo](#)

Mastodon: [@thilo@fromm.social](#)

Twitter: [ThiloFM](#)

Email: thilofromm@microsoft.com



Und Ihr?

Cluster-Admins / Verteilte Applikationen?



Und Ihr?



Cluster-Admins / Verteilte Applikationen?

Container-Applikationen/Workloads?



Und Ihr?



Cluster-Admins / Verteilte Applikationen?

Container-Applikationen/Workloads?

Kubernetes?





Container-optimiertes Linux

- Geschichte



Container-optimiertes Linux

CoreOS Container Linux (Okt. 2013 – Mai 2020)

(Basiert auf Chromium OS, welches auf Gentoo basiert)





Container-optimiertes Linux

CoreOS Container Linux (Okt. 2013 – Mai 2020)

(Basiert auf Chromium OS, welches auf Gentoo basiert)

→ Fedora CoreOS / Red Hat CoreOS (seit Juni 2018)

(basiert auf Fedora, eingeschränkt kompatibel zu CoreOS)





Container-optimiertes Linux

CoreOS Container Linux (Okt. 2013 – Mai 2020)

(Basiert auf Chromium OS, welches auf Gentoo basiert)

→ Fedora CoreOS / Red Hat CoreOS (seit Juni 2018)

(basiert auf Fedora, eingeschränkt kompatibel zu CoreOS)

→ Flatcar Container Linux (seit Nov. 2019)

("Friendly fork", vollständig mit CoreOS kompatibel)





Container-optimiertes Linux?





Container-optimiertes Linux?

Betriebssystem als Infrastruktur

Langweilig (nicht aufregend), “tut einfach”

Mehr Zeit für Applikationen / Business Logic

*Wenn Dein Lichtschalter ständig Aufmerksamkeit braucht,
kann das die Freude am Heimkino trüben*

Container-Isolation – auch vom Betriebssystem

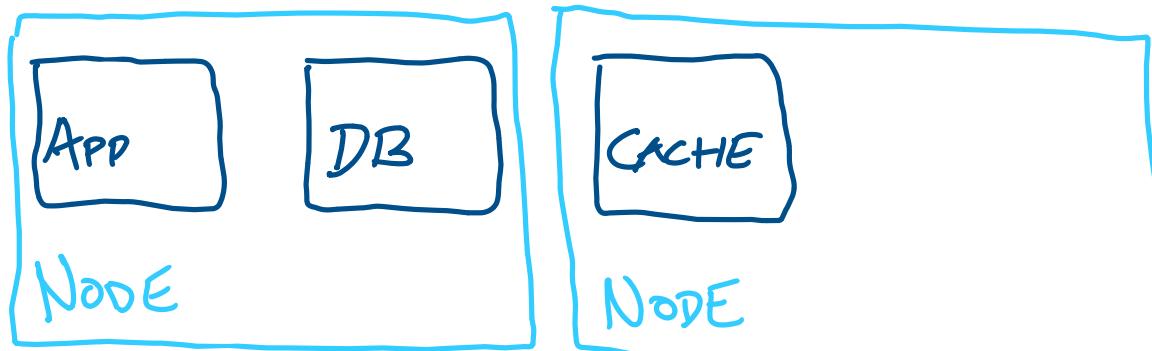
*Container-Applikationen sind voneinander isoliert.
Isoliert sie das nicht auch vom Betriebssystem?*





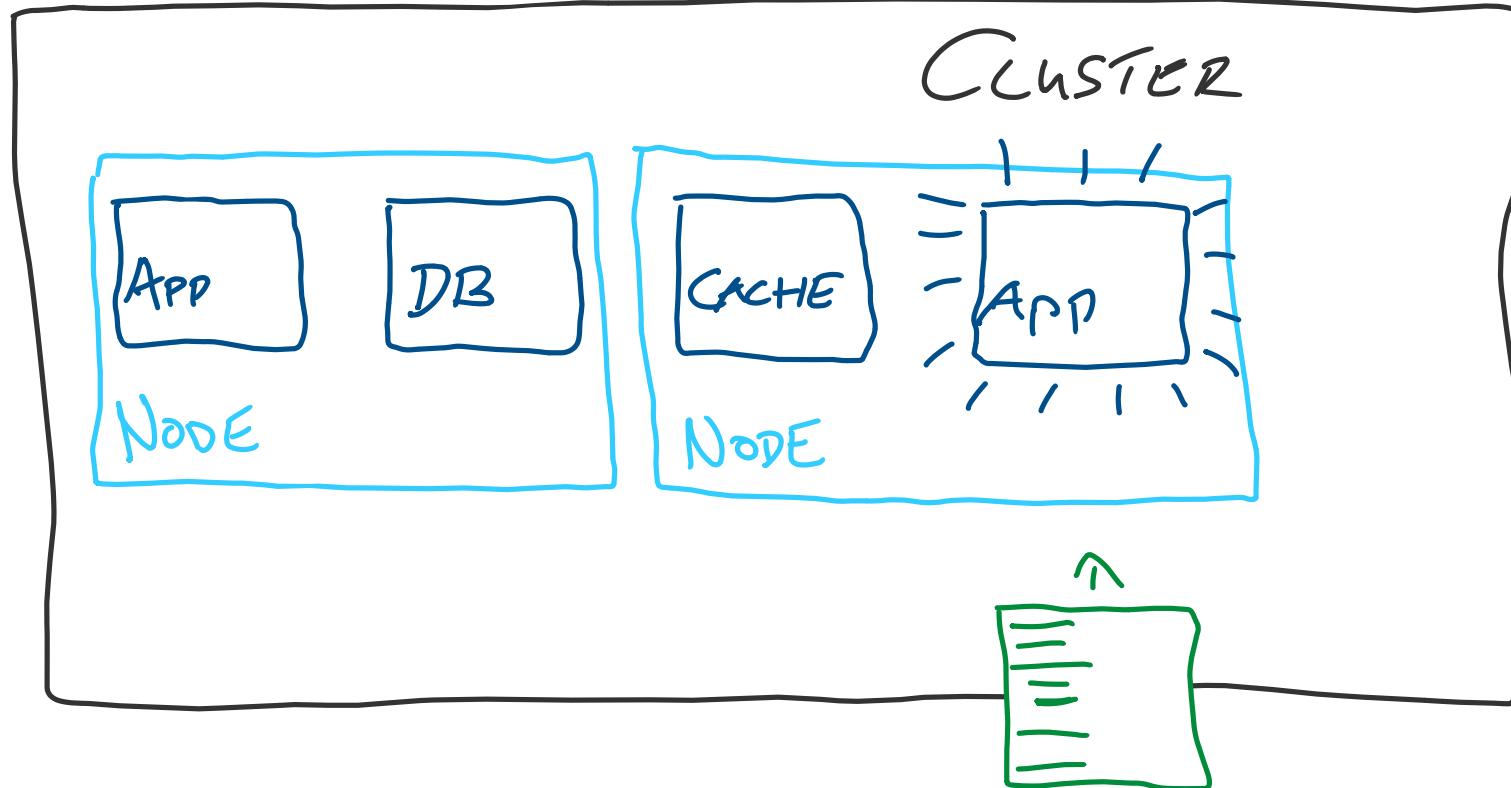
Container-optimiertes Linux?

CLUSTER



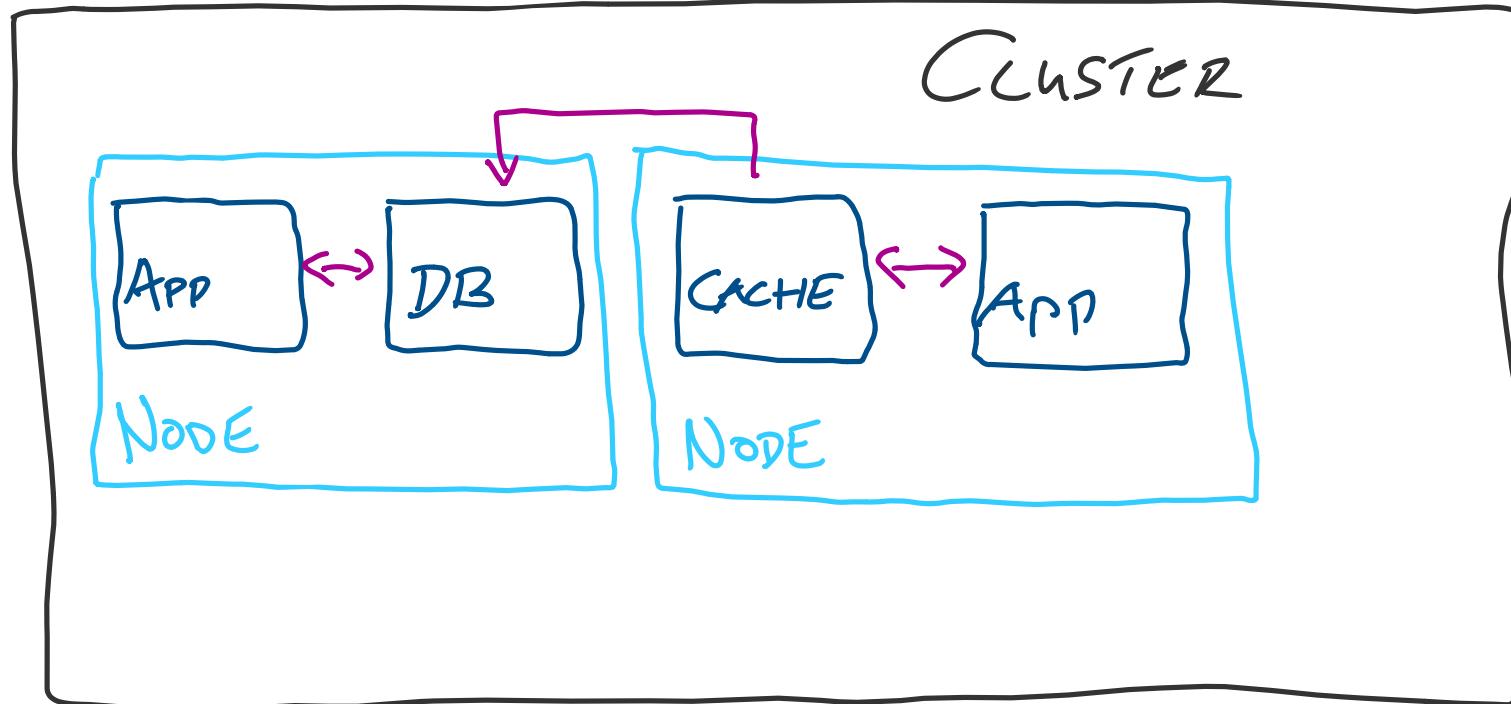


Container-optimiertes Linux?



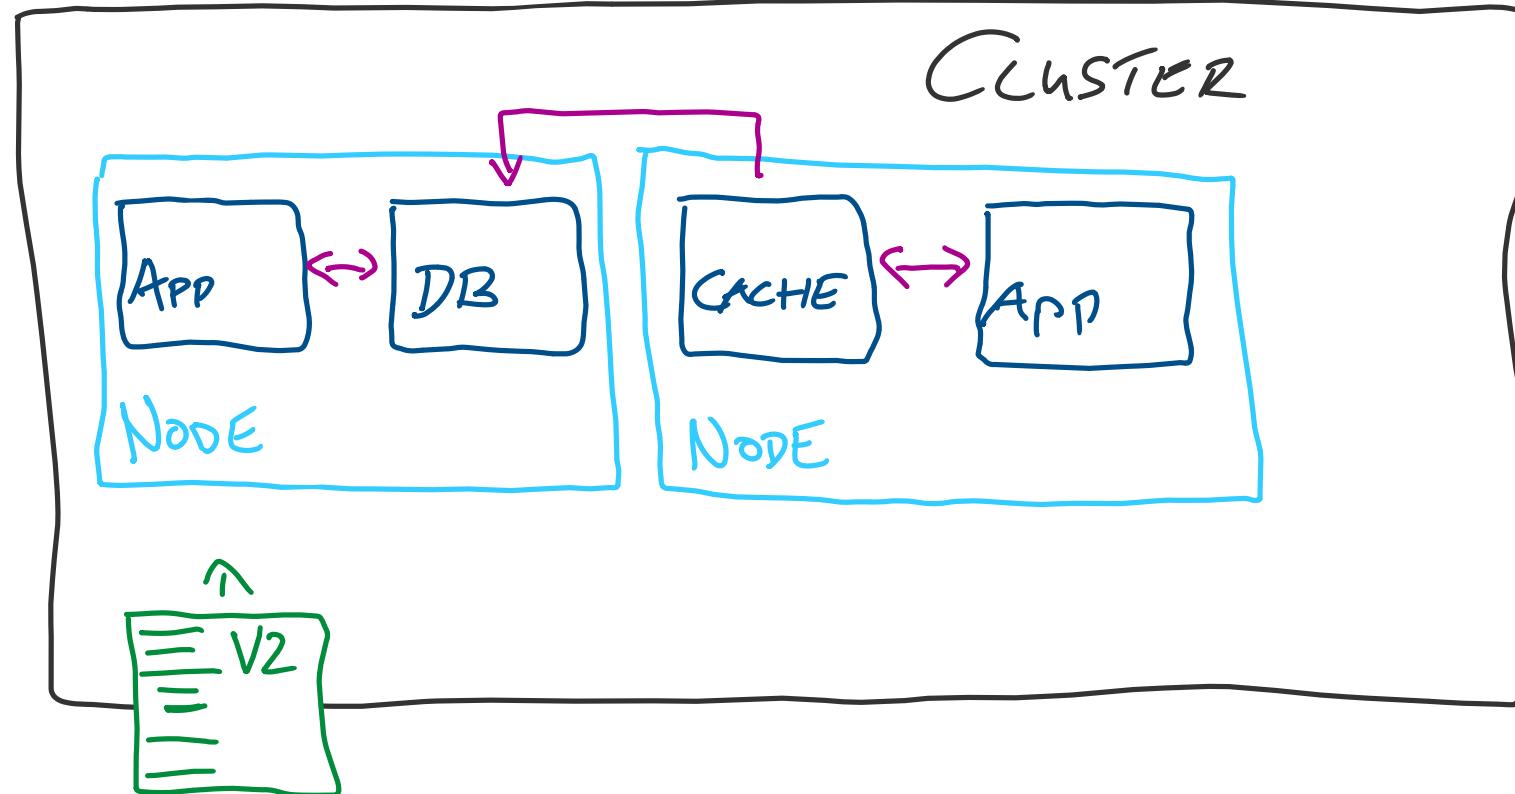


Container-optimiertes Linux?



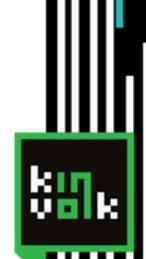
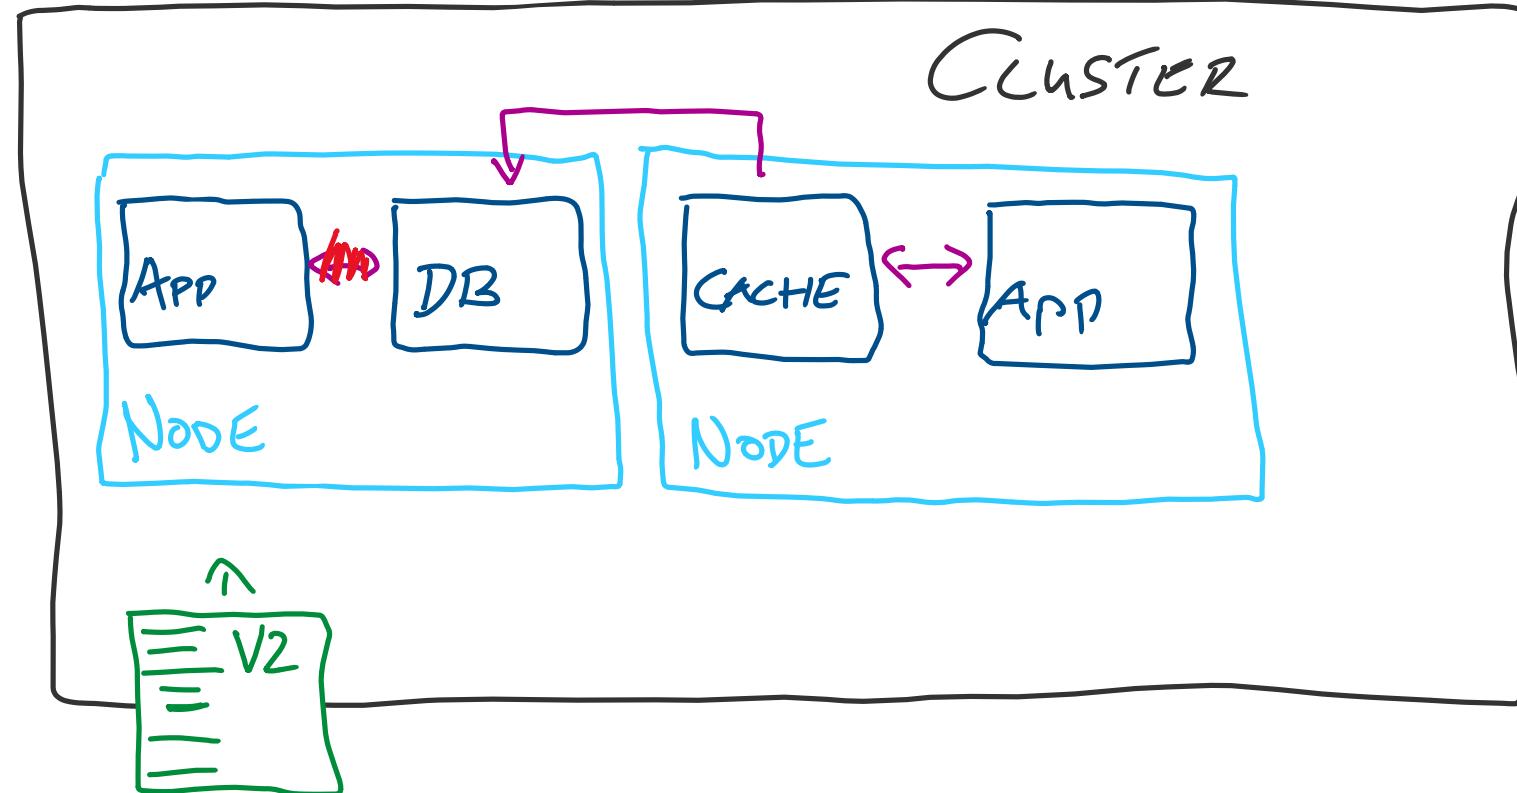


Container-optimiertes Linux?



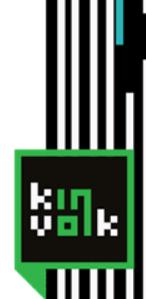
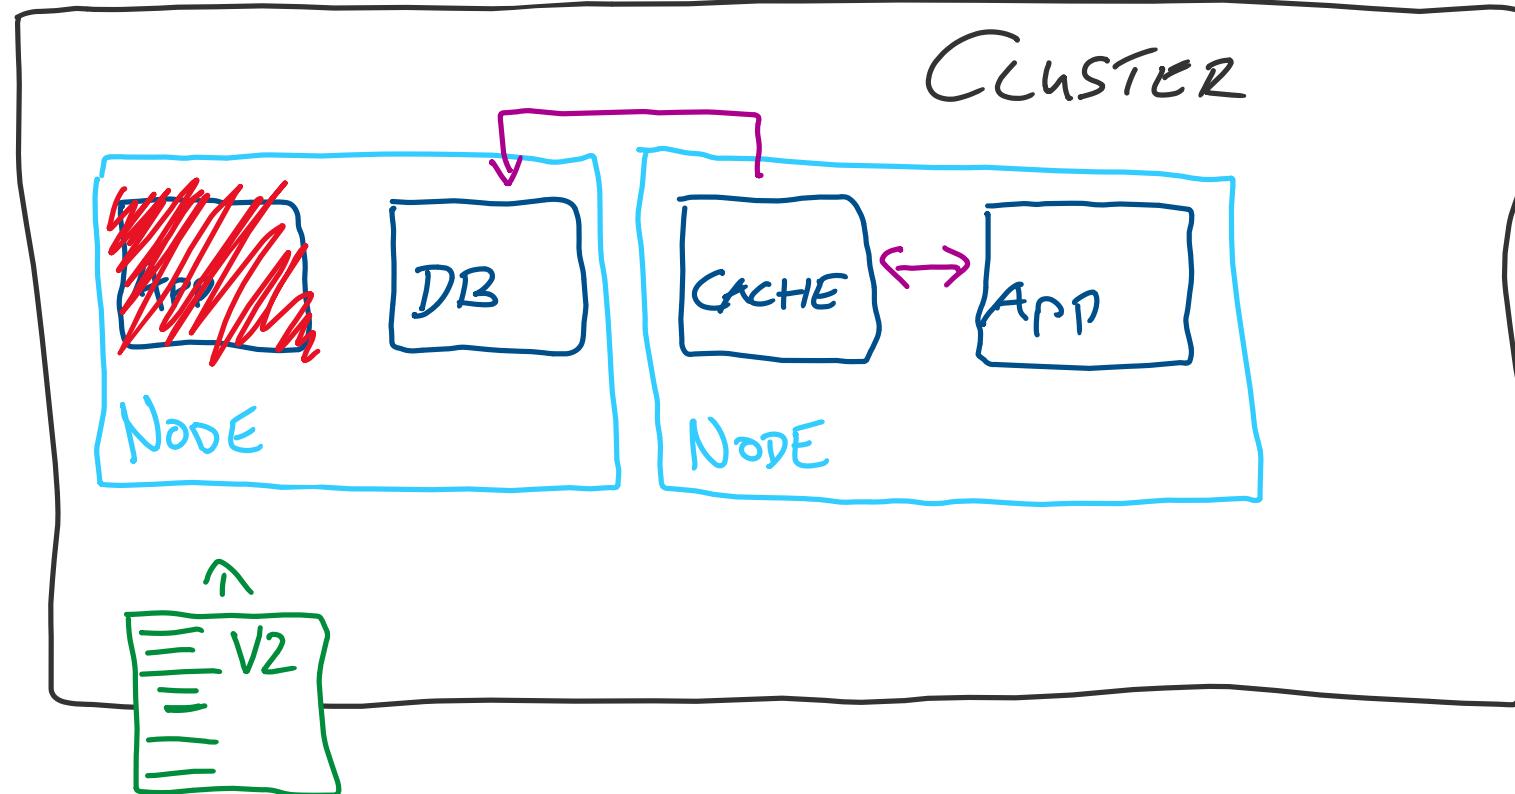


Container-optimiertes Linux?



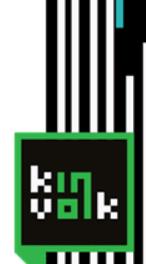
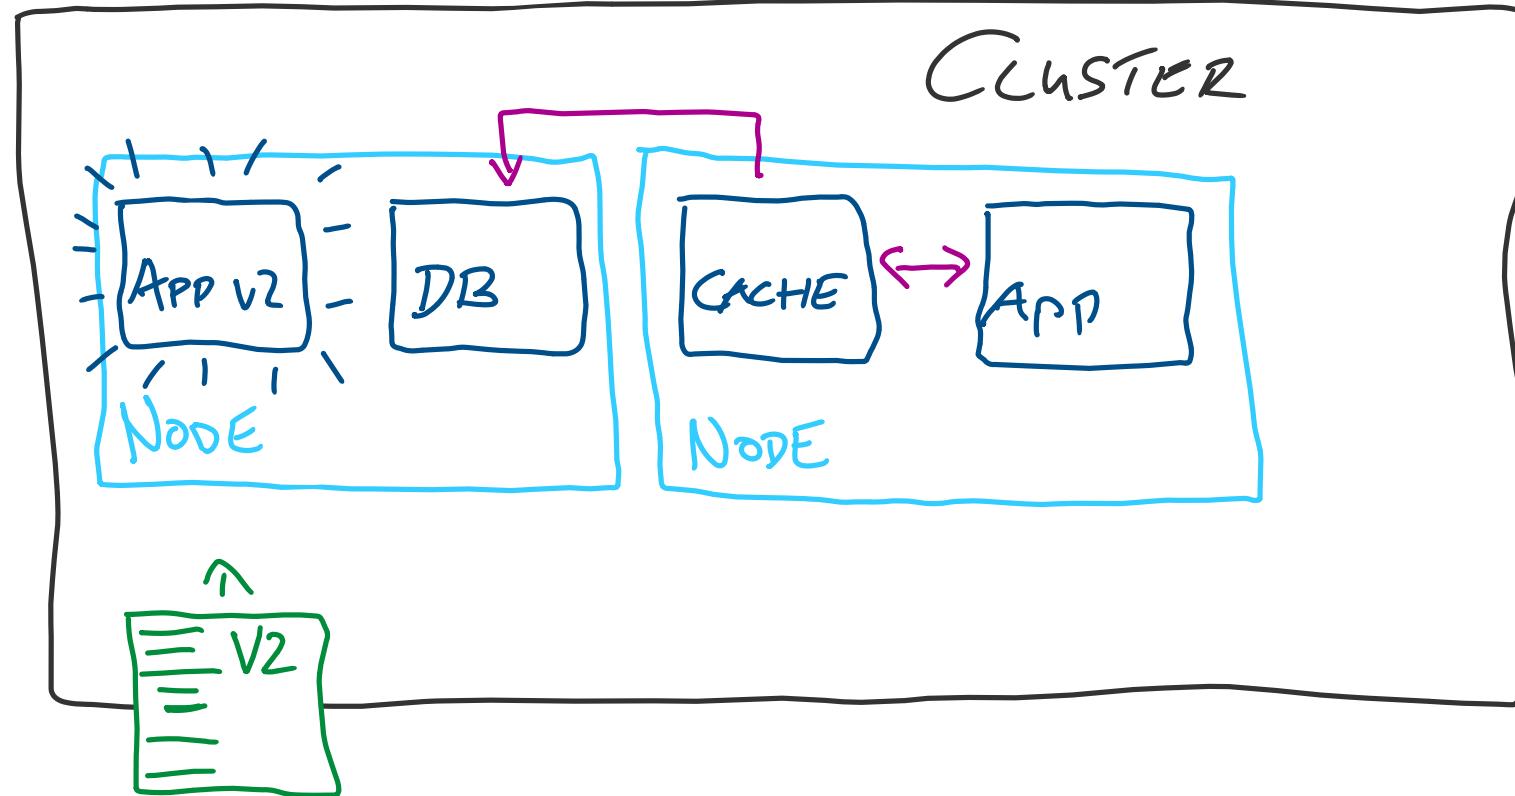


Container-optimiertes Linux?



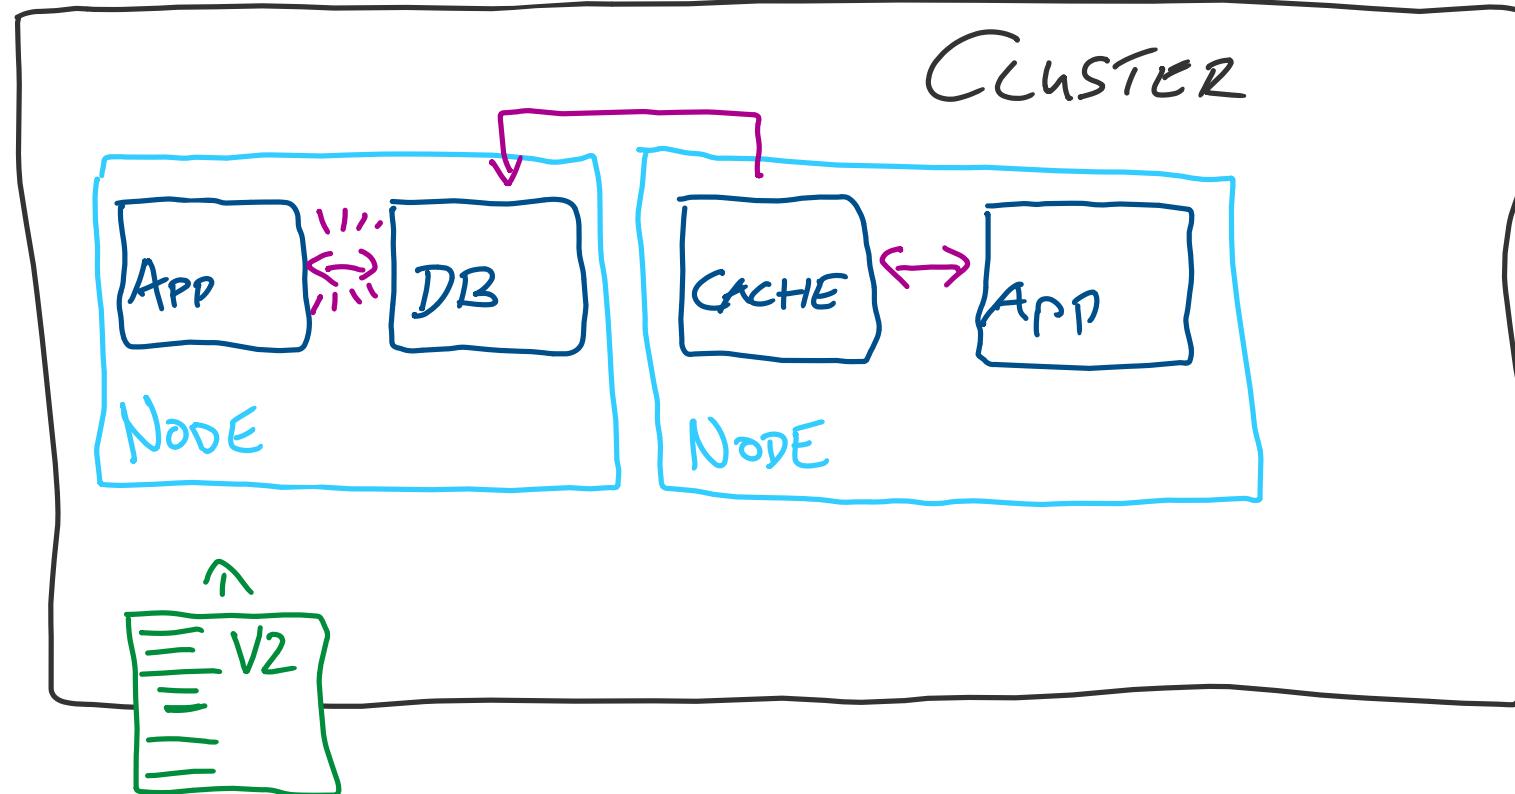


Container-optimiertes Linux?



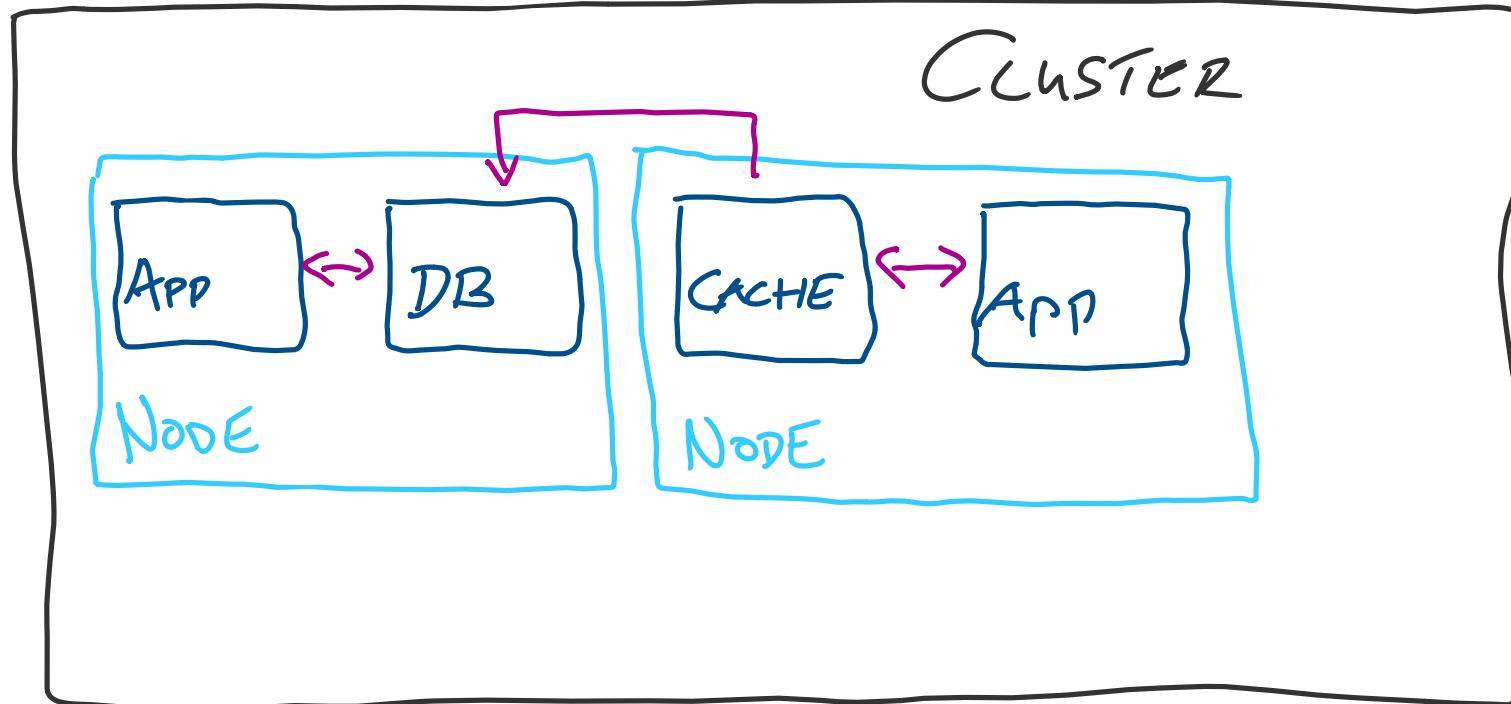


Container-optimiertes Linux?



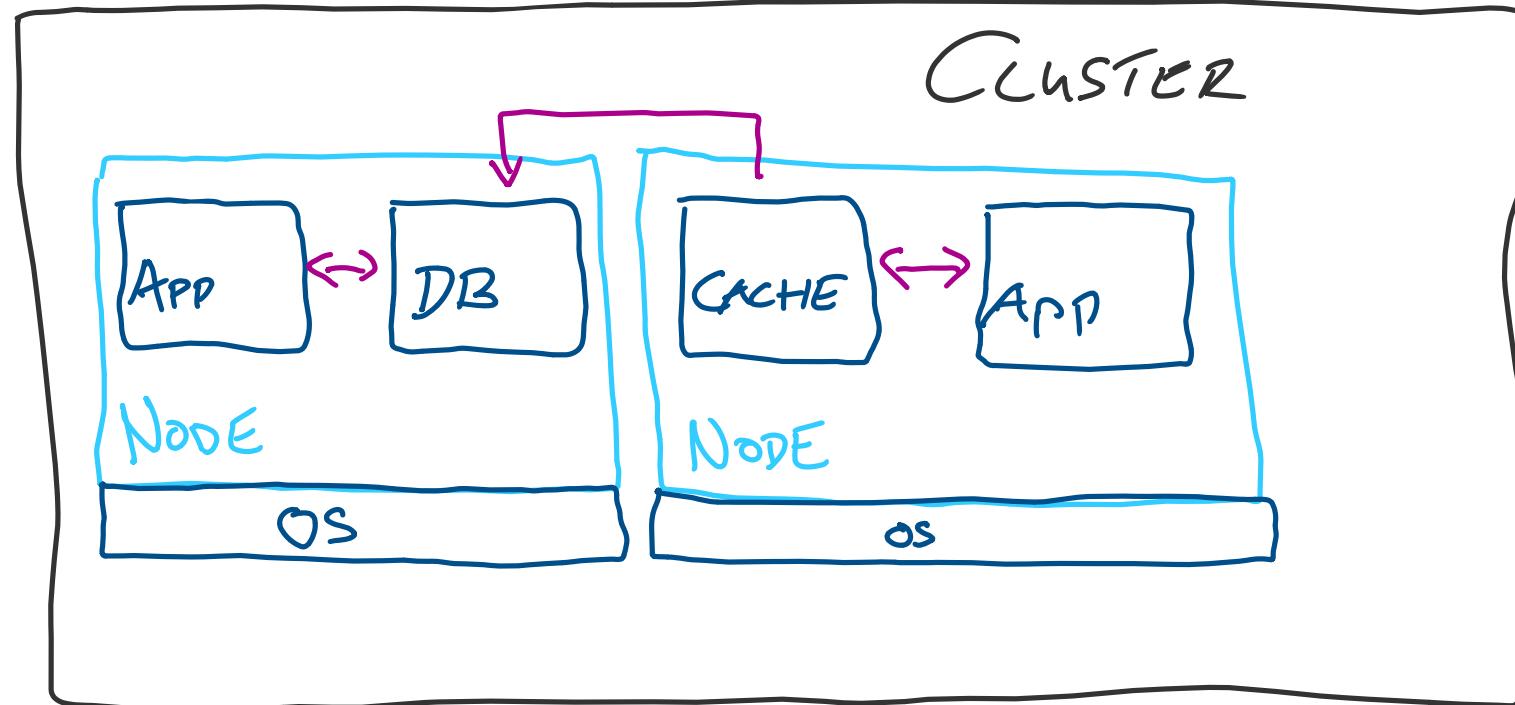


Container-optimiertes Linux?



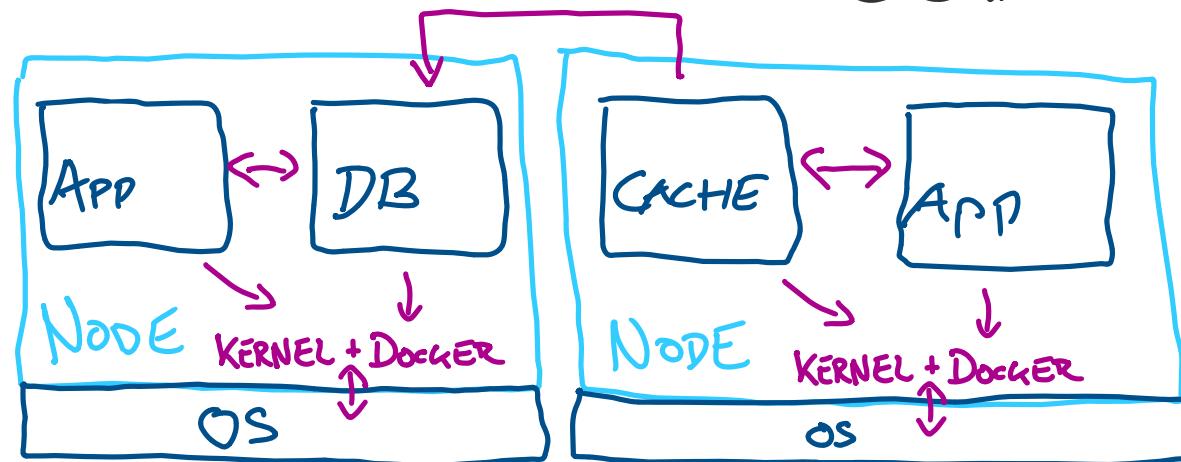


Container-optimiertes Linux?



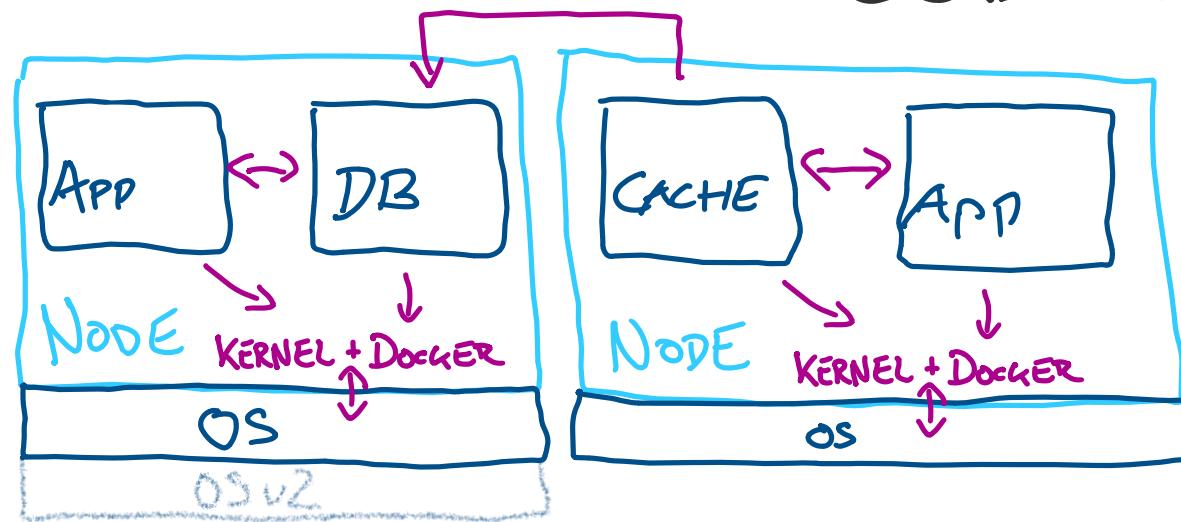
Container-optimiertes Linux?

CLUSTER

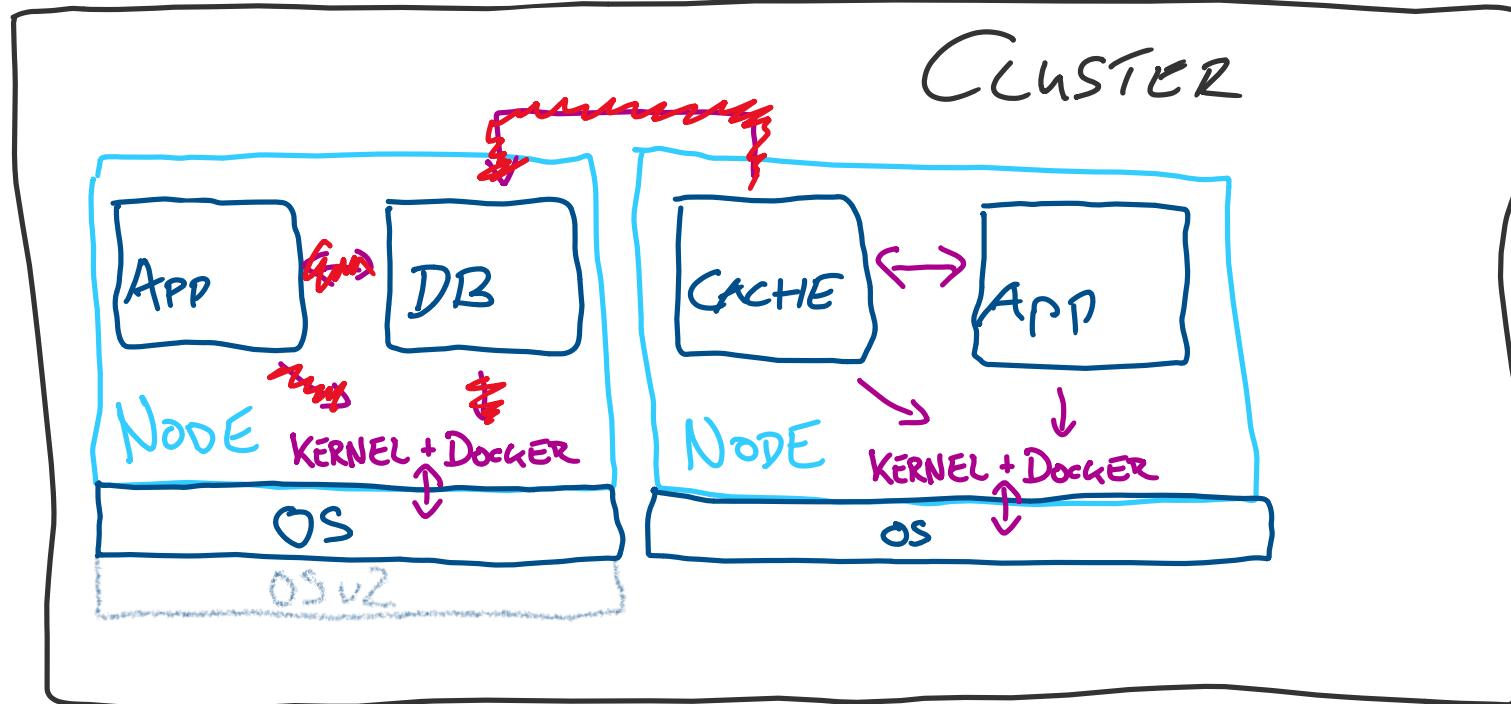


Container-optimiertes Linux?

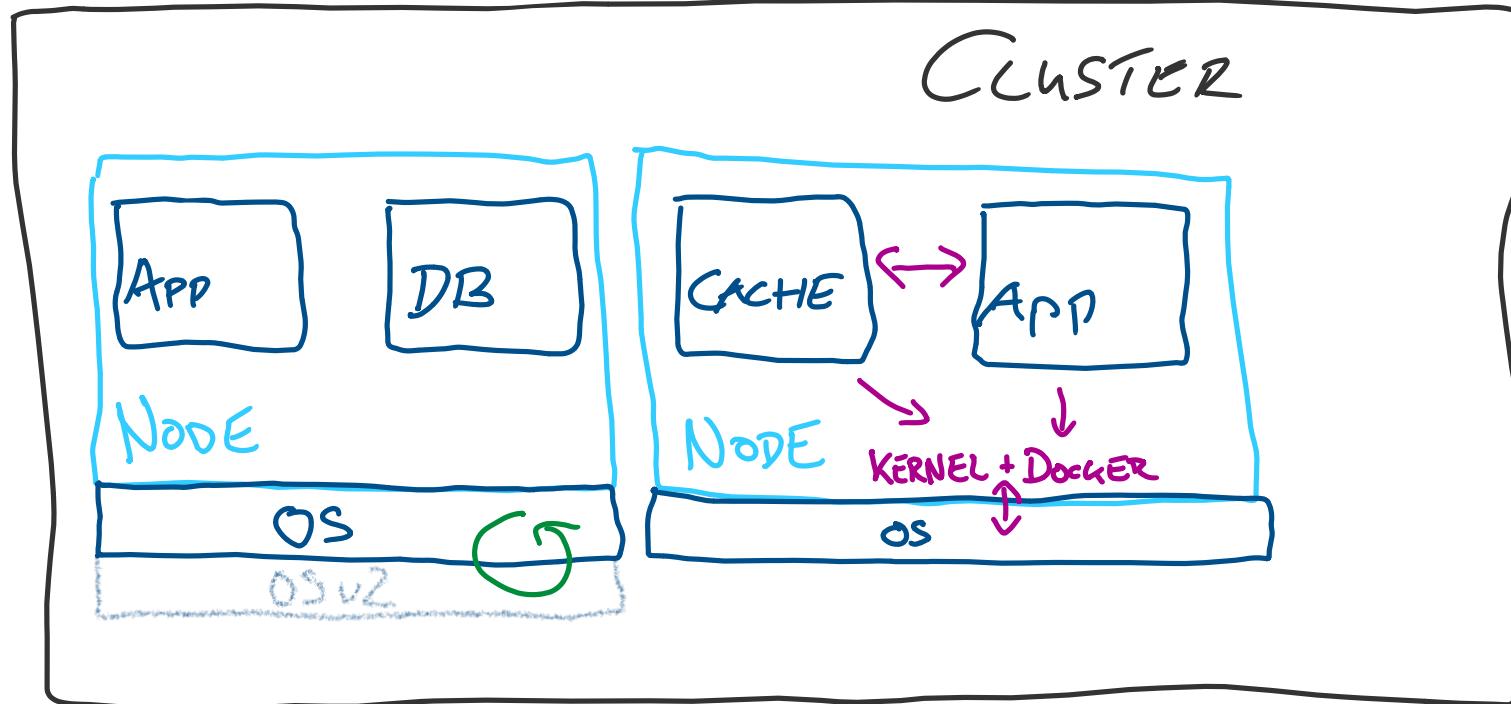
CLUSTER



Container-optimiertes Linux?

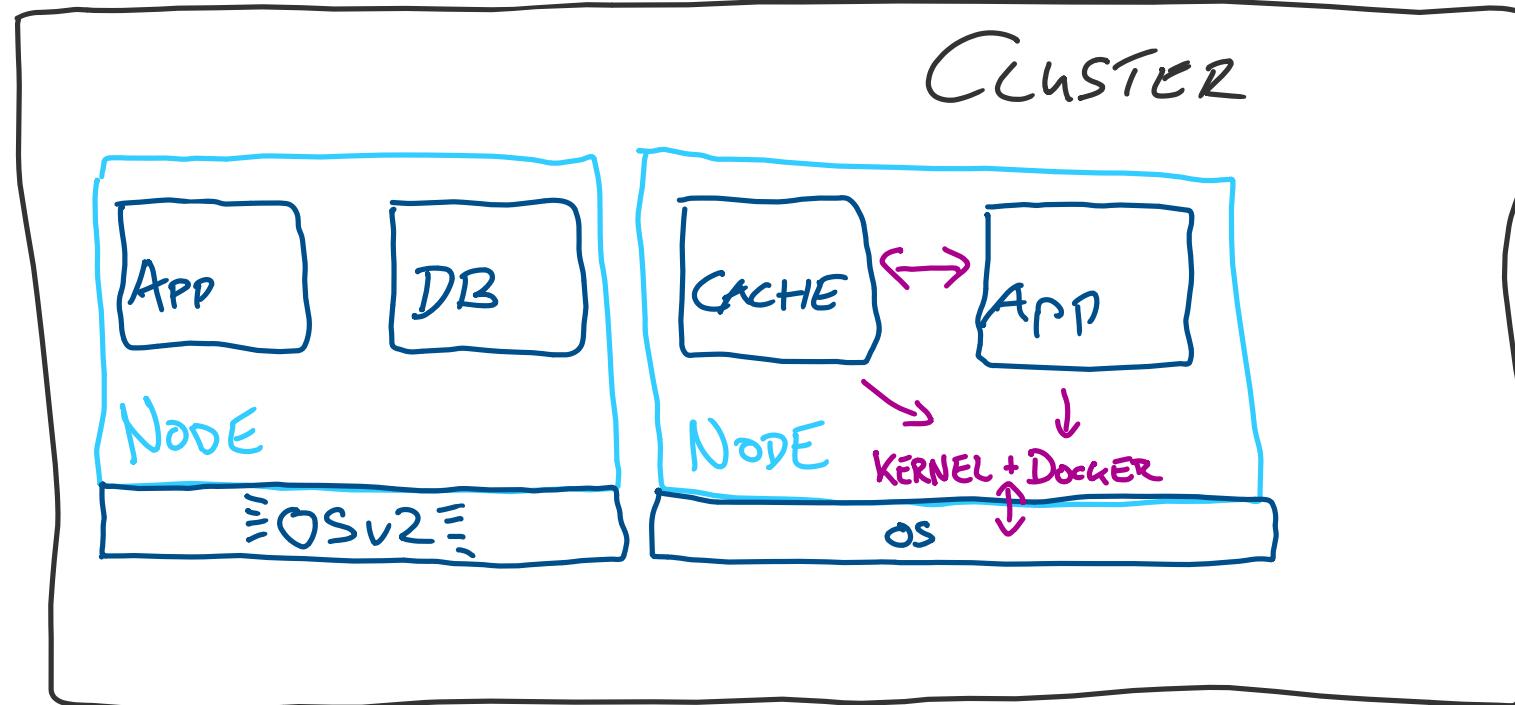


Container-optimiertes Linux?



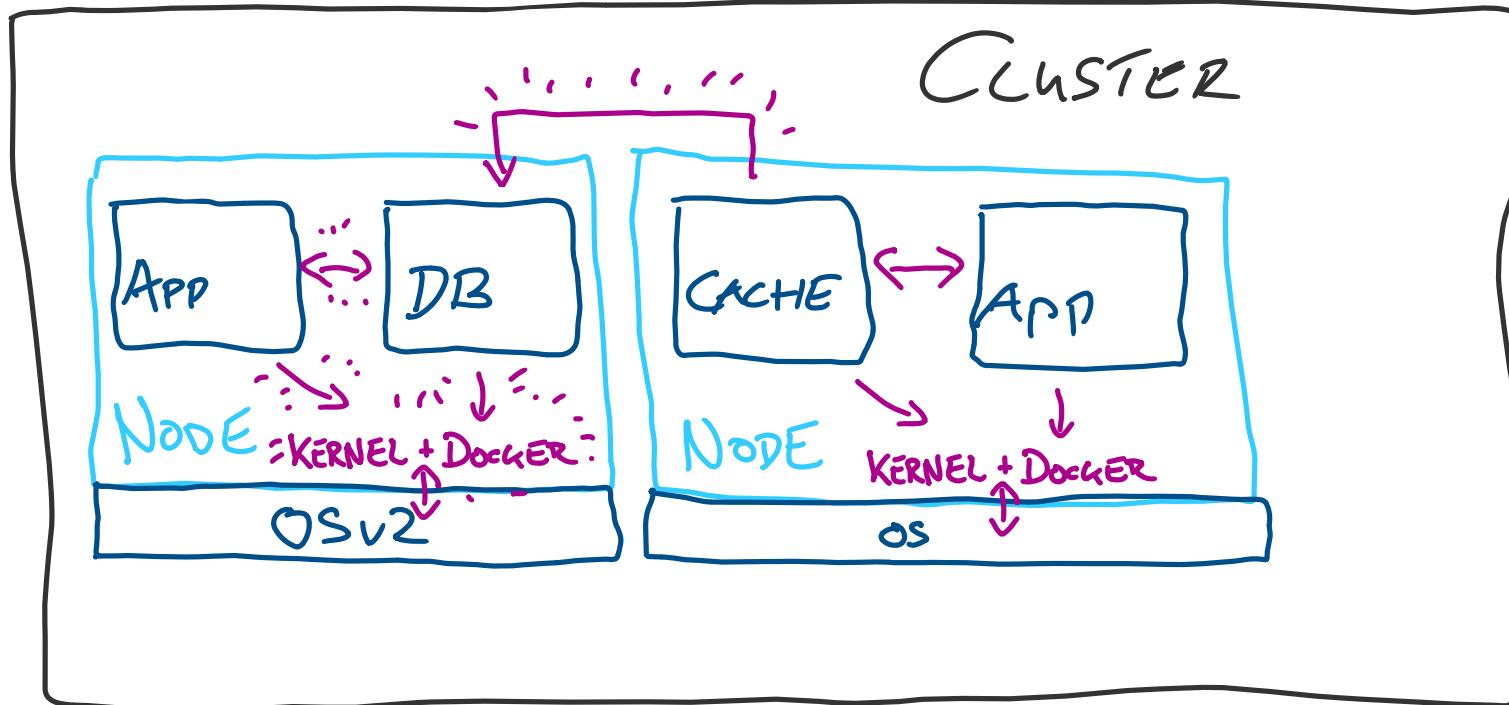


Container-optimiertes Linux?





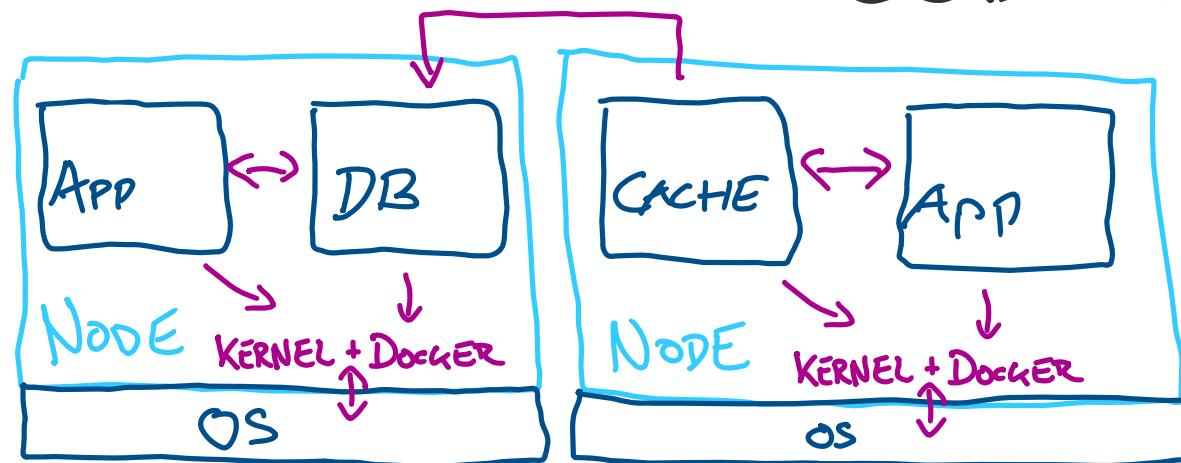
Container-optimiertes Linux?





Container-optimiertes Linux?

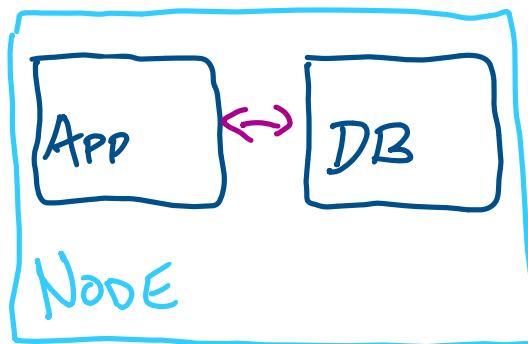
CLUSTER



Container-optimiertes Linux?

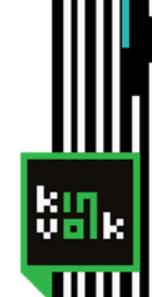
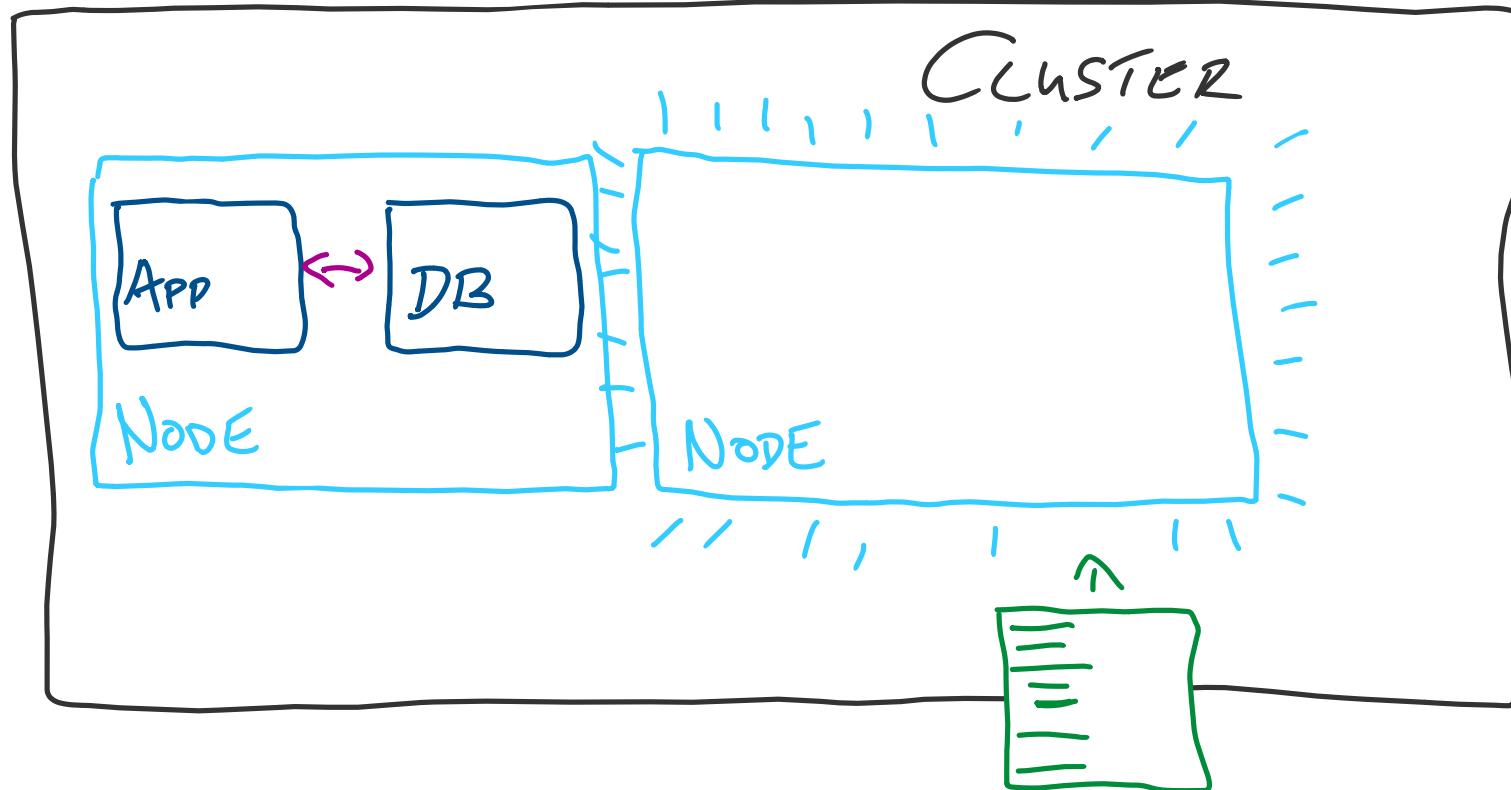


Cluster





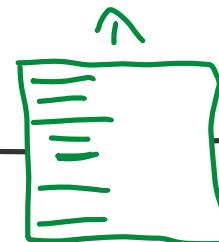
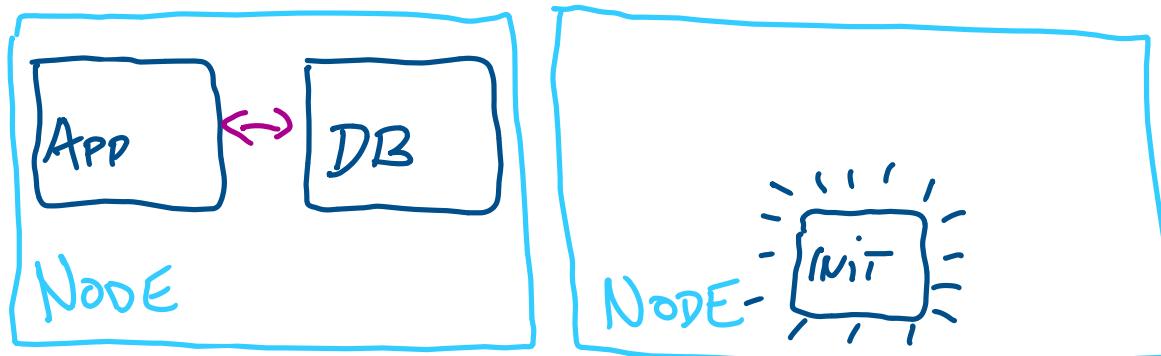
Container-optimiertes Linux?





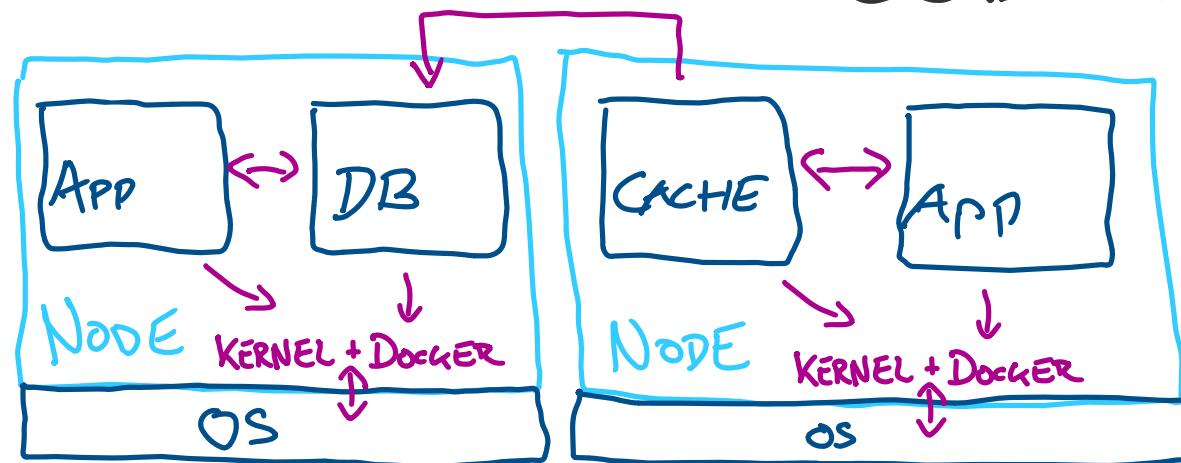
Container-optimiertes Linux?

CLUSTER



Container-optimiertes Linux?

CLUSTER





Container-optimiertes Linux - Basis

Isolation der Apps untereinander und zum OS

Wohldefinierte Schnittstellen zwischen App und OS

Deklarative, automatisierbare Installation

Konfiguration bei Installation, ohne dynamische Änderungen

Automatische updates, kontrolliertes roll-(out|back)

Bereistellen, wechseln, aktivieren (oder zurückrollen)

==> Das OS als zustandslose Applikation.





Isolation

Keine Abhängigkeiten zwischen OS und Apps

Keine gemeinsamen Bibliotheken, nur Kernel und Docker.

Saubere API zwischen OS und Apps

Netzwerk, Storage etc. in Konfiguration dokumentiert.

OS-Updates (und roll-backs) ohne Seiteneffekte

Gut testbar über z.B. “Canary”-Muster.





Automatisierte Installation

Deklarative Konfiguration (per Ignition)

Aktiviert bei Installation.

Vollständig automatisierbares roll-out

Umfassende Konfigurationsoptionen

Einfache Integration in Orchestrierung

GO-bindings, [Terraform](#) provider, CAPI support, etc.



```

passwd:
users:
- name: caddy
  no_create_home: true
  groups: [ docker ]

storage:
files:
- path: /srv/www/html/index.html
  mode: 0644
  user:
    name: caddy
  group:
    name: caddy
  contents:
    inline: |
      <html><body align="center">
        <h1>Hallo, FroSCon!</h1>
        
      </body></html>
- path: /srv/www/html/froscon.png
  mode: 0644
  user:
    name: caddy
  group:
    name: caddy
  contents:
    local: froscon_logo_print_color.png

systemd:
units:
- name: froscon-demo-webserver.service
  enabled: true
  contents: |
    [Unit]
    Description=FroSCon example static web server
    After=docker.service
    Requires=docker.service
    [Service]
    User=caddy
    TimeoutStartSec=0
    ExecStartPre=-/usr/bin/docker rm --force caddy
    ExecStart=docker run -i -p 80:80 --name caddy \
      -v /srv/www/html:/usr/share/caddy \
      docker.io/caddy caddy file-server \
      --root /usr/share/caddy --access-log
    ExecStop=/usr/bin/docker stop caddy
    Restart=always
    RestartSec=5s
    [Install]
    WantedBy=multi-user.target

```



```
passwd:
```

```
users:
```

```
  name: caddy  
  no_create_home: true  
  groups: [ docker ]
```

```
storage:
```

```
  files:
```

```
- path: /srv/www/html/index.html
```

```
  mode: 0644
```

```
  user:
```

```
    name: caddy
```

```
  group:
```

```
    name: caddy
```

```
  contents:
```

```
    inline: |
```

```
      <html><body align="center">  
      <h1>Hallo, FroSCon!</h1>  
        
      </body></html>
```

```
- path: /srv/www/html/froscon.png
```

```
  mode: 0644
```

```
  user:
```

```
    name: caddy
```

```
  group:
```

```
    name: caddy
```

```
  contents:
```

```
    local: froscon_logo_print_color.png
```

```
systemd:
```

```
  units:
```

```
- name: froscon-demo-webserver.service
```

```
  enabled: true
```

```
  contents: |
```

```
    [Unit]
```

```
    Description=FrOSCon example static web server
```

```
    After=docker.service
```

```
    Requires=docker.service
```

```
    [Service]
```

```
    User=caddy
```

```
    TimeoutStartSec=0
```

```
    ExecStartPre=-/usr/bin/docker rm --force caddy
```

```
    ExecStart=docker run -i -p 80:80 --name caddy \
```

```
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log
```

```
    ExecStop=/usr/bin/docker stop caddy
```

```
    Restart=always
```

```
    RestartSec=5s
```

```
    [Install]
```

```
    WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]  
  
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    inline: |  
      <html><body align="center">  
        <h1>Hallo, FroSCon!</h1>  
          
      </body></html>  
- path: /srv/www/html/froscon.png  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    local: froscon_logo_print_color.png
```

```
systemd:  
units:  
- name: froscon-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=FroSCon example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 --name caddy \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop caddy  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



//

```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]  
  
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    inline: |  
      <html><body align="center">  
        <h1>Hallo, FroSCon!</h1>  
          
      </body></html>  
- path: /srv/www/html/froscon.png  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    local: froscon_logo_print_color.png
```

```
systemd:  
units:  
- name: froscon-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=FroSCon example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 --name caddy \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop caddy  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



```
passwd:
```

```
users:
```

```
- name: caddy  
  no_create_home: true  
  groups: [ docker ]
```

```
storage:
```

```
files:
```

```
- path: /srv/www/html/index.html  
  mode: 0644
```

```
  user:
```

```
    name: caddy
```

```
  group:
```

```
    name: caddy
```

```
  contents:
```

```
    inline: |  
      <html><body align="center">  
        <h1>Hallo, FroSCon!</h1>  
          
      </body></html>
```

```
- path: /srv/www/html/froscon.png
```

```
  mode: 0644
```

```
  user:
```

```
    name: caddy
```

```
  group:
```

```
    name: caddy
```

```
  contents:
```

```
    local: froscon_logo_print_color.png
```



```
systemd:
```

```
units:
```

```
- name: froscon-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=FrOSCon example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 --name caddy \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop caddy  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```

```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]  
  
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    inline: |  
      <html><body align="center">  
        <h1>Hallo, FroSCon!</h1>  
          
      </body></html>  
- path: /srv/www/html/froscon.png  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    local: froscon_logo_print_color.png
```

```
systemd:  
units:  
- name: froscon-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=FroSCon example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 --name caddy \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop caddy  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]  
  
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    inline: |  
      <html><body align="center">  
        <h1>Hallo, FroSCon!</h1>  
          
      </body></html>  
- path: /srv/www/html/froscon.png  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    local: froscon_logo_print_color.png
```

```
systemd:  
units:  
- name: froscon-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=FroSCon example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 --name caddy \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop caddy  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



```
passwd:
```

```
users:
```

```
- name: caddy  
  no_create_home: true  
  groups: [ docker ]
```

```
storage:
```

```
files:
```

```
- path: /srv/www/html/index.html
```

```
  mode: 0644
```

```
  user:
```

```
    name: caddy
```

```
  group:
```

```
    name: caddy
```

```
  contents:
```

```
    inline: |
```

```
      <html><body align="center">  
      <h1>Hallo, FroSCon!</h1>  
        
      </body></html>
```

```
- path: /srv/www/html/froscon.png
```

```
  mode: 0644
```

```
  user:
```

```
    name: caddy
```

```
  group:
```

```
    name: caddy
```

```
  contents:
```

```
    local: froscon_logo_print_color.png
```



```
systemd:
```

```
units:
```

```
- name: froscon-demo-webserver.service
```

```
  enabled: true
```

```
  contents: |
```

```
    [Unit]
```

```
    Description=FroSCon example static web server
```

```
    After=docker.service
```

```
    Requires=docker.service
```

```
    [Service]
```

```
    User=caddy
```

```
    TimeoutStartSec=0
```

```
    ExecStartPre=-/usr/bin/docker rm --force caddy
```

```
    ExecStart=docker run -i -p 80:80 --name caddy \
```

```
        -v /srv/www/html:/usr/share/caddy \
```

```
        docker.io/caddy caddy file-server \
```

```
        --root /usr/share/caddy --access-log
```

```
    ExecStop=/usr/bin/docker stop caddy
```

```
    Restart=always
```

```
    RestartSec=5s
```

```
    [Install]
```

```
    WantedBy=multi-user.target
```



Image-basiertes Betriebssystem

Zustandslose Installation

Betriebssystem-Austausch ohne Seiteneffekte

Unveränderbare (read-only) OS-Partition

Binaries nicht änderbar (aus Versehen oder Absicht)

Verifizierbar / Attestierbar

OS ändert sich nicht und kann fortlaufend attestiert werden





Updates

Atomare Updates

Download + Bereitstellen im Hintergrund

Reboot wechselt in die neue Version

Automatisierte, konfigurierbare roll-backs

Wichtige Dienste können Rollback auslösen

Gesteuertes, kontrolliertes Ausrollen

Cluster-weite Reboot-Koordination





Updates – Stabilisierungsprozess

Alpha -> Beta -> Stable (->LTS)

Alpha für Entwickler, Beta für canaries, und Stable für Prod

Weitreichende Testszenarien für jedes Release

Komplexe Tests, z.B. Kubernetes workloads, CNI, etc.

Alle Releases gehen stets durch alle Tests

Beta-Canaries für Workload-spezifische Tests





Große Cluster / “at scale”

(Community) LTS

Wenn jegliche Wartung Schmerzen bereitet.

FOSS update server ([Nebraska](#))

Updates selbst hosten und roll-out an Gruppen managen

Nutzerzentriert, Community-getrieben, Open Source

Schau zu, mach mit, und bring Dich ein!





Community-getrieben

Historisch stets FOSS, aber firmengetrieben

Direkt nach dem “Friendly Fork”

Umorientierung auf Community-Projekt

Beiträge, Koordination, Steuerung

Microsoft unterstützt als Sponsor

Öffentliche Verpflichtung bei Aquise



Microsofts Verpflichtung zur Unterstützung der Community



<https://azure.microsoft.com/en-us/blog/microsoft-acquires-kinvolk-to-accelerate-containeroptimized-innovation/>

“Flatcar Container Linux has a sizeable community of users on Azure, as well as other clouds, and on-premises. We know the CoreOS community has been on a winding journey over the years—we want to assure the Flatcar community that Microsoft and the Kinvolk team will continue to collaborate with the larger Flatcar community on the evolution of Flatcar Container Linux. Microsoft is committed to Flatcar Container Linux community development and will invest in working with the Flatcar community to create a growth path forward together. We’ll have our first meeting with the community within the coming weeks and invite anyone interested to attend and join the conversation.”



Microsofts Verpflichtung zur Unterstützung der Community



<https://azure.microsoft.com/en-us/blog/microsoft-acquires-kinvolk-to-accelerate-containeroptimized-innovation/>

“Flatcar Container Linux has a sizeable community of users on Azure, as well as other clouds, and on-premises. We know the CoreOS community has been on a winding journey over the years—we want to assure the Flatcar community that Microsoft and the Kinvolk team will continue to collaborate with the larger Flatcar community on the evolution of Flatcar Container Linux. Microsoft is committed to Flatcar Container Linux community development and will invest in working with the Flatcar community to create a growth path forward together. We’ll have our first meeting with the community within the coming weeks and invite anyone interested to attend and join the conversation.”



Arbeit und Kommunikation öffentlich



Tagesgeschäft

Matrix, Slack – Maintainer-Interaktion, Release-Chat

Kurzzeit- und Langzeitplanung, Feedback

“Office hours” und “Release planning” meetings

Projekt-Boards

“Bug smashing” und “docs writing” - Tage

Mit Livestream / Videocall und Matrix chat!



Was ist grad so los? - Projekt-Boards



Implementierung

Tasks, Bugs, Features

Release

Wann ist Feature X verfügbar?

Roadmap

Epics und subtasks

A screenshot of several GitHub Project Boards side-by-side. From left to right, they represent different stages: Upcoming / Backlog, Blocked, Planned / To Do, In Progress, and Testing / In Review. Each board lists specific tasks or bugs with their descriptions and labels. For example, the 'Planned / To Do' board has a task for 'Flatcar #645: [RFE] Keep packages in portage-stable updated'. The 'In Progress' board shows tasks like 'Flatcar #762: [RFE] Add supply chain provenance releases' and 'Flatcar #646: [RFE] Move out customizations out of portage-stable'. The 'Testing / In Review' board includes tasks such as 'Flatcar #790: [RFE] Trigger the scripts CI GitHub Actions workflow for coreos-overlay/portage' and 'Flatcar #799'. The 'Release' boards show scheduled releases for 2022-07-18 and 2022-06-20, each with its own list of tasks and labels like 'kind/feature', 'kind/test', and 'kind/rhsmapi'.

Upcoming / Backlog	Blocked	Planned / To Do	In Progress	Testing / In Review	Release: 2022-07-18	Release: 2022-06-20
Flatcar #645: [RFE] Handle configuration options in coreos-overlay	Flatcar #646: [RFE] Add supply chain provenance releases	Flatcar #762: [RFE] Keep packages in portage-stable updated	Flatcar #762: [RFE] Keep packages in portage-stable updated	Flatcar #790: [RFE] Trigger the scripts CI GitHub Actions workflow for coreos-overlay/portage	Draft Alpha 3304.0 Beta 3277.1.0 Stable 3227.0 LTS-2022.3033.3 LTS-2021.2605.30.1	Draft Alpha 3277.0.0 Beta 3227.1.1 Stable 3139.2.3 LTS-2022.3033.3.2 LTS-2021.2605.29.1





Beiträge zu anderen Projekten

Gentoo

Package updates, Stabilisierungen, bug fixes

Ignition, Butane

Erweiterungen und bug fixes

Linux Kernel, Cilium, Falco

Bug fixes, Interoperationsverbesserungen



Mach mit!



Gib Feedback, melde Bugs, oder erstelle Features

Deine Meinung ist uns wichtig!

Bug smashing / doc writing - Tage

Mit Livestream und live chat

Package updates oder OS Image - Erweiterungen

Werde Flatcar-Maintainer!





Bleib in Verbindung

Office hours

Jeder zweite Dienstag im Monat um 17:30 Uhr (CE[S]T)

Matrix

Tägliche Kommunikation mit Maintainern und Entwicklern

Slack

siehe Matrix ;)



Demo time!

Wir nutzen die bereits gezeigte Konfiguration
(Nicht ganz. Wir schalten update reboots ab.)

Wir machen auch ein OS-Update.
Einfach, schnell und langweilig. Wie ein Lichtschalter.

Demo-Quellen sind hier verfügbar:

<https://github.com/flatcar-linux/flatcar-demos/tree/main/froscon-2022-08-20>



Thank you 🚗



Website - <https://www.flatcar.org>

office hours - <https://github.com/flatcar-linux/Flatcar/#monthly-office-hours-and-release-planning>

Mach mit - <https://github.com/flatcar-linux/Flatcar>

Melde Dich - <https://app.element.io/#/room/#flatcar:matrix.org>
- <https://kubernetes.slack.com/archives/C03GQ8B5XNJ>

