



# "Day 2 Ops"

Linux for Kubernetes and Container Workloads

Hello, I'm

*Thilo*



**Thilo Fromm**

Flatcar Maintainer

Github: [t-lo](https://github.com/t-lo)

Mastodon: [@thilo@fromm.social](https://mastodon.social/@thilo@fromm.social)

Email: [thilofromm@microsoft.com](mailto:thilofromm@microsoft.com)

"Day 2"?

## Visionary Ideas / Brainstorming Phase

Large parts don't exist yet, focus on Key Feature(s)

Very limited exposure to users ("closed Alpha")

Code has "Proof of Concept" quality

Operations are manual, no time for automation

Things constantly break



Project goes live / emerges from "stealth mode"

Main features are implemented

Service meets reality as users on-board

Spotty, gappy Automation

Much firefighting

"Test in Production", but not in a good way



Project is Mature, in steady state

Deterministic feature lifecycle

No user-visible service degradations / interruptions

Operations are well-defined, extensively automated

Service impacts are well managed

"Boring", in a good way







"Day 2" OS

- Deployments are fully automated and reproducible
- Reliable update / lifecycle management
- Robust and safe, proven support of your target environment
- Solid configuration management / inventory management (SBOM)
- Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*





# Fully Automated Deployments



# Kubernetes automates service deployments



Configure

->

Deploy

->

Operate

Sane defaults  
- no boiler plate.

Integration in  
cluster environment.

Customisation.

`kubectl apply`

Automated

- Config management
- Version management
- Lifecycle management

my-service.yaml

# Flatcar automates Node deployments



Configure

->

Deploy

->

Operate

Sane defaults  
- no boiler plate.

Integration in  
ops environment.

Customisation.

Butane-config.yaml

Custom data,  
User data,  
Http, or  
[i]PXE

ignition.json

Automated

- Self-configuration
- Unattended updates

# Declarative configuration, before provisioning



```
1 variant: flatcar
2 version: 1.0.0
3
4 passwd:
5   users:
6     - name: caddy
7       no_create_home: true
8       groups: [ docker ]
9
10 storage:
11   files:
12     - path: /srv/www/html/index.html
13       mode: 0644
14       user:
15         name: caddy
16       contents:
17         inline: |
18           <html><body align="center">
19             <h1>Hallo FrOSCon 2024!</h1>
20             </html>
22     - path: /srv/www/html/froscon_logo_print_color.png
23       mode: 0644
24       user:
25         name: caddy
26       group:
27         name: caddy
28       contents:
29         local: froscon_logo_print_color.png
30
31 systemd:
32   units:
33     - name: update-engine.service
34       mask: true
35     - name: Froscon-demo-webserver.service
36       enabled: true
37       contents: |
38         [Unit]
39         Description=FrOSCon example static web server
40         After=docker.service
41         Requires=docker.service
42         [Service]
43         User=caddy
44         TimeoutStartSec=0
45         ExecStartPre=/usr/bin/docker rm --force caddy
46         ExecStart=/usr/bin/docker run -i -p 80:80 --name caddy \
47             -v /srv/www/html:/usr/share/caddy \
48             docker.io/caddy caddy file-server \
49             --root /usr/share/caddy --access-log
50         ExecStop=/usr/bin/docker stop caddy
51         Restart=always
52         RestartSec=5s
53         [Install]
54         WantedBy=multi-user.target
```

# Provisioning Demo

Reproducible, Cloud Agnostic

Zero touch, Automatable

Supports complex deployments (See e.g. the [Flatcar Jitsi server](#))

No config drift

Integrates into existing automation (Go, Terraform/OpenTOFU, ClusterAPI)





Fully automated, easy to deploy

Reliable update / lifecycle management

Robust and safe, proven support of your target environment

Solid configuration management / inventory management (SBOM)

Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*



# Reliable Update / Lifecycle Management



# Container Apps are isolated

## No Dependencies to OS Libraries

Container images are self-sufficient

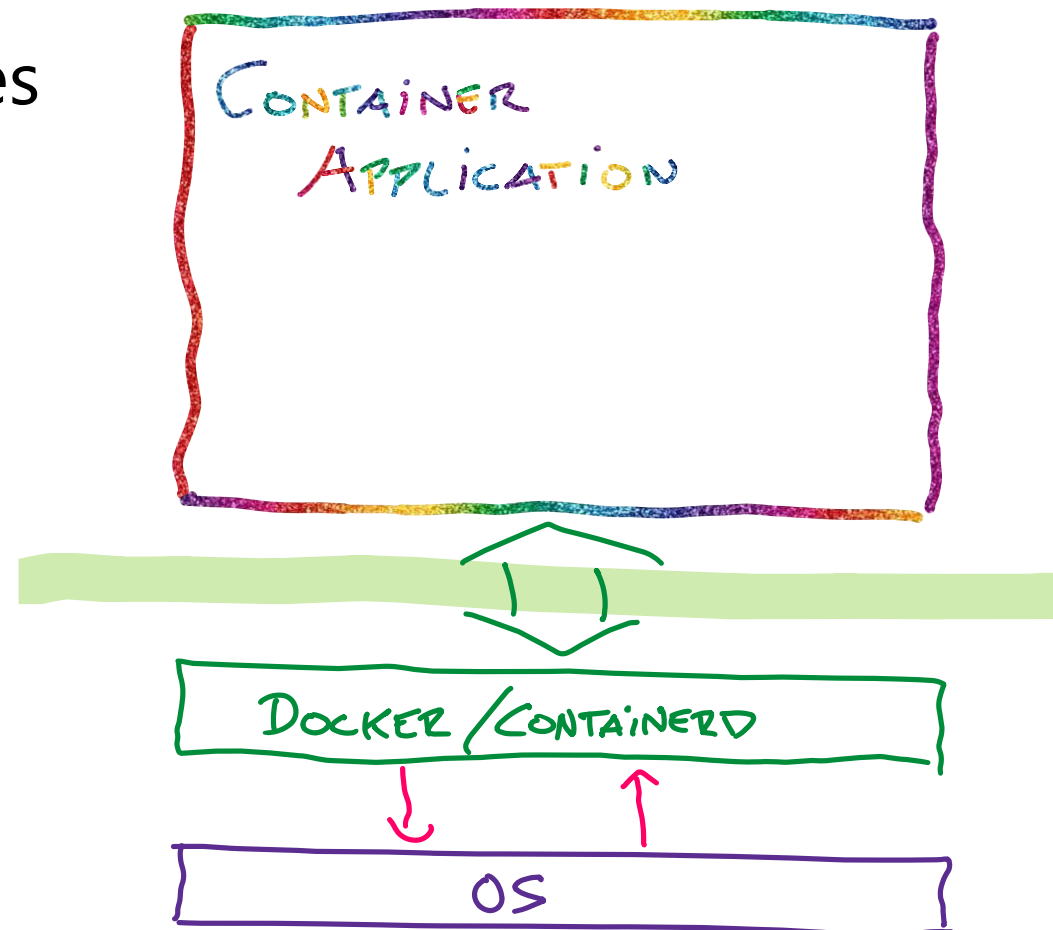
## No Dependencies to OS Tools and Services

Only require Container Runtime

## No Dependencies to OS Configuration

App config cleanly separated from OS

## ➔ Portable Applications



# And so is the OS!

Well-defined interfaces OS  $\leftrightarrow$  App

Very few components, easy to test thoroughly

No other inter-dependencies

Container apps isolate from the OS

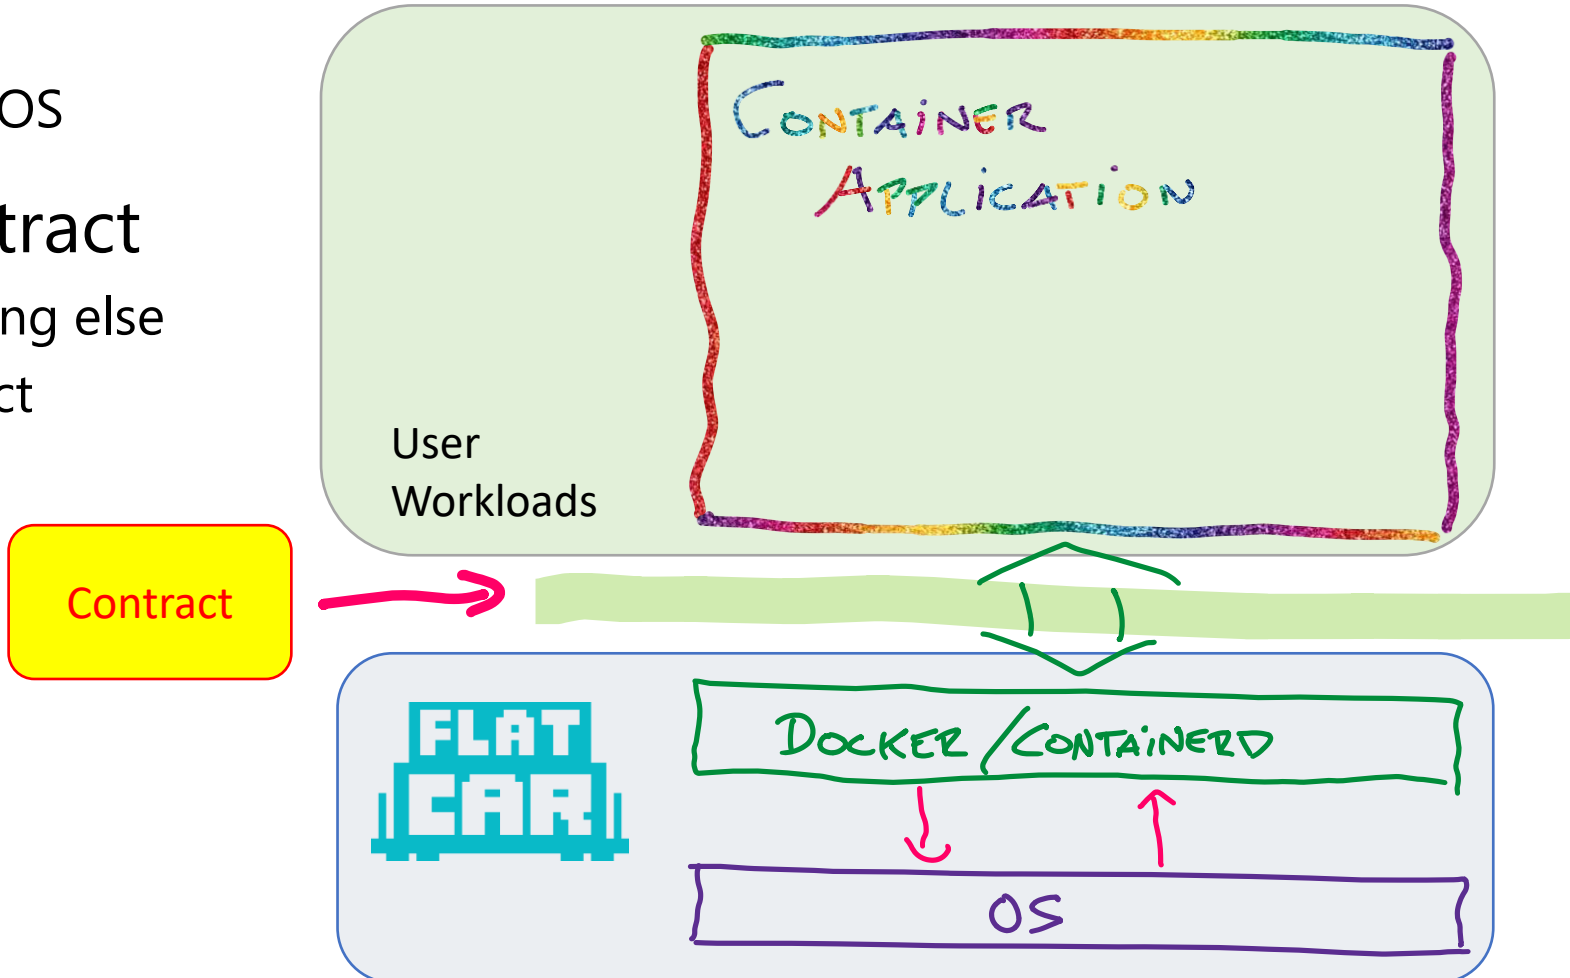
Runtime + Kernel == Contract

App relies on contract and nothing else

OS guarantees and fulfils contract

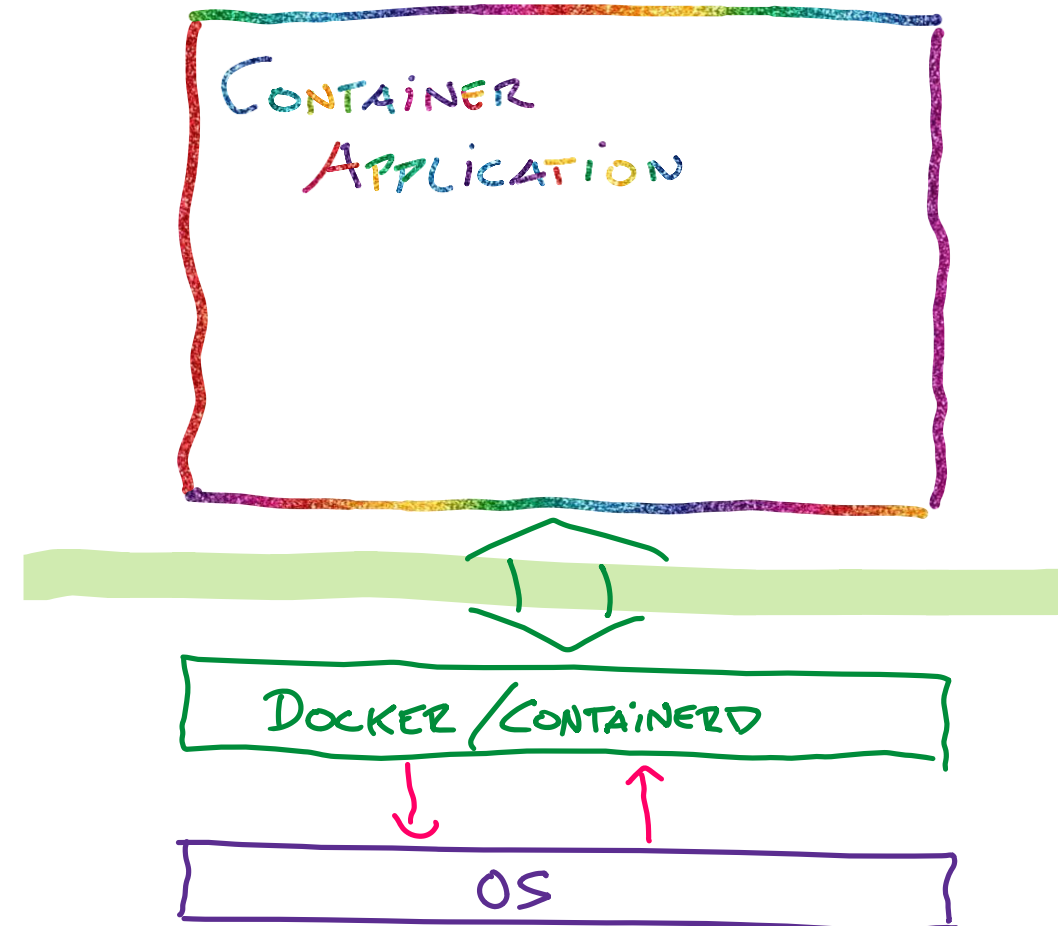
➔ Interchangeable OS

Just uphold the contract!



# Staying up to date

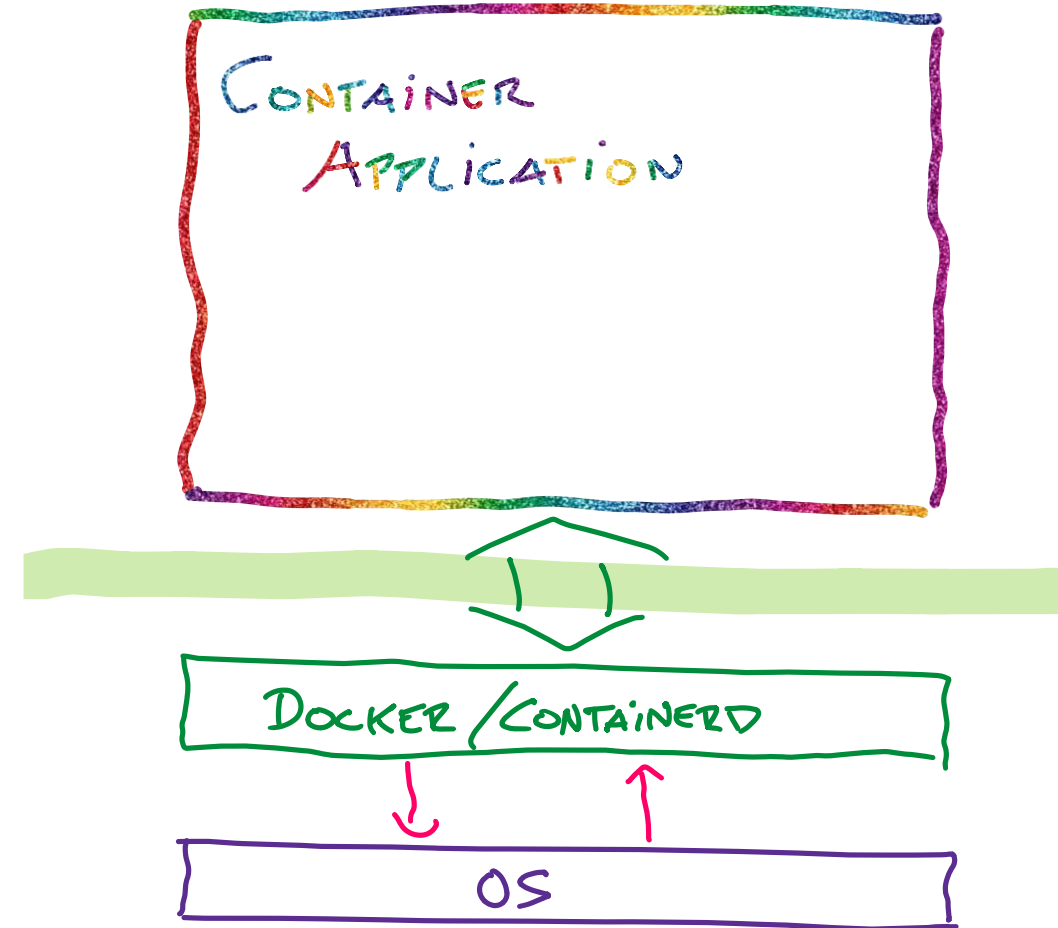

## Atomic In-Place Updates



# Staying up to date

## Atomic In-Place Updates

### 1. Stage

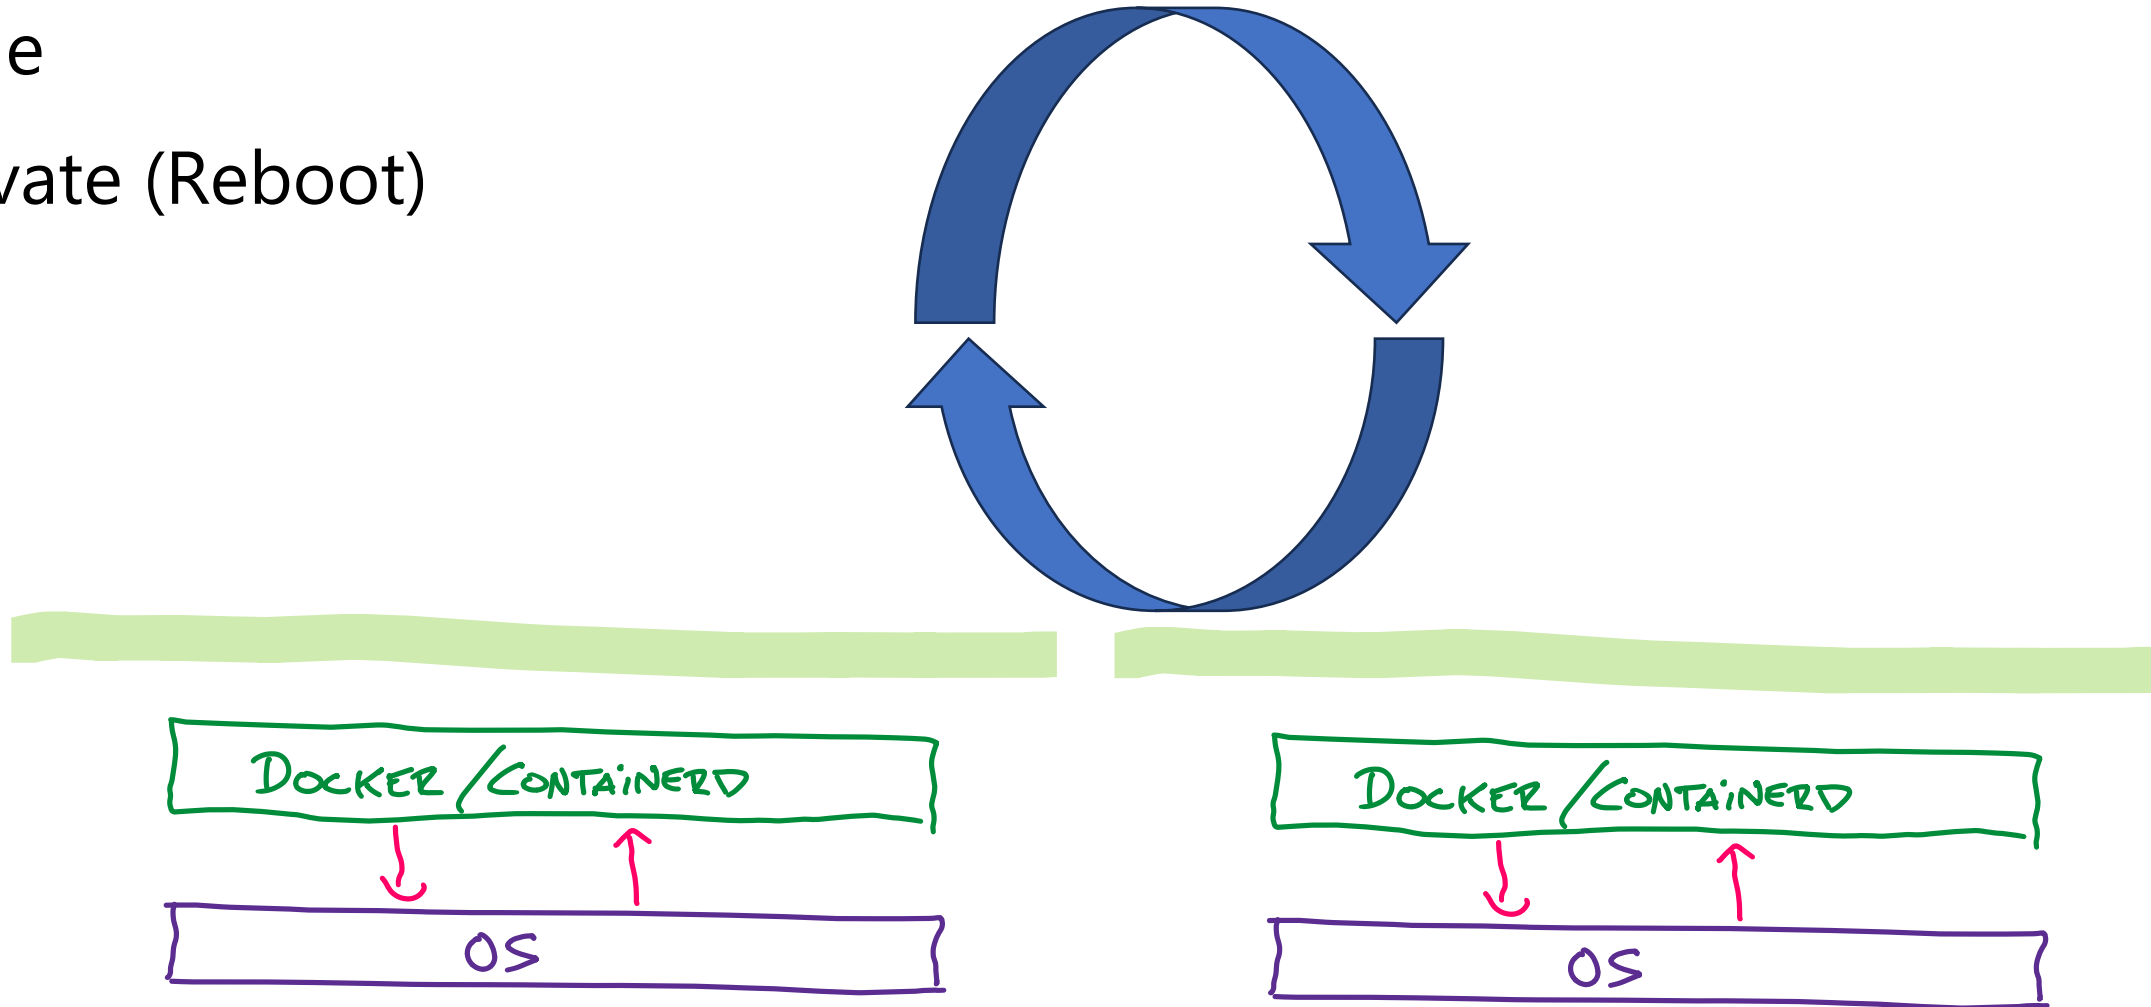




# Staying up to date

## Atomic In-Place Updates


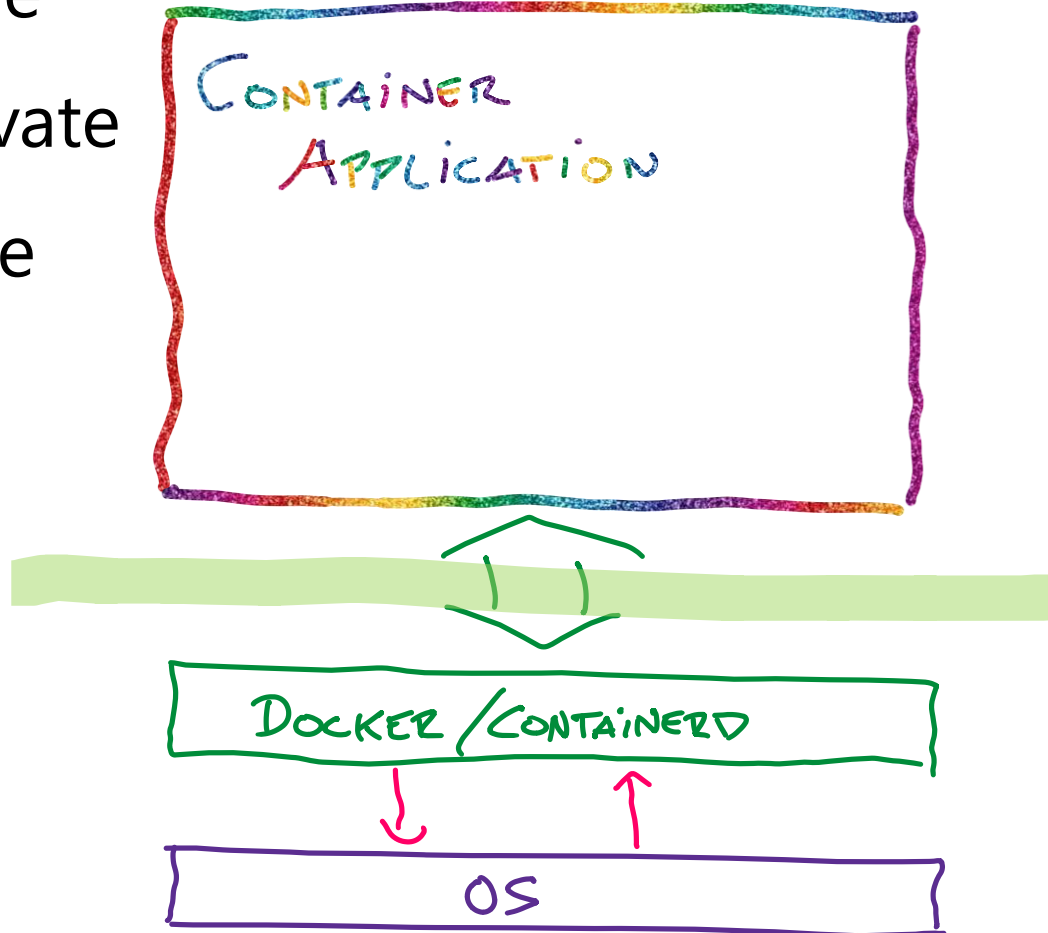
1. Stage
2. Activate (Reboot)



# Staying up to date

## Atomic In-Place Updates


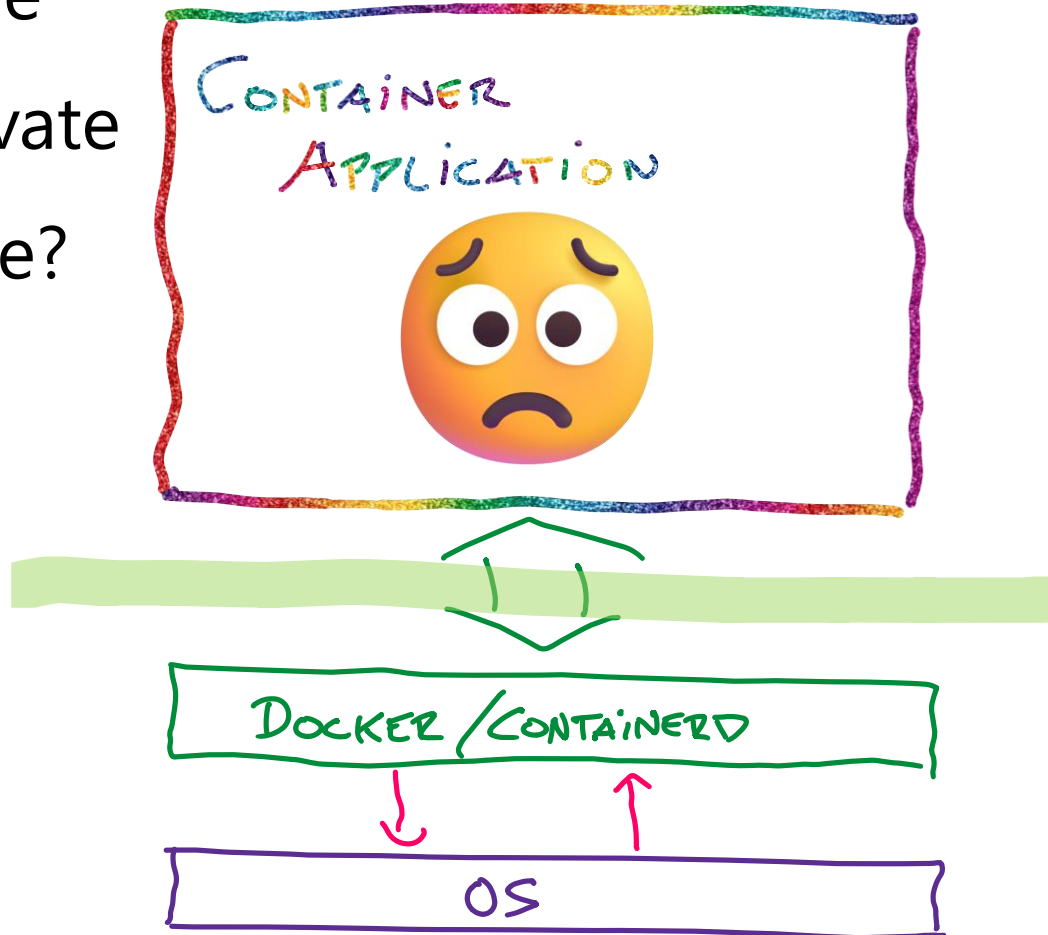
1. Stage
2. Activate
3. Done



# Staying up to date

## Atomic In-Place Updates

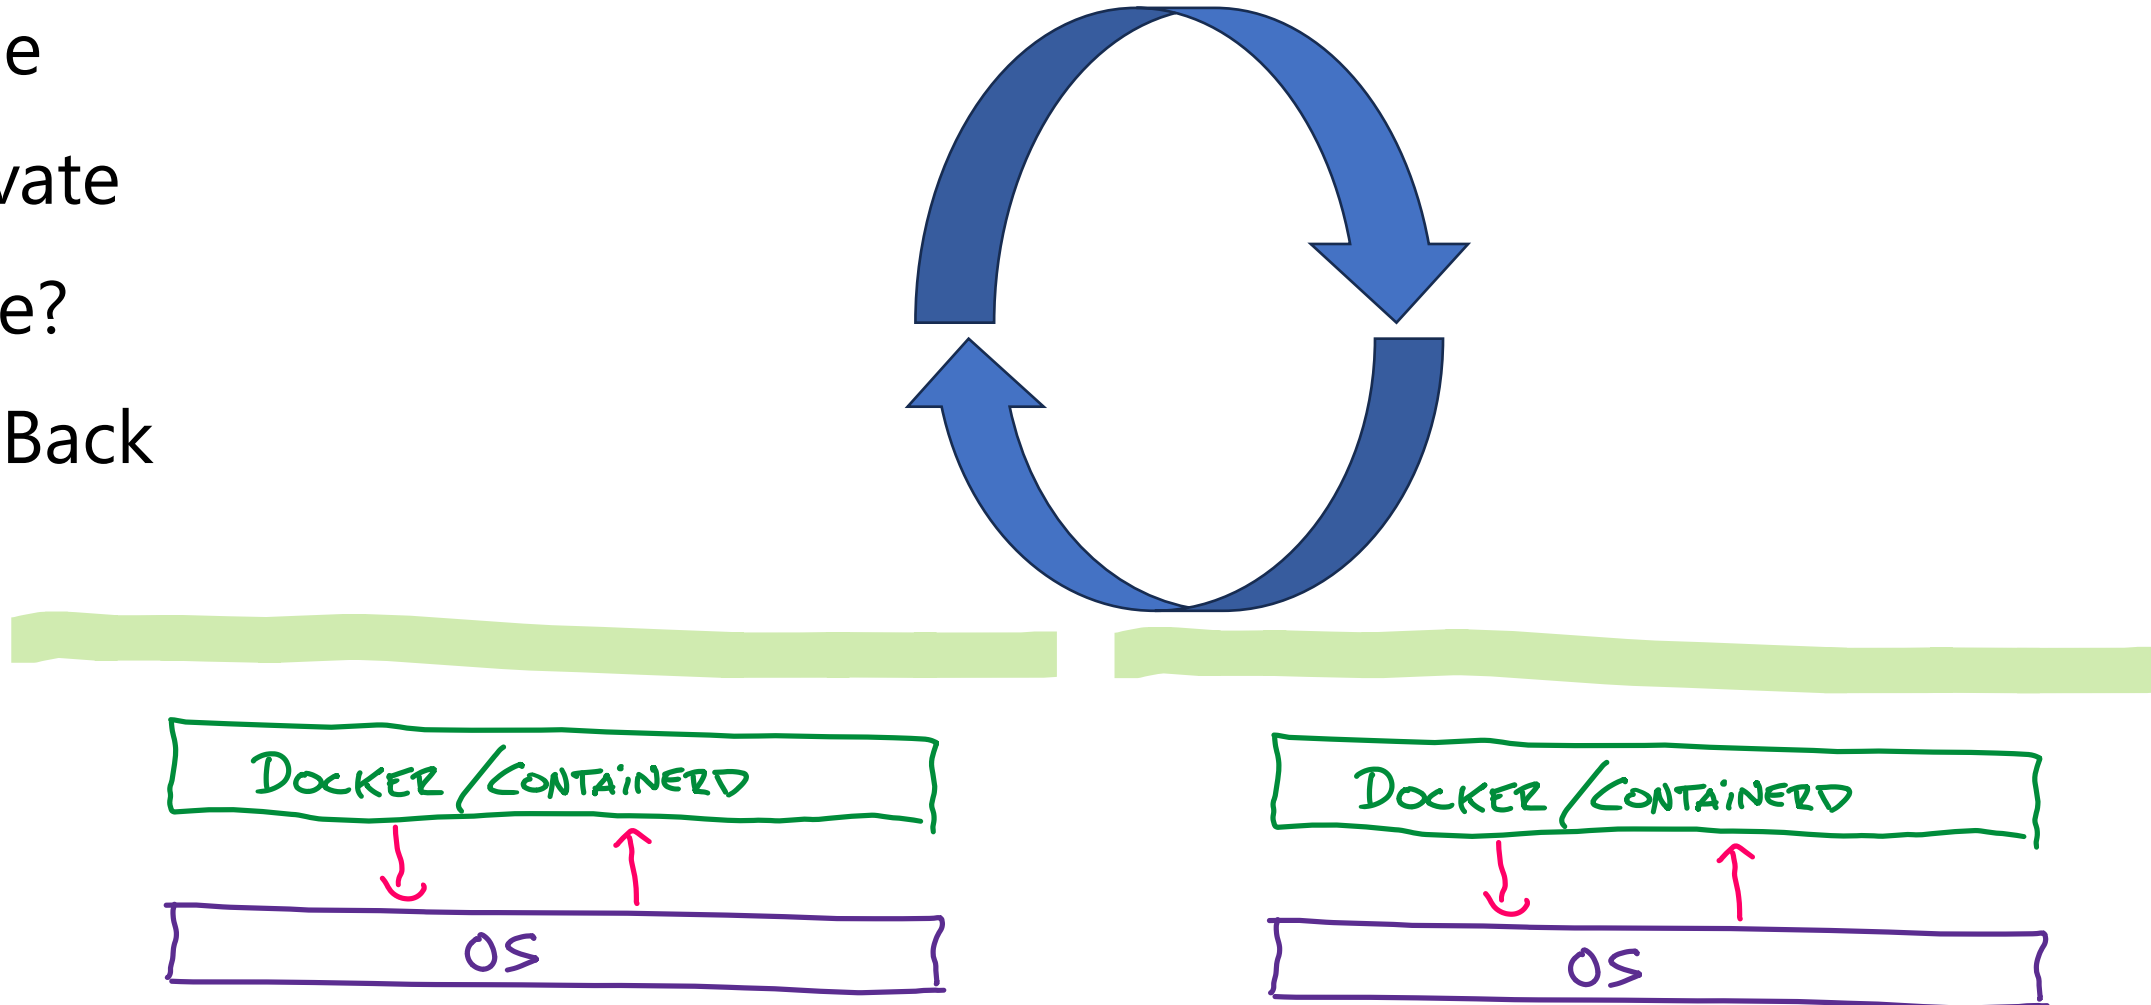
1. Stage
2. Activate
3. Done?



# Staying up to date

## Atomic Roll-Backs

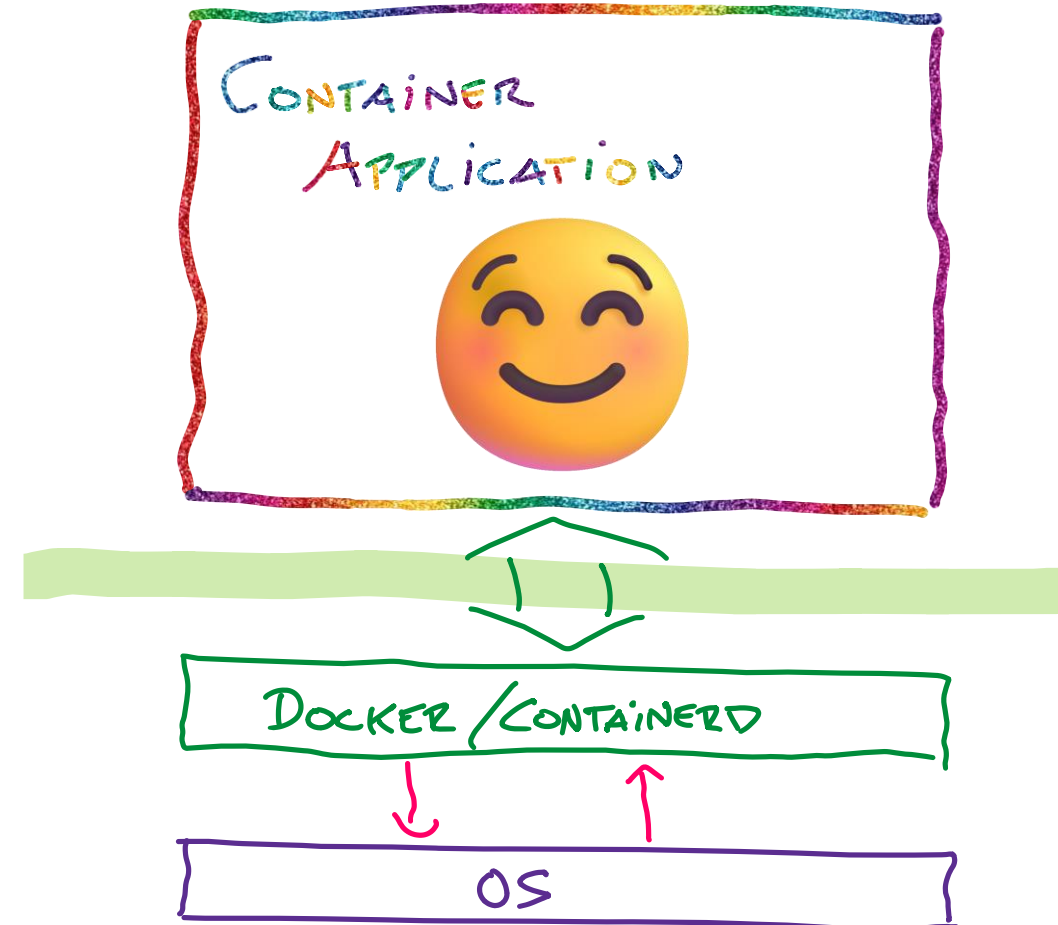

1. Stage
2. Activate
3. Done?
4. Roll Back



# Staying up to date

## Atomic Roll-Backs

1. Stage
2. Activate
3. Done?
4. Roll Back  
to known-good state



Update Demo



# Rolling out Updates

Defined Maintenance Windows

Update Operator for larger Clusters

Use Beta Canaries

Host-your-own stateful FOSS update server (Nebraska sub-project)

- ✓ Fully automated, easy to deploy
- ✓ Reliable update / lifecycle management
- Robust and safe, proven support of your target environment
- Solid configuration management / inventory management (SBOM)
- Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*

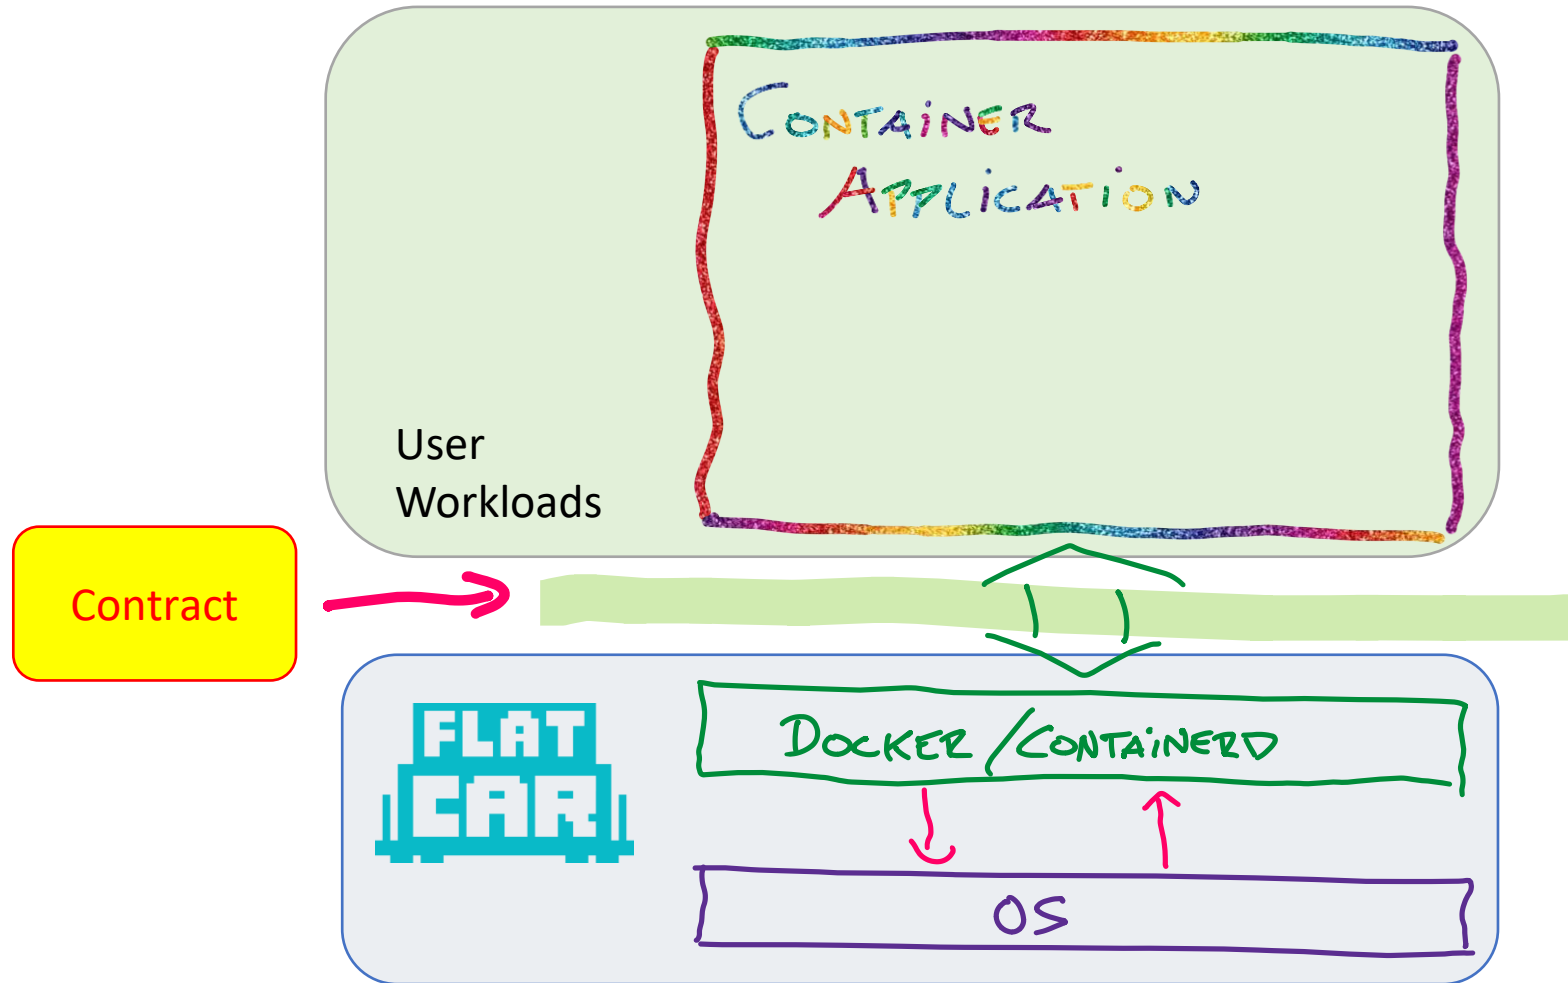


**Robust and Safe**



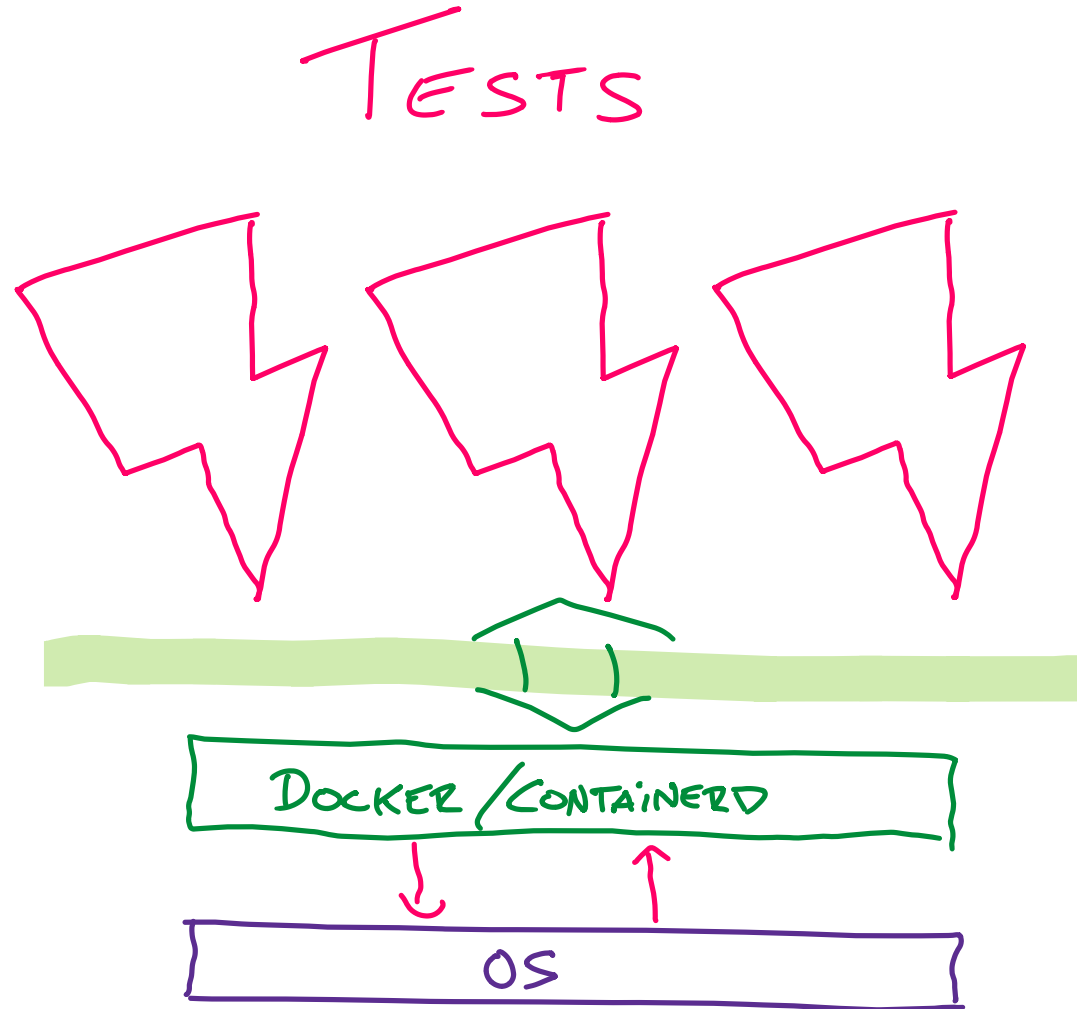
# Remember the Contract?

Container Runtime + Kernel == Contract



# Use the same contract for tests

Contract is well-testable (and rigorously tested)



# What's in a test run?

More than 100 tests, including complex Kubernetes scenarios like CNIs

Run on every release (all vendors), every CI build (QEMu) for almost every commit to the OS repo.

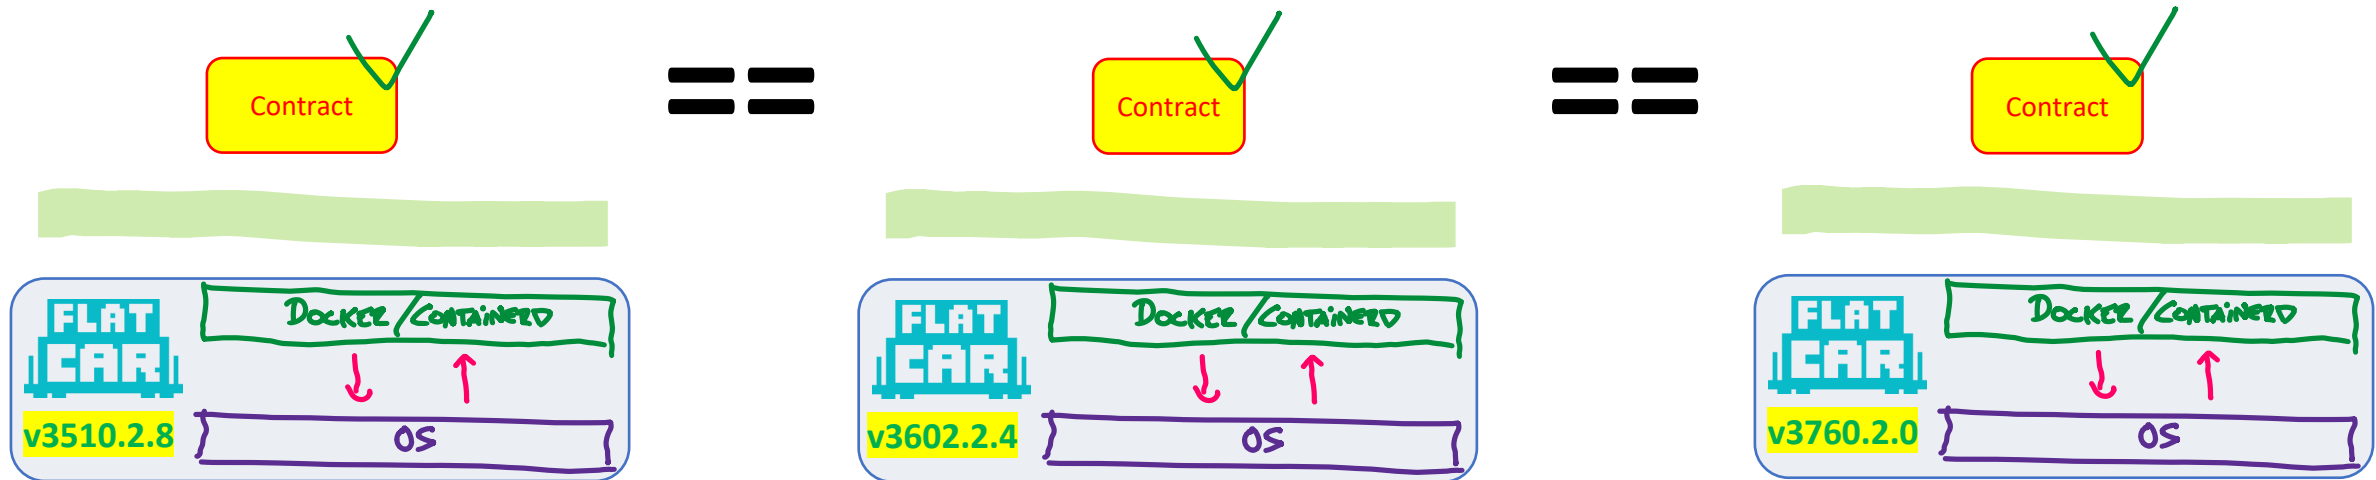
<a href="#">github-actions</a> <a href="#">bot</a> <a href="#">commented last week</a>				
<b>Test report for 4027.0.0-nightly-20240710-2100 / amd64 arm64</b>				
<b>Platforms tested:</b> qemu_ufi-amd64 qemu_update-amd64 qemu_ufi-arm64 qemu_update-arm64				
<a href="#">hpf.execonsoop</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.oem.indirect.new</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.misc.falco</a> Succeeded: qemu_ufi-amd64 (1)	<a href="#">coreos.ignition.groups</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">extra-test.[first_dual].cl.update.docker-btrfs-compat</a> Succeeded: qemu_update-amd64 (1) qemu_update-arm64 (1)
<a href="#">hpf.local.gadget</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.oem.regular</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.network.intramfs.second-boot</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.once</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">extra-test.[first_dual].cl.update.payload</a> Succeeded: qemu_update-amd64 (1) qemu_update-arm64 (1)
<a href="#">cl.basic</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.oem.regular.new</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.network.listeners</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.resource.local</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.28.7.calico.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.cgroupv1</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.oem.reuse</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.network.wireguard</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.resource.remote</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.28.7.calico.cgroupv1.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.cloudinit.basic</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.oem.wipe</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.osreset.ignition-rerun</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.security.s3.versioned</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.28.7.cilium.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.cloudinit.multipart-mime</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.partition_on_boot.disk</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.overlay.cleanup</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.security.tls</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.28.7.cilium.cgroupv1.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.cloudinit.script</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.symblink</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.swap.activation</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.sethostname</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.28.7.flannel.cgroupv1.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.disk.raido.data</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.translation</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.sysboot</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.ignition.system.enable-service</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.28.7.flannel.flannel.cgroupv1.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.disk.raido.root</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.v1.btrfsroot</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.sysfsboot</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.locksmith.reboot</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.29.2.calico.base</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)
<a href="#">cl.disk.raido1.data</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.ignition.v1.ext4root</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">cl.sysfsfallbackdownload</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">coreos.locksmith.tls</a> Succeeded: qemu_ufi-amd64 (1) qemu_ufi-arm64 (1)	<a href="#">kubeadm.v1.29.2.cilium.base</a>



# Contract is upheld across releases

Contract is well-tested

Always upheld across releases



- ✓ Fully automated, easy to deploy
- ✓ Reliable update / lifecycle management
- ? Robust and safe, proven support of your target environment
- Solid configuration management / inventory management (SBOM)
- Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*

- ✓ Fully automated, easy to deploy
- ✓ Reliable update / lifecycle management
- ? Robust and safe, proven support of your target environment
- Solid configuration management / inventory management (SBOM)
- Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

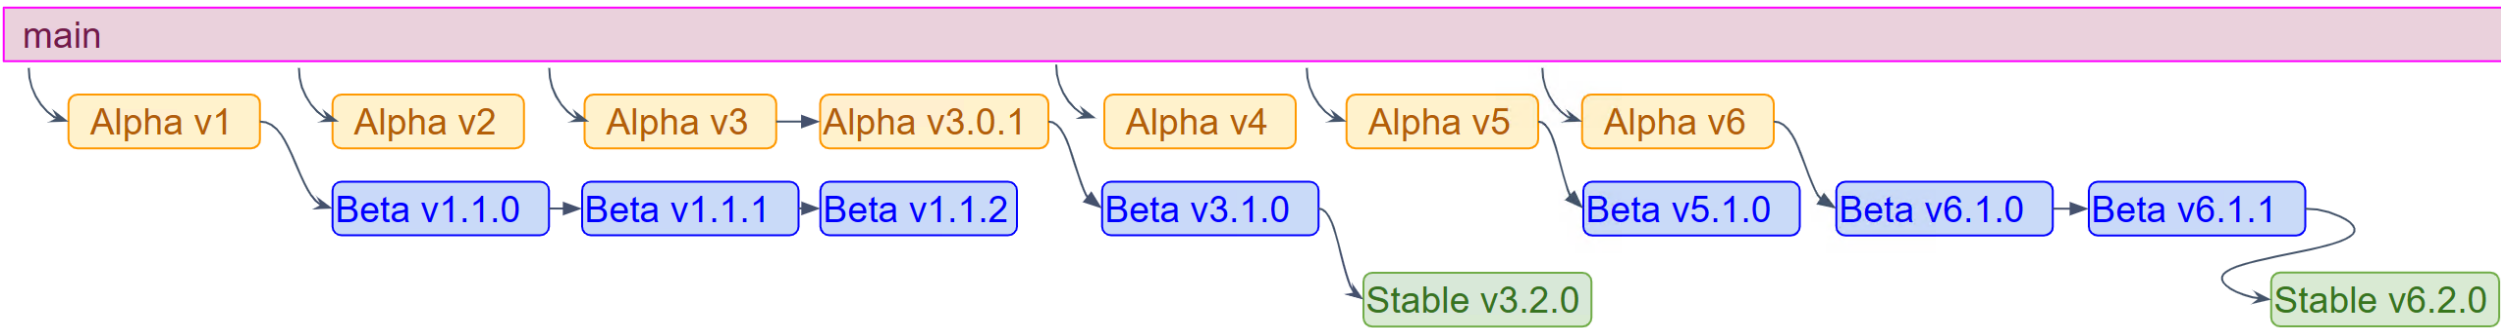
*Can we translate this philosophy to operating systems?*

# Major updates: Channel Progression

## Release Channels

<b>Stable</b> 3760.2.0 <a href="#">amd64</a> <a href="#">arm64</a> Release Date: Jan 18, 2024  The Stable channel is intended for use in production clusters. Versions of Flatcar Container Linux have been tested as they move through Alpha and Beta channels before being promoted to stable.  <b>containerd</b> - 1.7.7 <b>docker</b> - 20.10.24 <b>ignition</b> - 2.15.0 <b>kernel</b> - 6.1.73 <b>systemd</b> - 252	<b>Beta</b> 3815.1.0 <a href="#">amd64</a> <a href="#">arm64</a> Release Date: Jan 18, 2024  The Beta channel is where Flatcar Container Linux stability is solidified. We encourage including some beta machines in production clusters in order to catch any issues that may arise with your setup.  <b>containerd</b> - 1.7.10 <b>docker</b> - 24.0.6 <b>ignition</b> - 2.15.0 <b>kernel</b> - 6.1.73 <b>systemd</b> - 252	<b>Alpha</b> 3850.0.0 <a href="#">amd64</a> <a href="#">arm64</a> Release Date: Jan 18, 2024  The Alpha channel follows a more frequent release cadence and is where new updates are introduced. Users can try the new versions of the Linux kernel, systemd and other core packages.  <b>containerd</b> - 1.7.11 <b>docker</b> - 24.0.6 <b>ignition</b> - 2.15.0 <b>kernel</b> - 6.6.12 <b>systemd</b> - 252	<b>LTS</b> 3510.3.1 <a href="#">amd64</a> <a href="#">arm64</a> Release Date: Oct 25, 2023  LTS release streams will be maintained for an extended lifetime of 18 months. The yearly LTS streams have an overlap of 6 months.  <b>containerd</b> - 1.6.16 <b>docker</b> - 20.10.23 <b>ignition</b> - 2.14.0 <b>kernel</b> - 5.15.136 <b>systemd</b> - 252
---	---	---	---

# Ensuring *safe* updates through channel progression

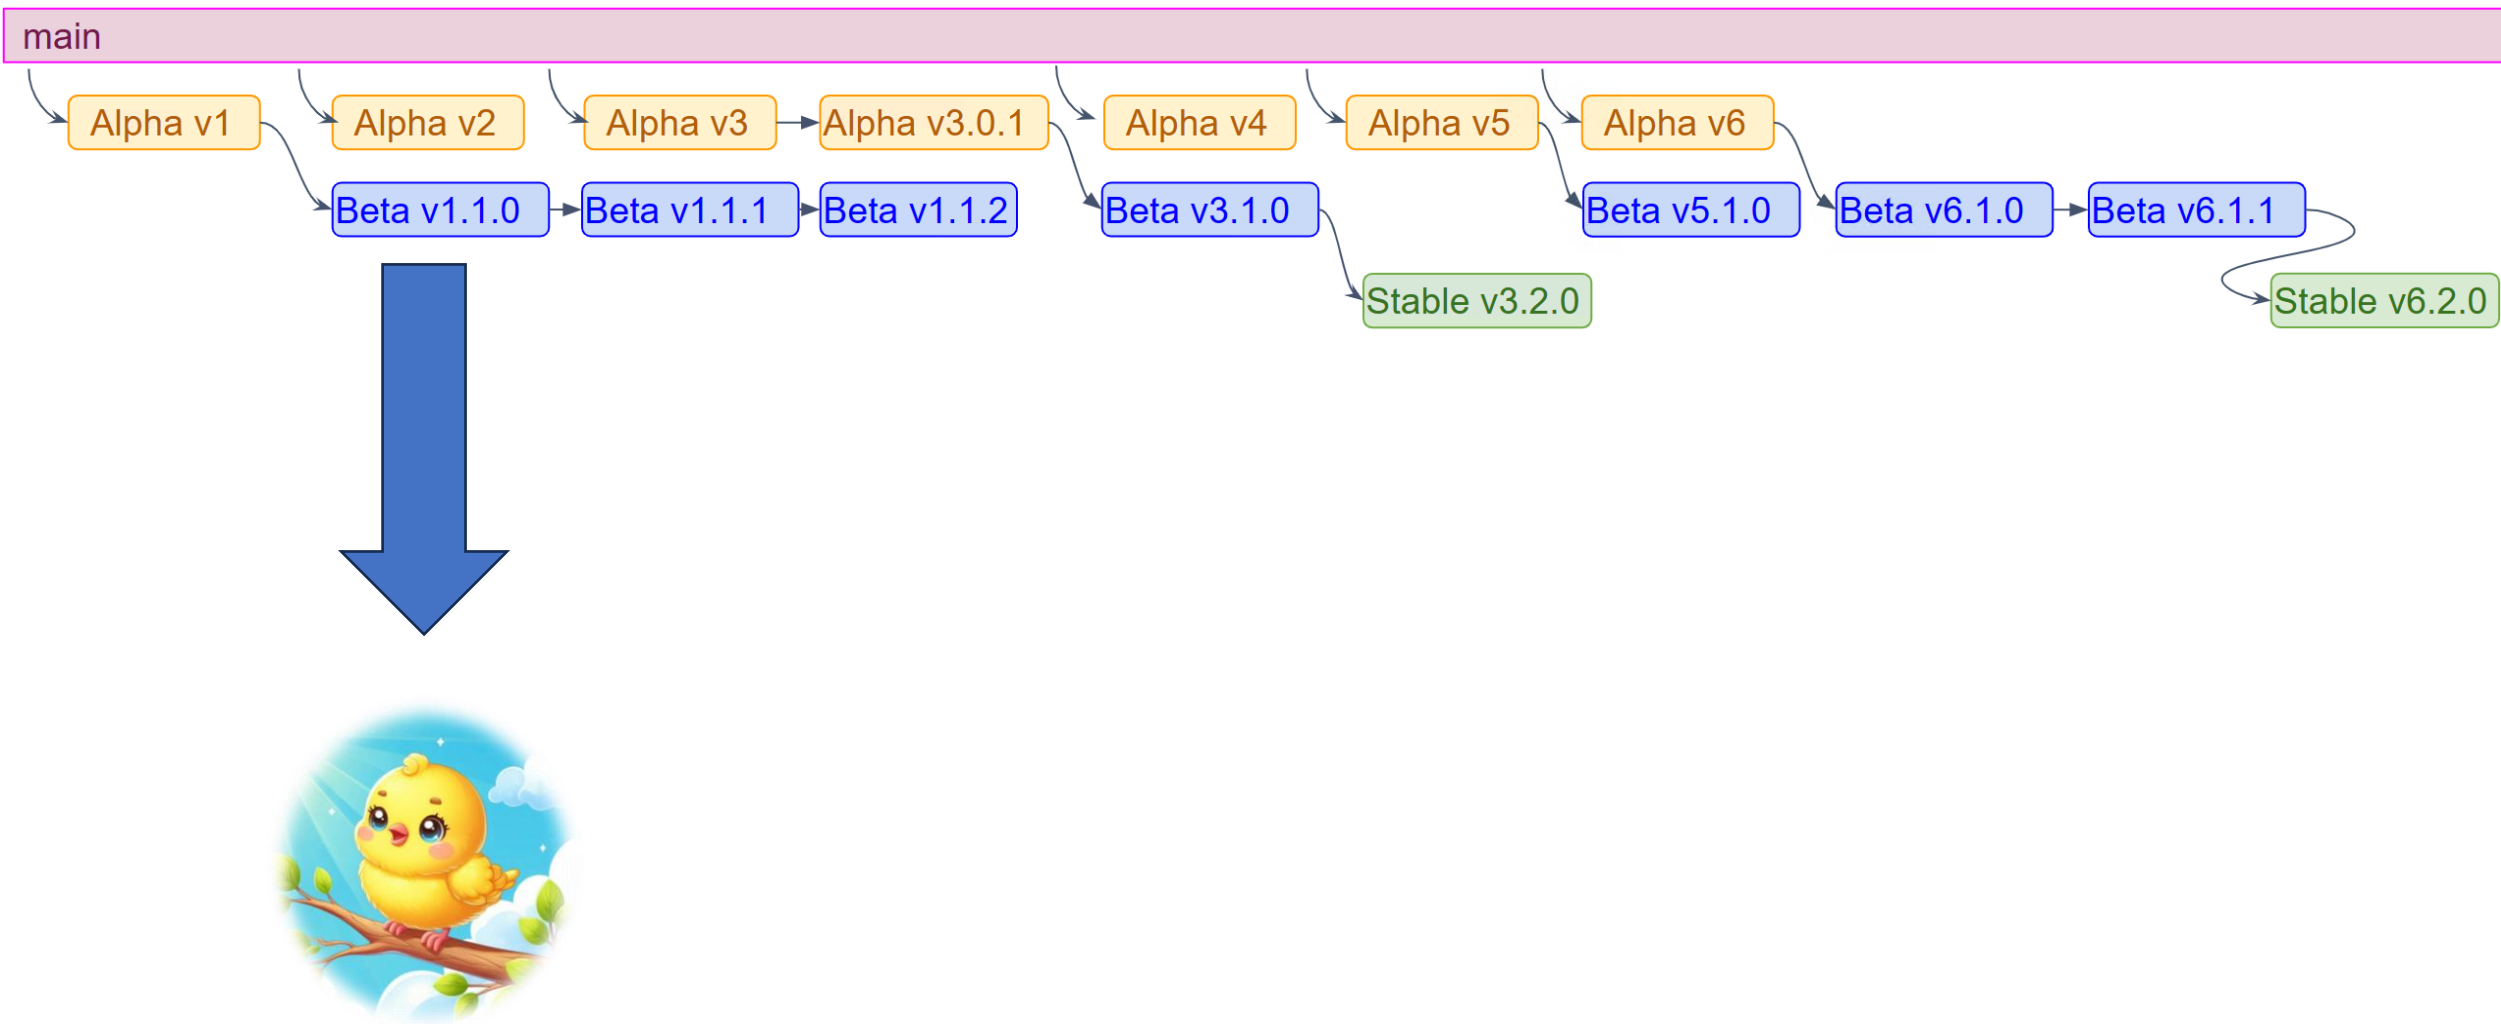


## Major OS release stabilisation milestones:

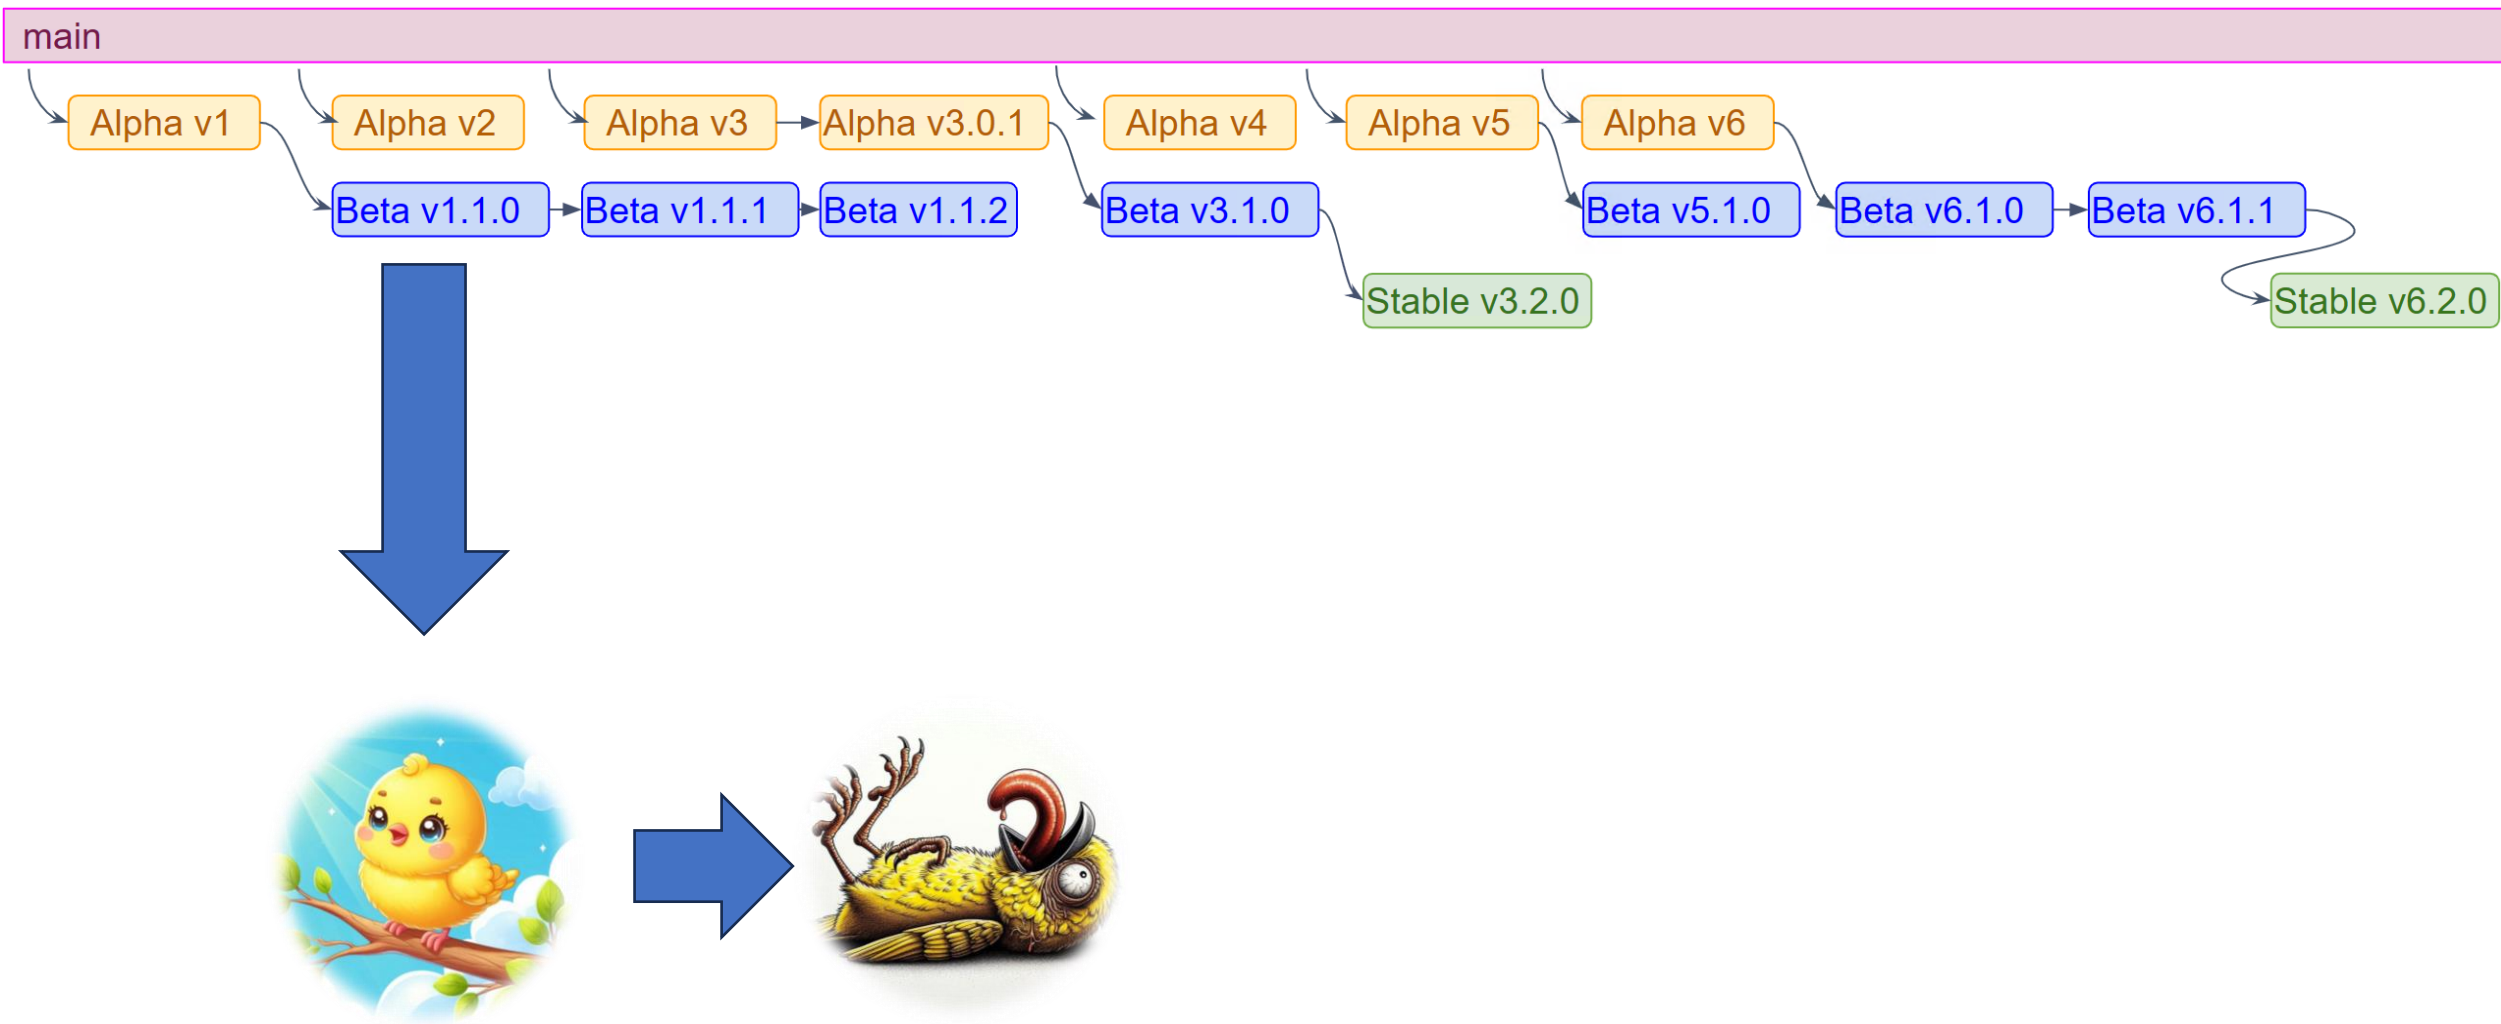
- "Alpha" Fully tested but may contain incomplete features. For developers.
- "Beta" Fully tested for production use. Recommended for canaries.
- "Stable" For widespread production use.  
Additional stabilisation through user feedback from Beta canaries.

Deployments defaults to "stable" but can be customised to any channel.

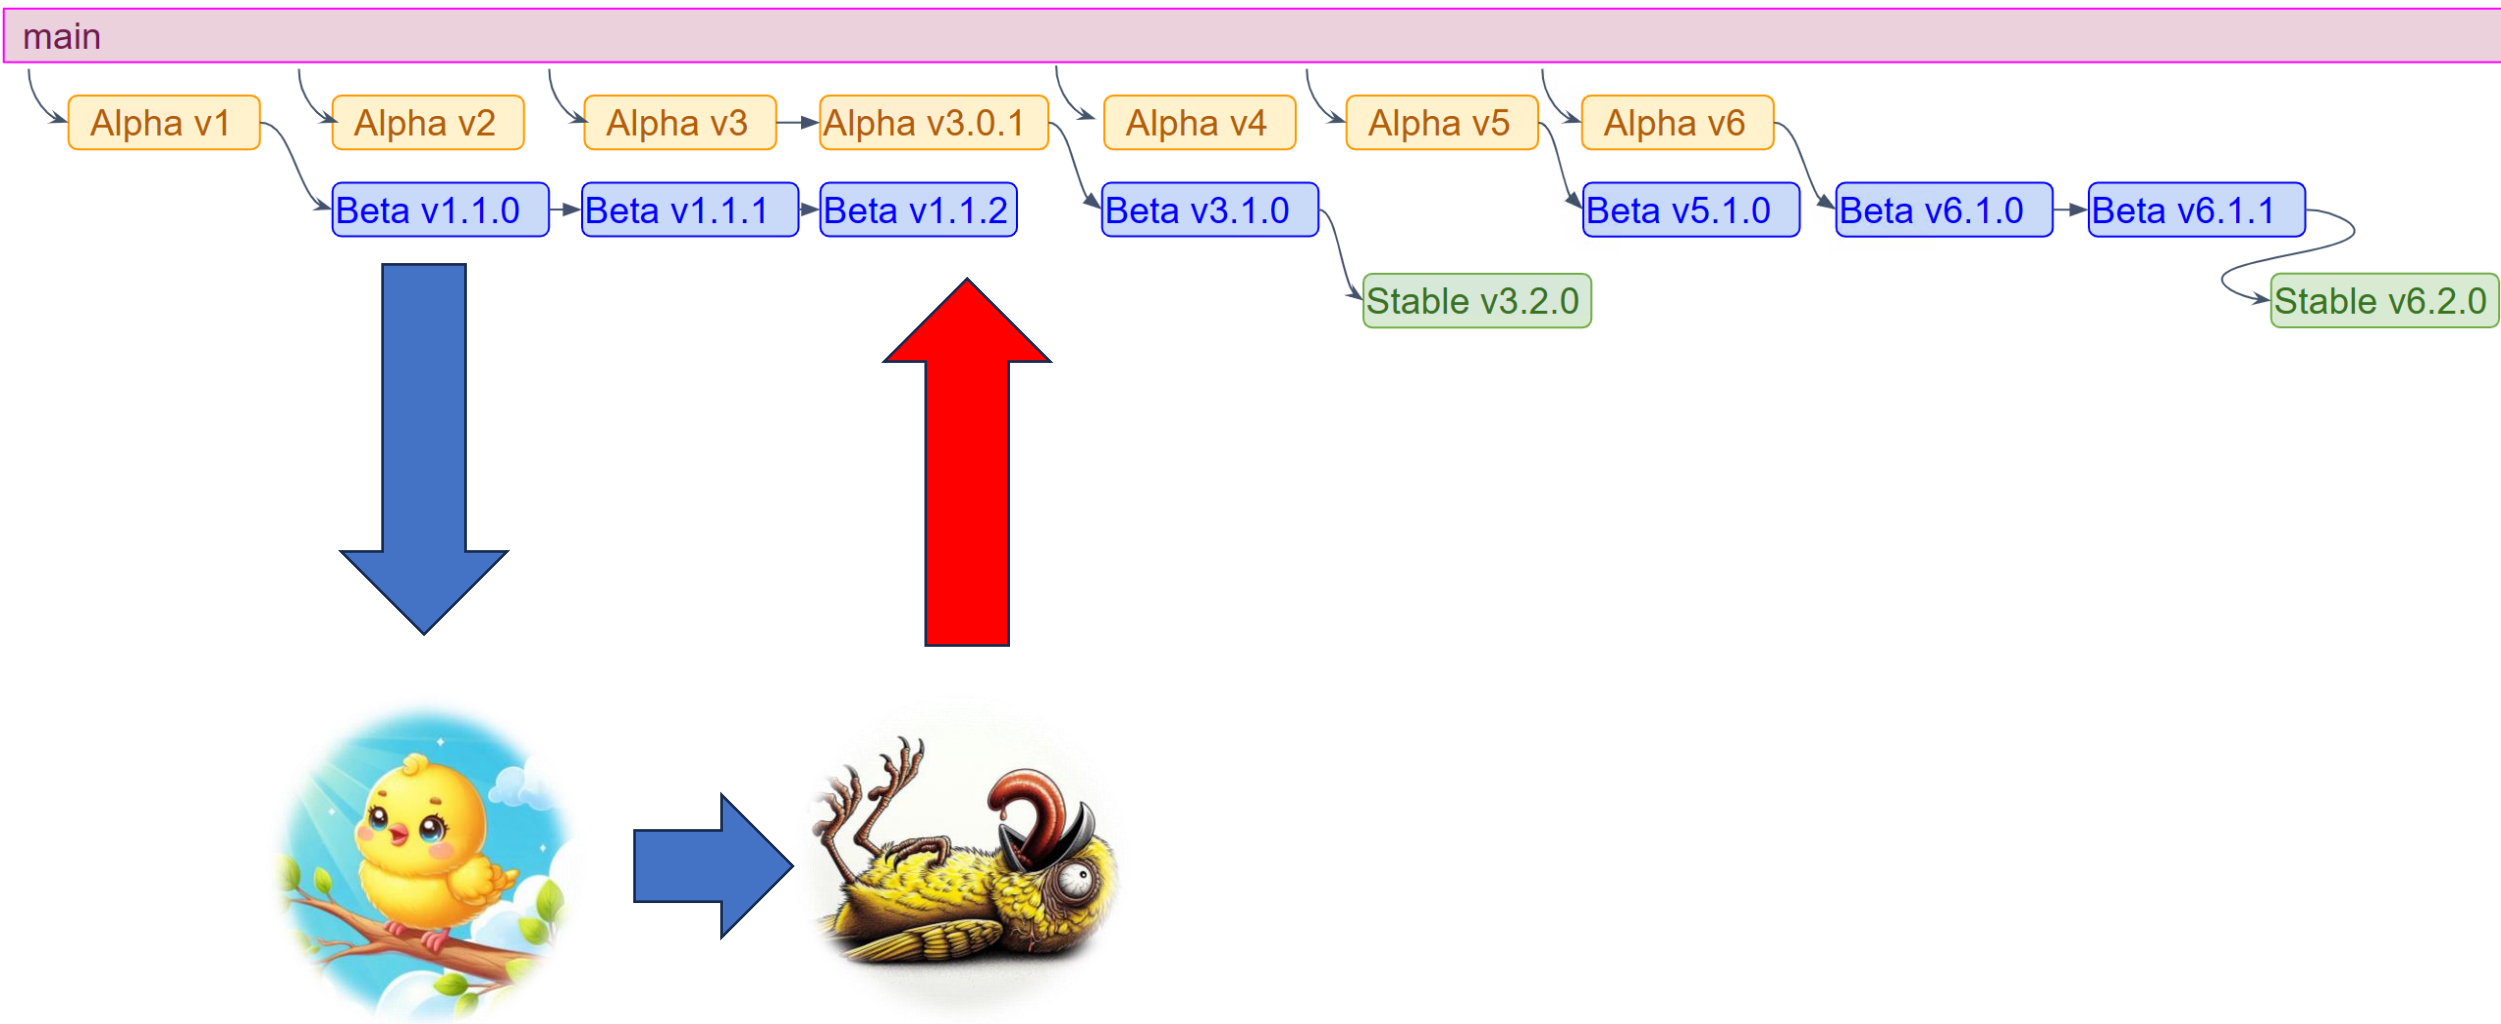
# Ensuring *safe* updates through Beta canaries



# Ensuring *safe* updates through Beta canaries



# Ensuring *safe* updates through Beta canaries





- ✓ Fully automated, easy to deploy
  - ✓ Reliable update / lifecycle management
  - ✓ Robust and safe, proven support of your target environment
- Solid configuration management / inventory management (SBOM)
- Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*



# Configuration and Inventory management



# Configuration Management



Configure

->

Deploy

->

Operate

Sane defaults  
- no boiler plate.

Integration in  
ops environment.

Customisation.

Butane-config.yaml

Custom data,  
User data,  
Http, or  
[i]PXE

ignition.json

**Applied ONCE.**  
**At provisioning time.**

Automated

- Self-configuration
- Unattended updates

# Configuration Management



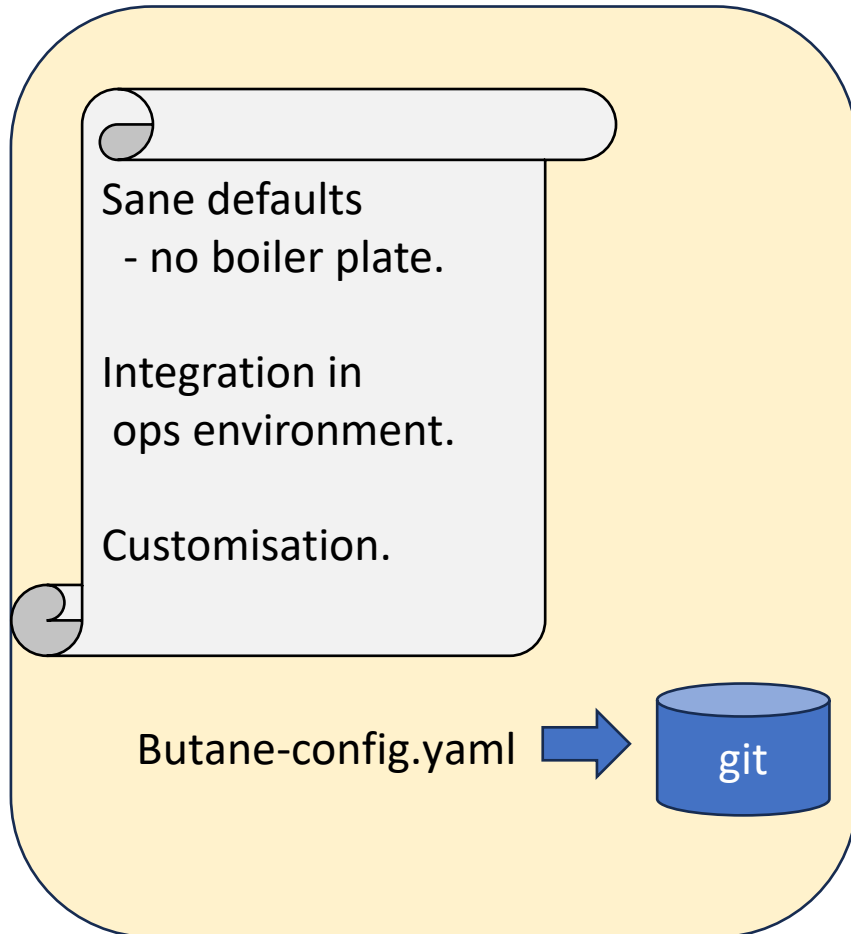
Configure

->

Deploy

->

Operate



Custom data,  
User data,  
Http, or  
[i]PXE

ignition.json

Automated

- Self-configuration
- Unattended updates

# Software Inventory Management

# Software Inventory Management

## Image-Based OS:

- Provisioned from a full disk image.
- Updated by a full OS partition image.
- Immutable: No package manager, no post-deploment installations.

# Software Inventory Management

## Image-Based OS:

- Provisioned from a full disk image.
- Updated by a full OS partition image.
- Immutable: No package manager, no post-deploment installations.

All packages and all files are documented for each release.

- No version drift.
- OS software inventory is pinned to OS release.

- ✓ Fully automated, easy to deploy
  - ✓ Reliable update / lifecycle management
  - ✓ Robust and safe, proven support of your target environment
  - ✓ Solid configuration management / inventory management (SBOM)
- Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*





# Immutable OS Extensions



# OS-level extensibility via Systemd Sysext

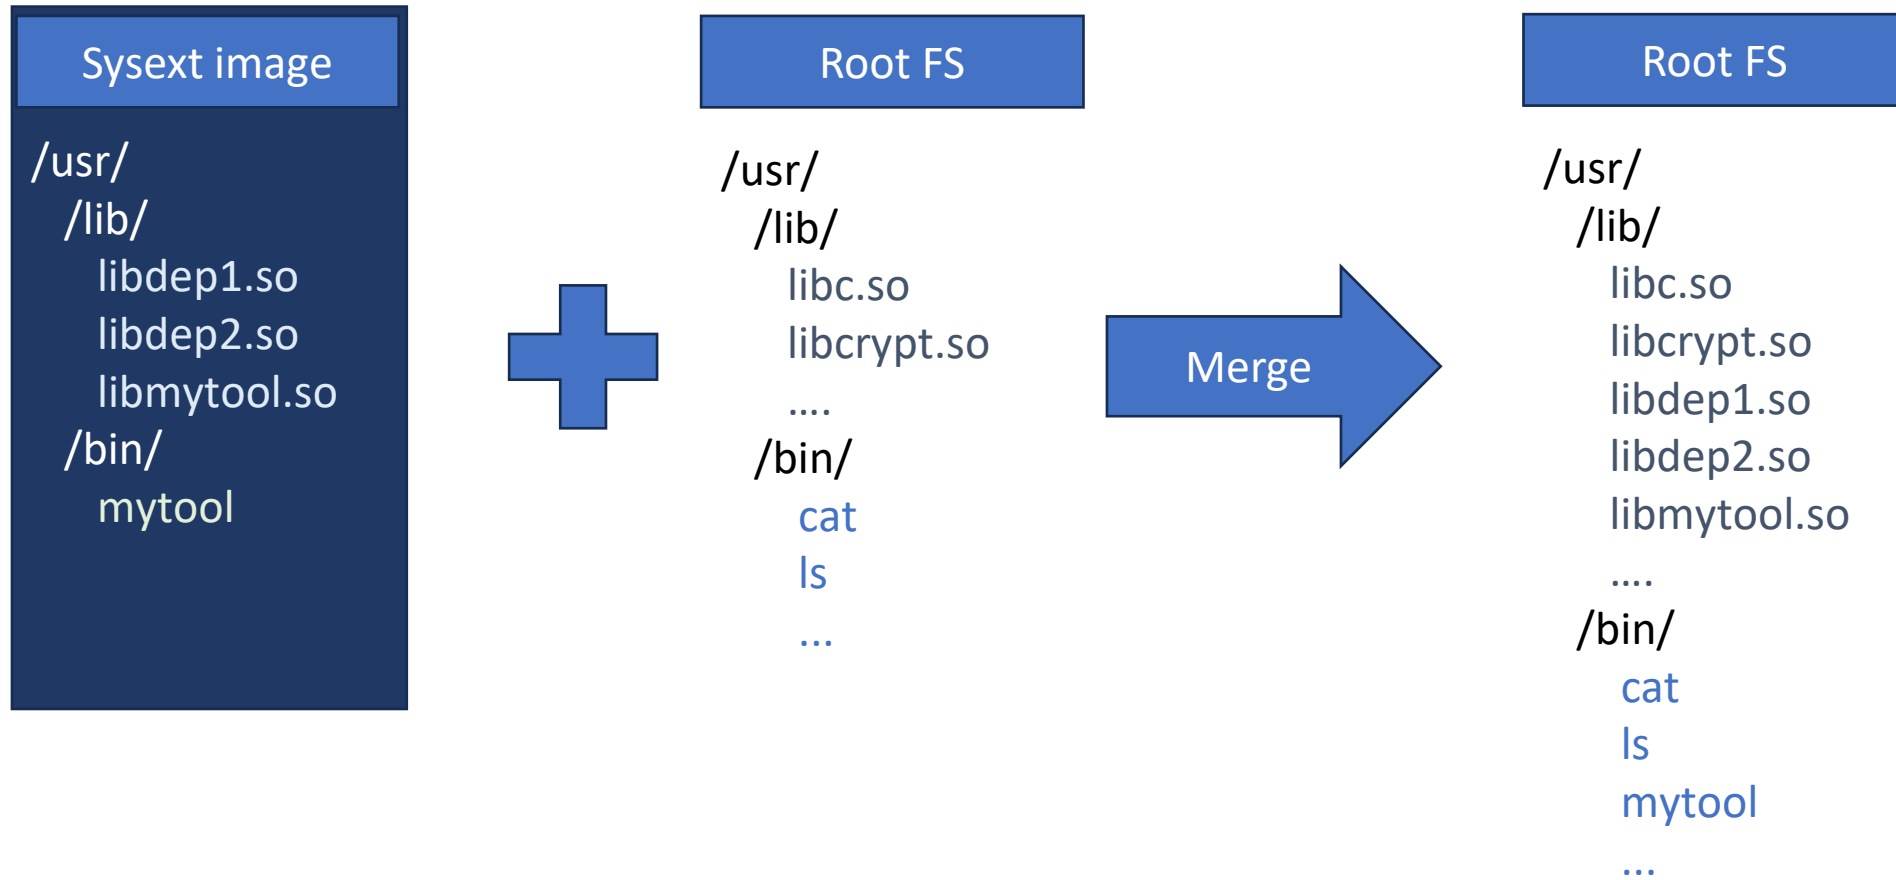
Nice set of tools, but I need podman/Kubernetes/WASM/...

Extensible via systemd-sysext:

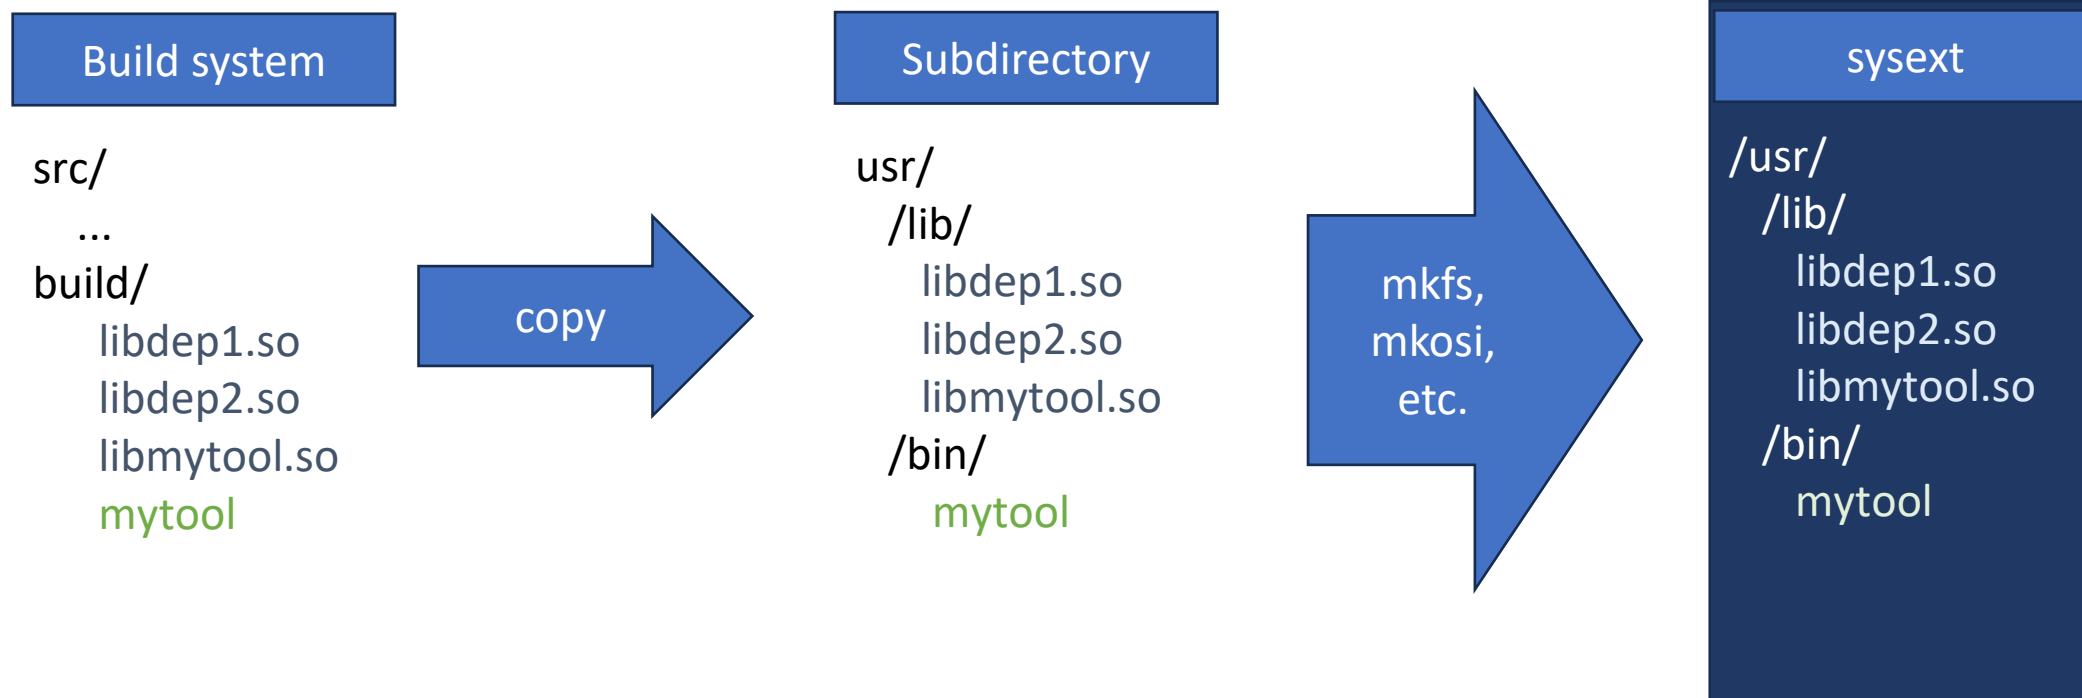
- Image-based, immutable extensions

Solid update story separate from OS updates

# Using Sysexts (Merging)



# Building Sysexts



# Sysexts Day #2

Clean separation between OS + customisations

Solid provisioning story: supply update config via Ignition

Solid update story: update customisations independent of the OS

Straightforward packaging: build output to FS image

Easy publishing of updates: Serve via HTTPS, white-list in index file

# Image composability in Flatcar

System dependent sysexts

- OEM / Vendor support (guest tools etc.)

- Optional OS features (python, zfs, podman, ...)

System independent sysexsts

- Kubernetes (e.g. for ClusterAPI support), Docker compose

- Tailscale, Wasmcloud + Wasmtime

- Cri-o, k3s, Rke2 (next-gen Rancher)

# Sysexr Demo

- ✓ Fully automated, easy to deploy
- ✓ Reliable update / lifecycle management
- ✓ Robust and safe, proven support of your target environment
- ✓ Solid configuration management / inventory management (SBOM)
- ✓ Integrates well into operational processes

*On Kubernetes, we operate workloads as interchangeable, replaceable commodities ("cattle, not pets").*

*Can we translate this philosophy to operating systems?*



# Flatcar Community

## Community-driven FOSS project

No single vendor, full community stewardship

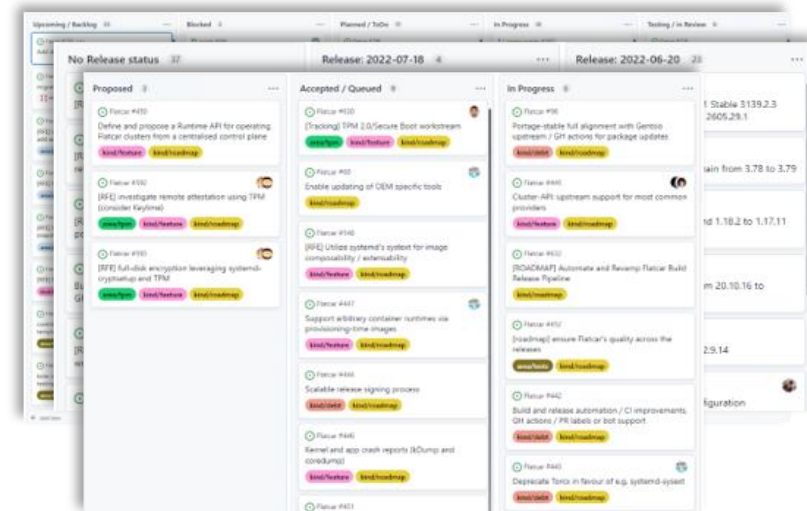
Submitted to the CNCF as incubation project (ongoing)

[Matrix](#), [Slack](#) - Our day-to-day comms

[Office hours](#) - Every 2nd Wednesday, 2:30pm UTC

[Dev Sync](#) - Every 4th Wednesday, 2:30pm UTC

[Roadmap](#), [Implementation](#), [Releases](#)





Thank you

The Community's  
Container Linux

