



Allet juht, Atze!

Kubernetes & Chill with automated infrastructure



KCD Berlin | 2022-06-30

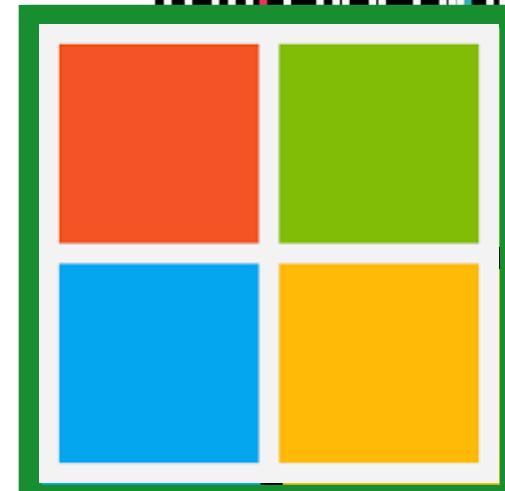
Hi, I'm Danielle



Danielle Tal
Program Manager,
Microsoft

Github: [miao0miao](https://github.com/miao0miao)

Email: danittal@microsoft.com



Hi, I'm Thilo



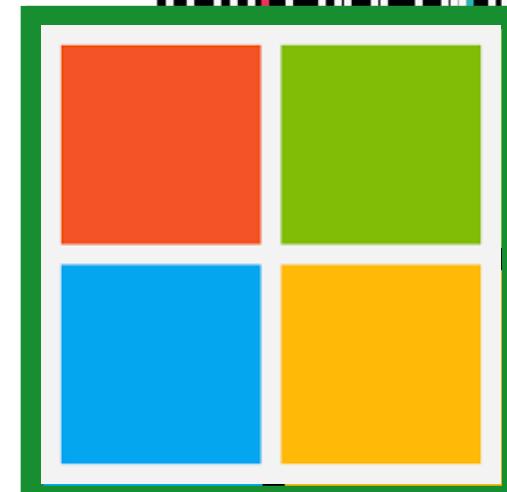
Thilo Fromm
Engineering manager,
Microsoft

Github: [t-lo](#)

Mastodon: [@thilo@fromm.social](#)

Twitter: [ThiloFM](#)

Email: thilofromm@microsoft.com



Fei|er|abend



1) Free time following a day of work.

"To enjoy one's Feierabend."

2) End of work / closing time.

"I'm calling Feierabend, heading home!"

3) Vigorous shut-down of an ongoing activity.

"Stop it, no more, that's it - Feierabend!"





Container-optimised Linux History

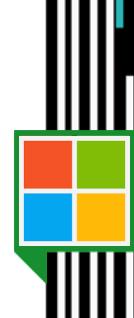


Container-optimised Linux History



CoreOS Container Linux (Oct. 2013 – May 2020)

(derived from Chromium OS, which was derived from Gentoo)





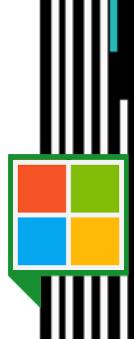
Container-optimised Linux History

CoreOS Container Linux (Oct. 2013 – May 2020)

(derived from Chromium OS, which was derived from Gentoo)

→ Fedora CoreOS / Red Hat CoreOS (June 2018)

(based on Fedora, limited compatibility to CoreOS)





Container-optimised Linux History

CoreOS Container Linux (Oct. 2013 – May 2020)

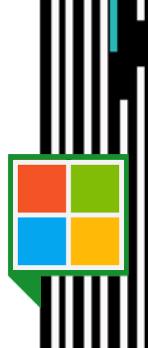
(derived from Chromium OS, which was derived from Gentoo)

→ Fedora CoreOS / Red Hat CoreOS (June 2018)

(based on Fedora, limited compatibility to CoreOS)

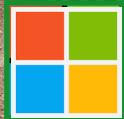
→ Flatcar Container Linux (Nov. 2019)

(Friendly fork, fully compatible to CoreOS)





Container-optimised Linux?





Container-optimised Linux?

Your OS is infrastructure

Should be boring (not exciting), just work

Let you focus on your business logic

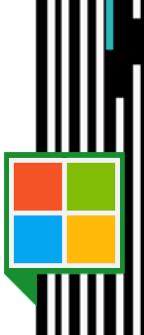
If a light switch needs constant attention,

are you ever going to enjoy your fancy home cinema?

Container isolation applied to the OS?

We isolate apps from each other, and from the OS.

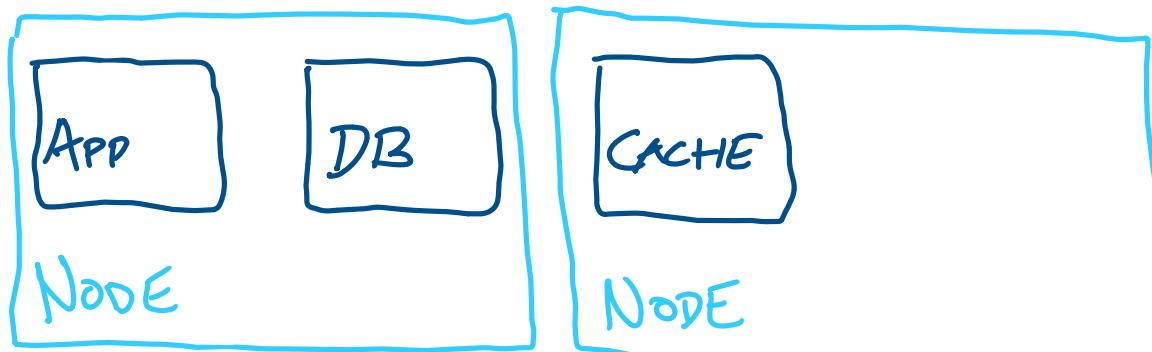
Why not use that to make the OS yet another app?





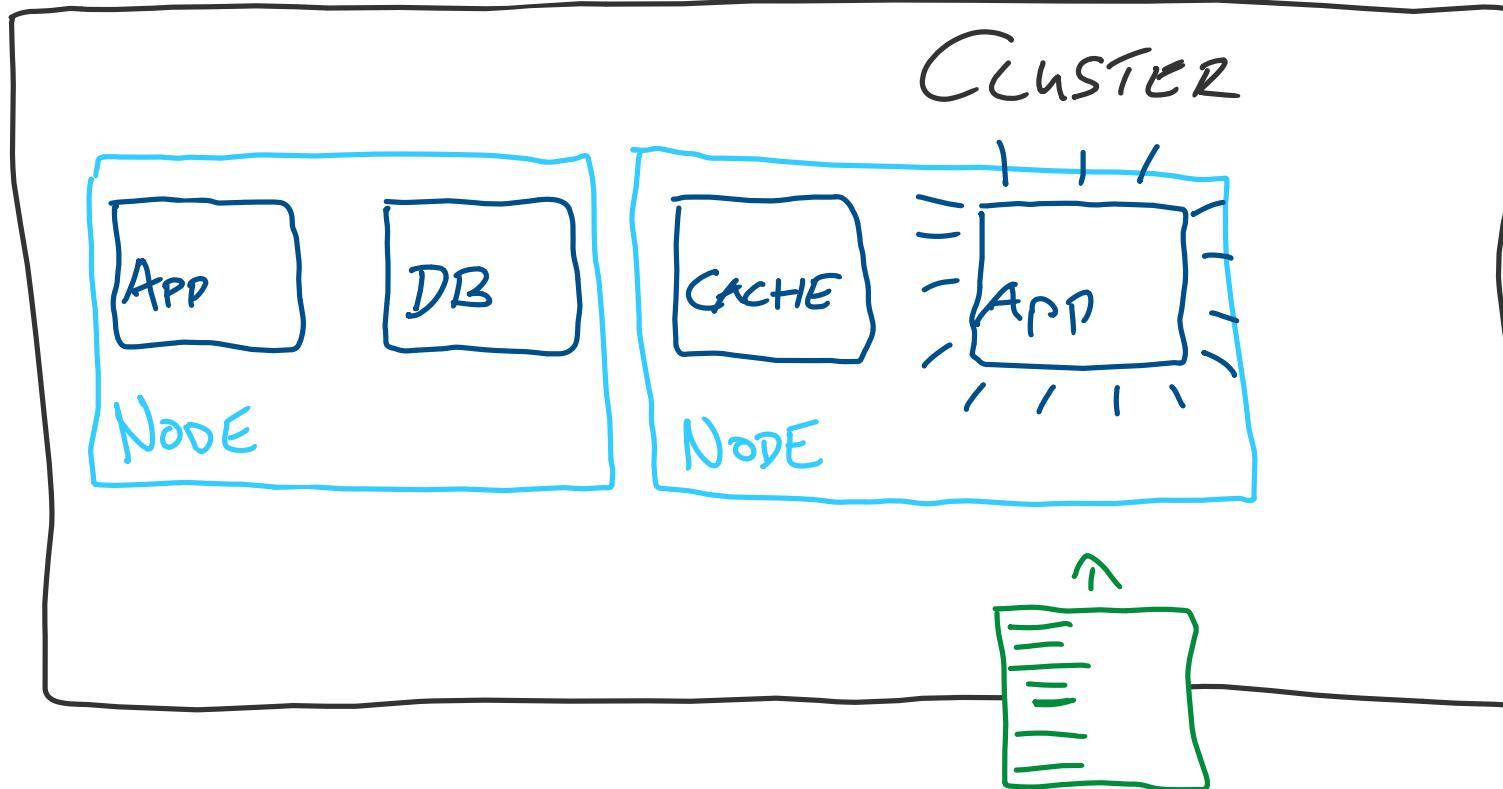
Container-optimised Linux?

CLUSTER



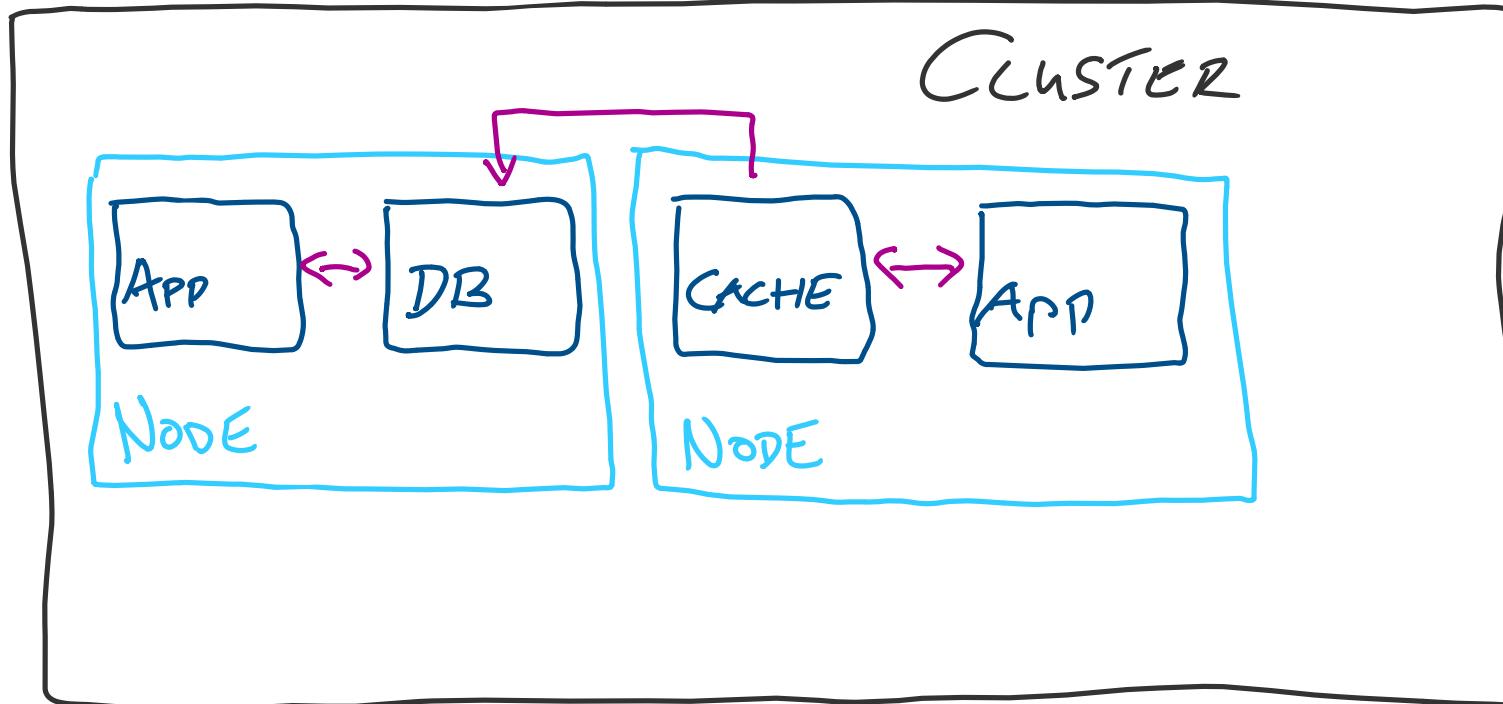


Container-optimised Linux?



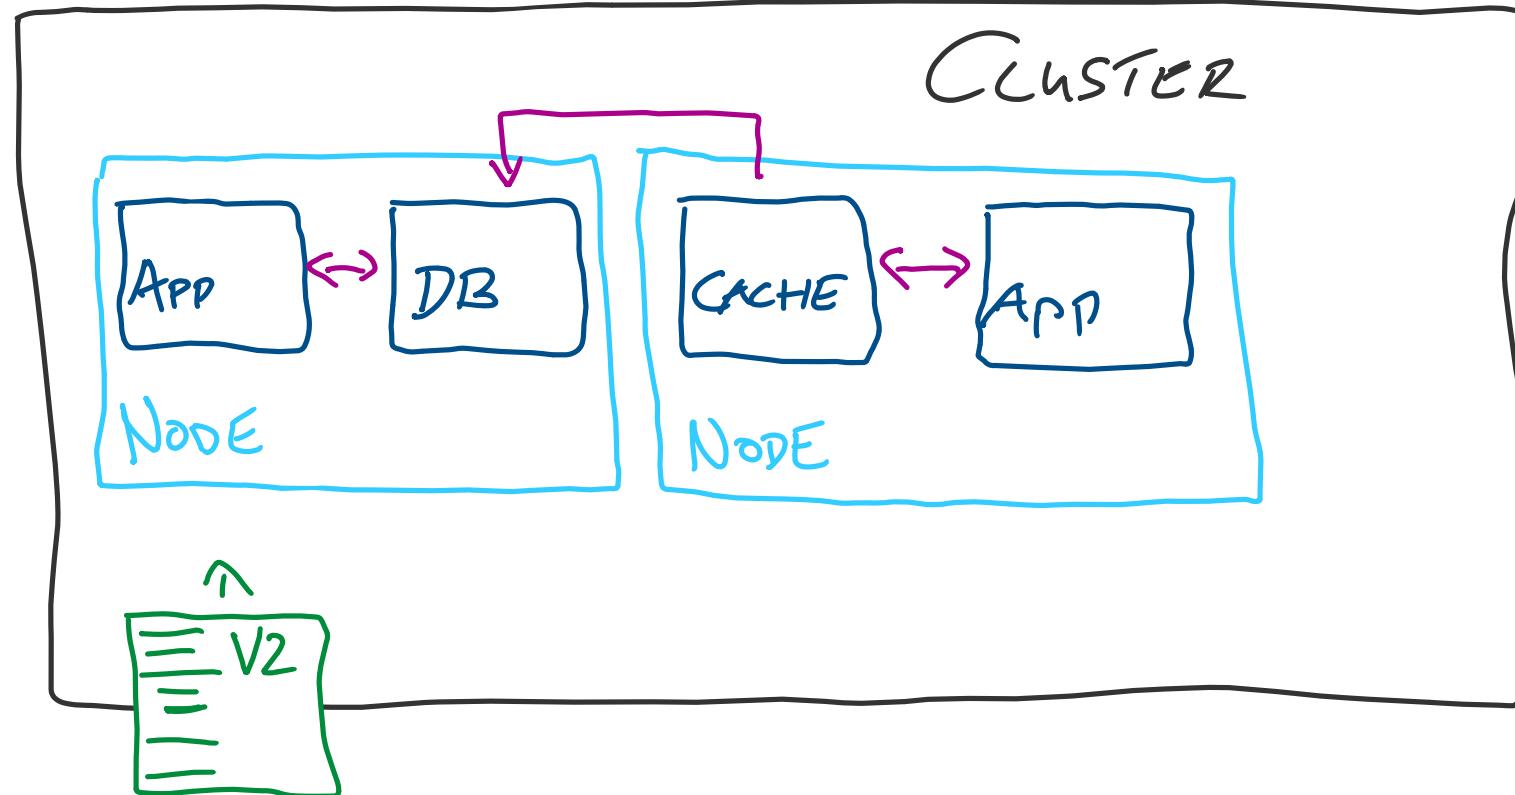


Container-optimised Linux?



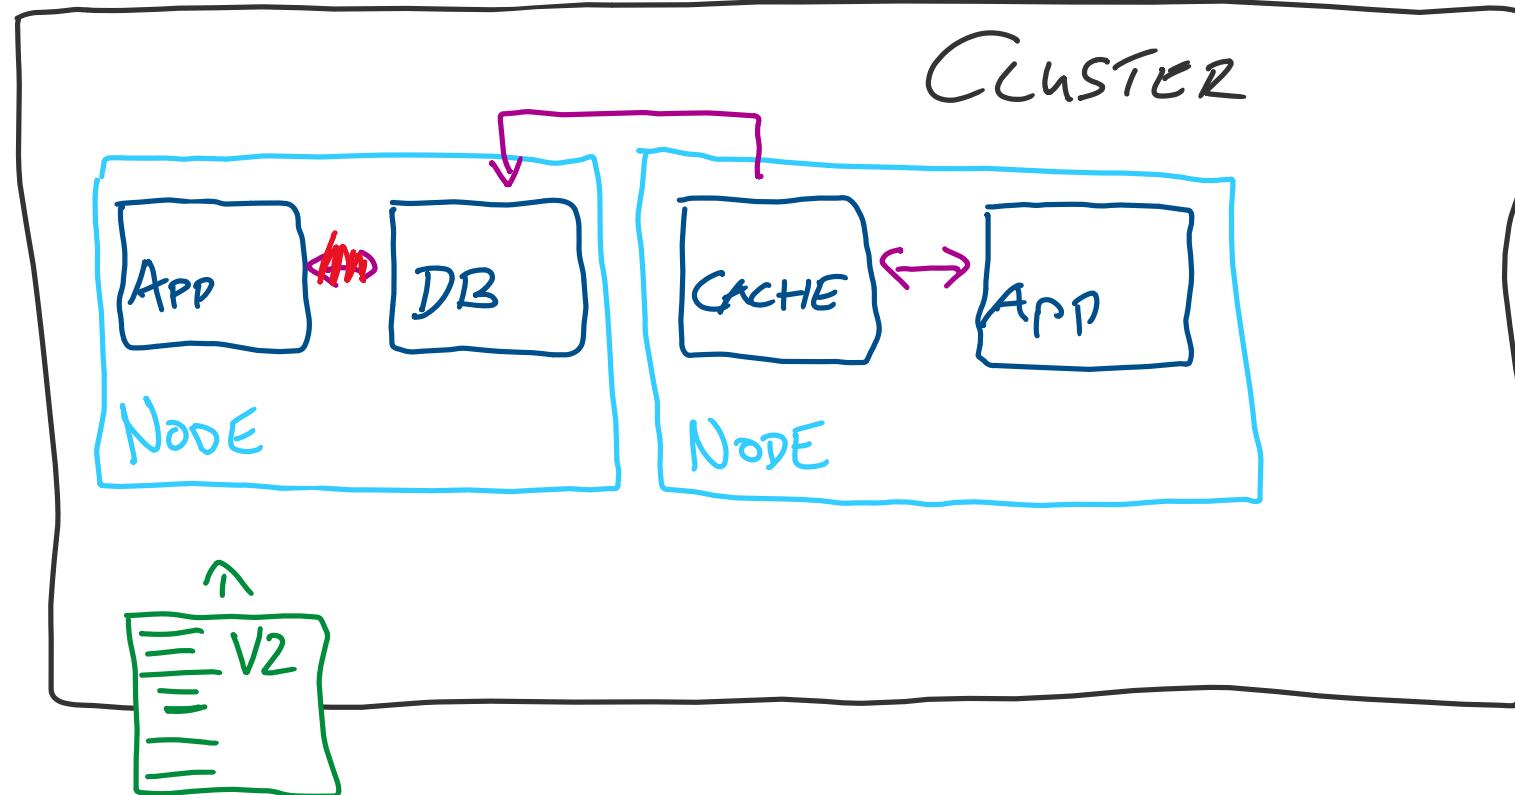


Container-optimised Linux?



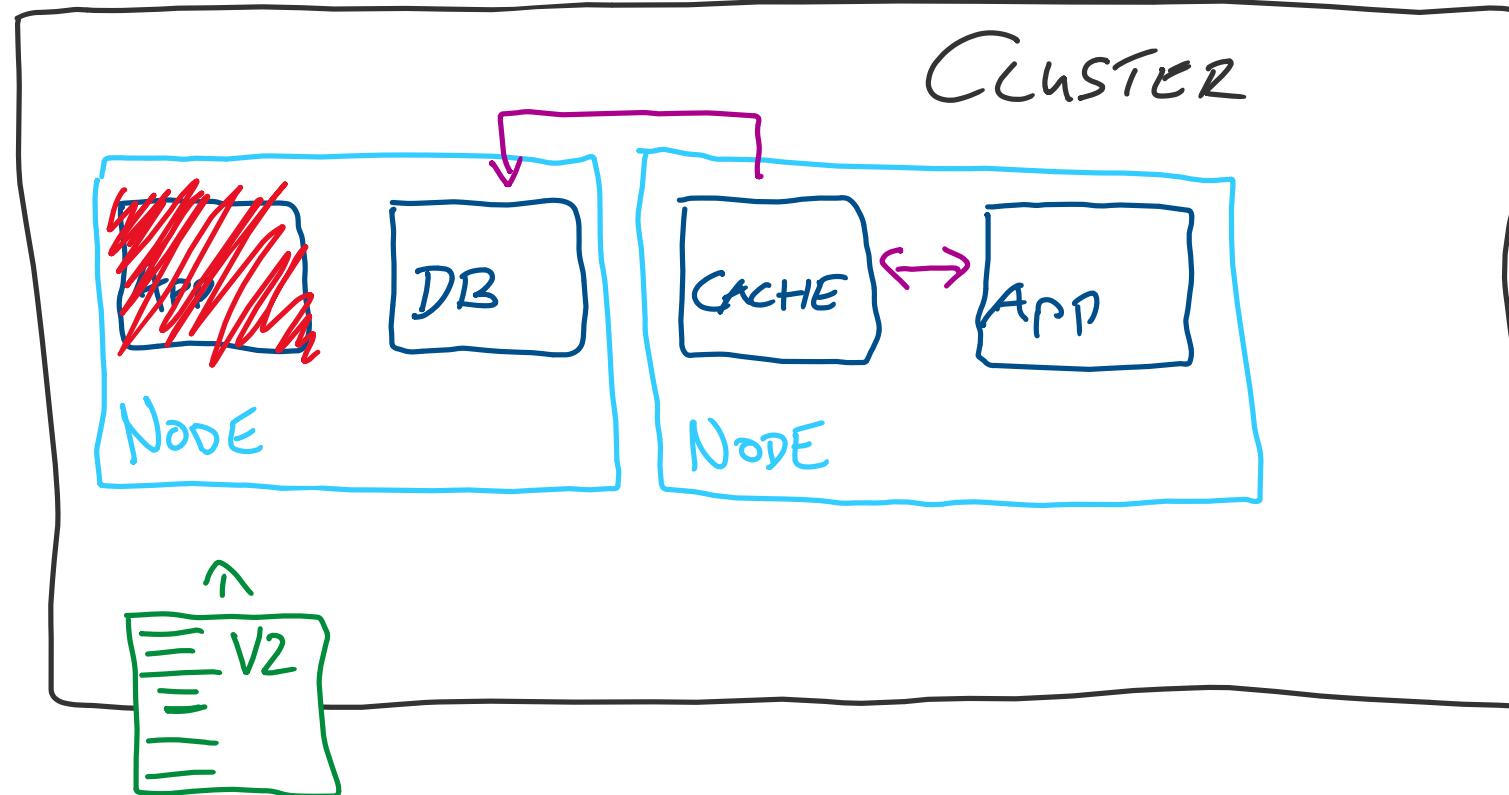


Container-optimised Linux?



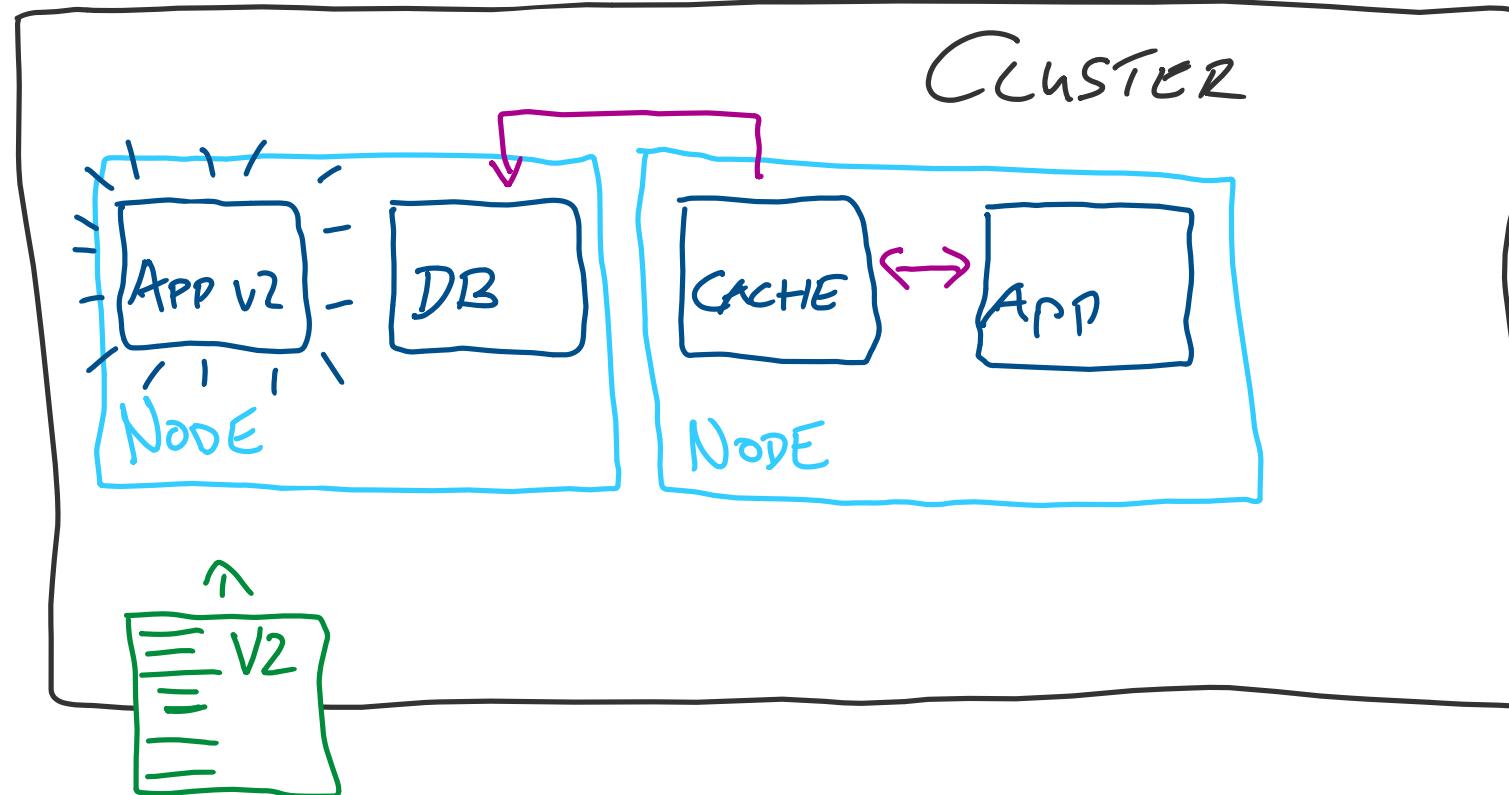


Container-optimised Linux?



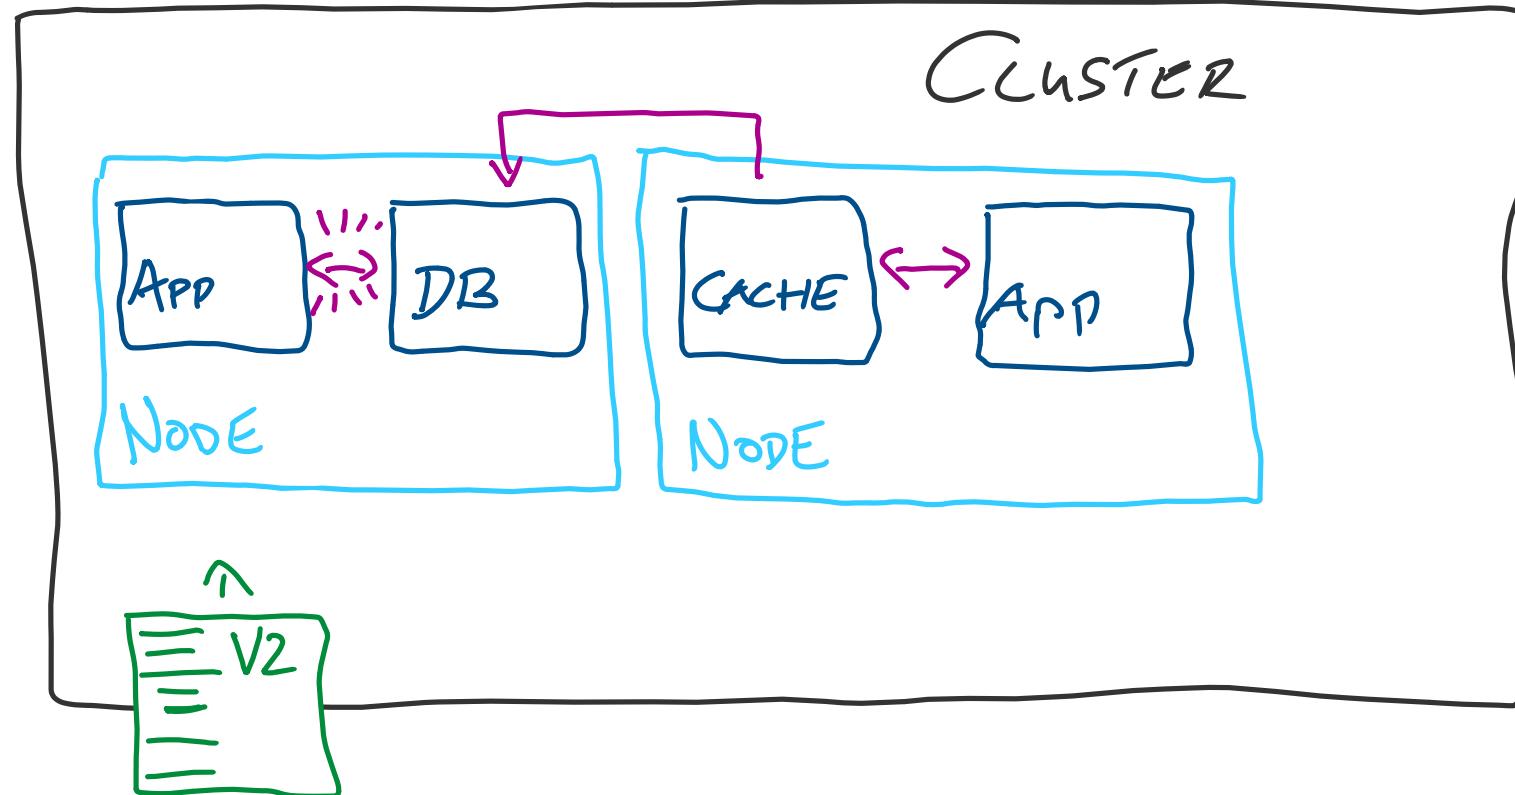


Container-optimised Linux?



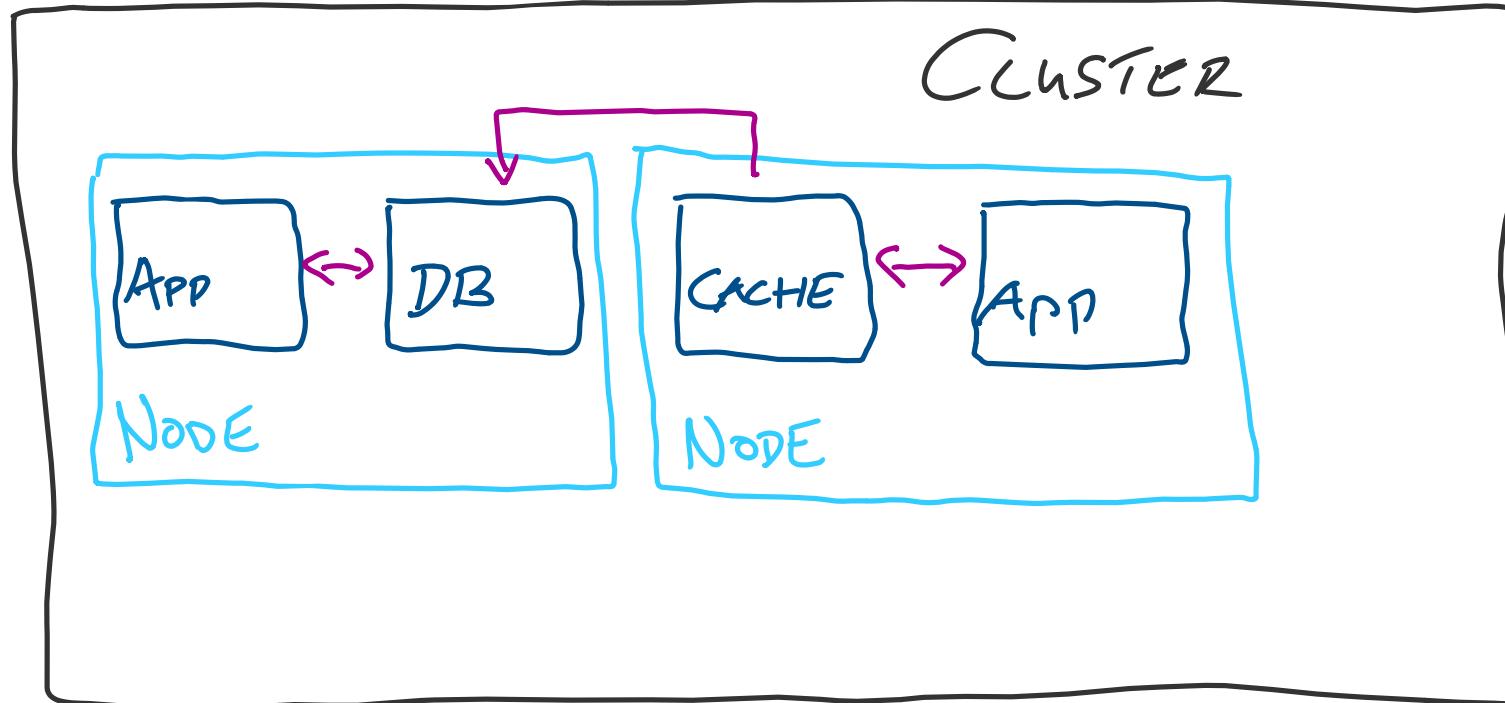


Container-optimised Linux?



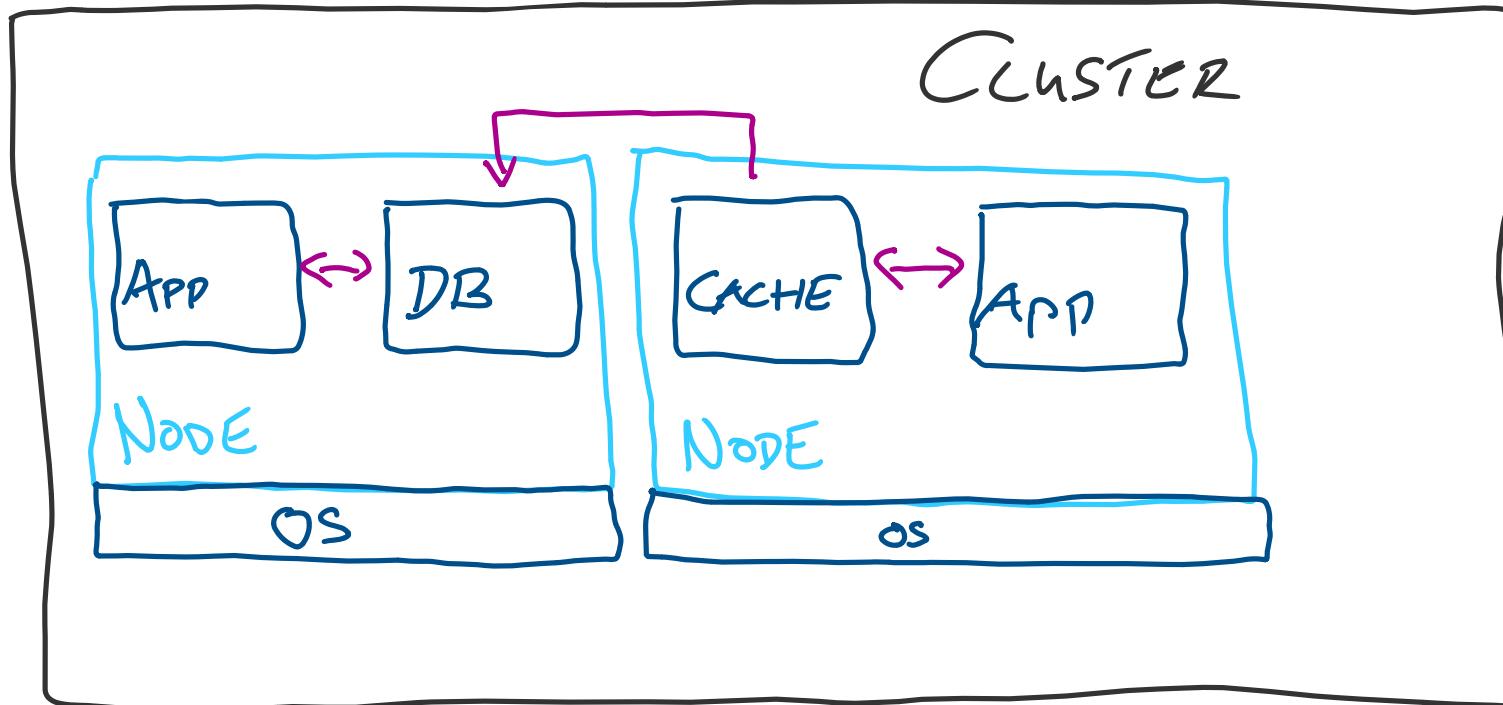


Container-optimised Linux?





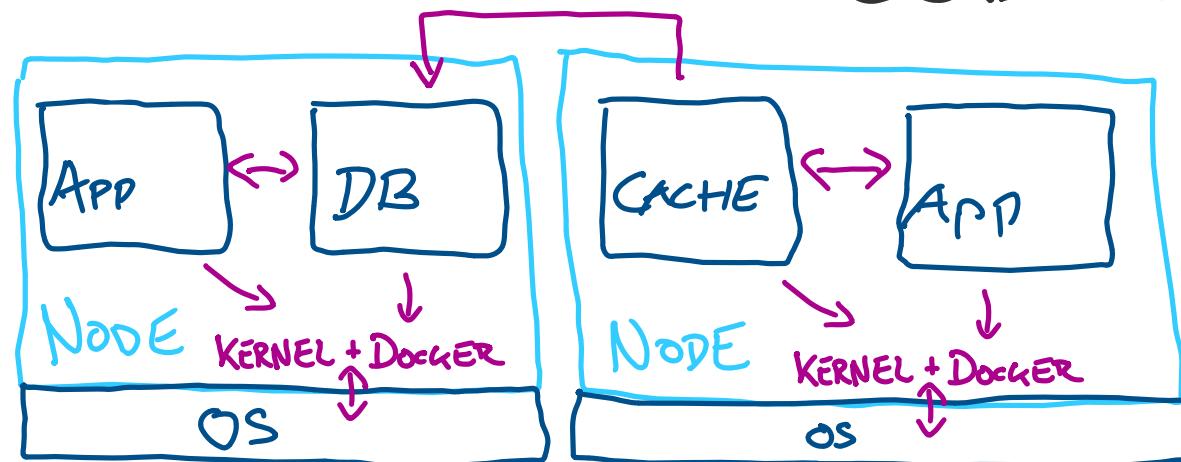
Container-optimised Linux?





Container-optimised Linux?

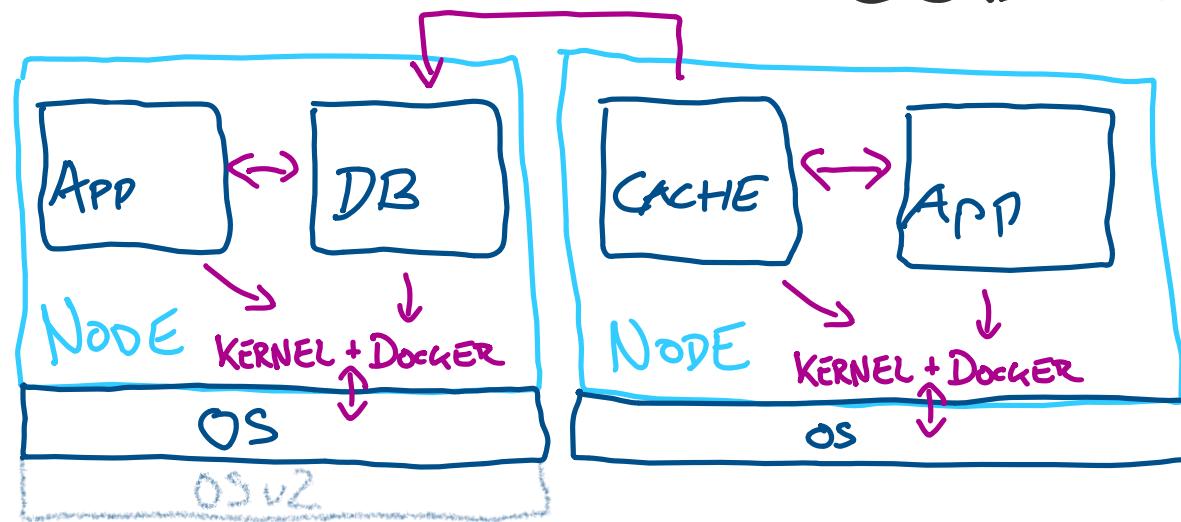
CLUSTER





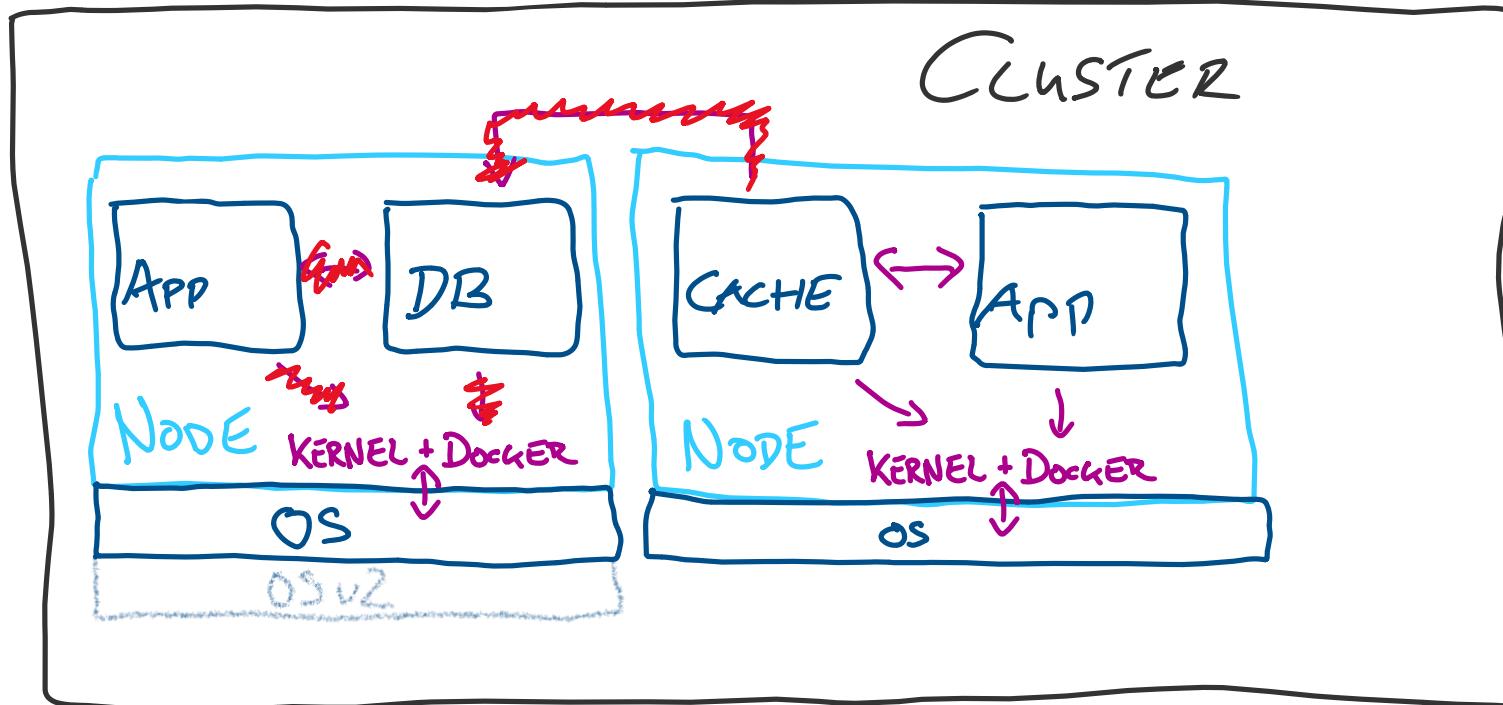
Container-optimised Linux?

CLUSTER



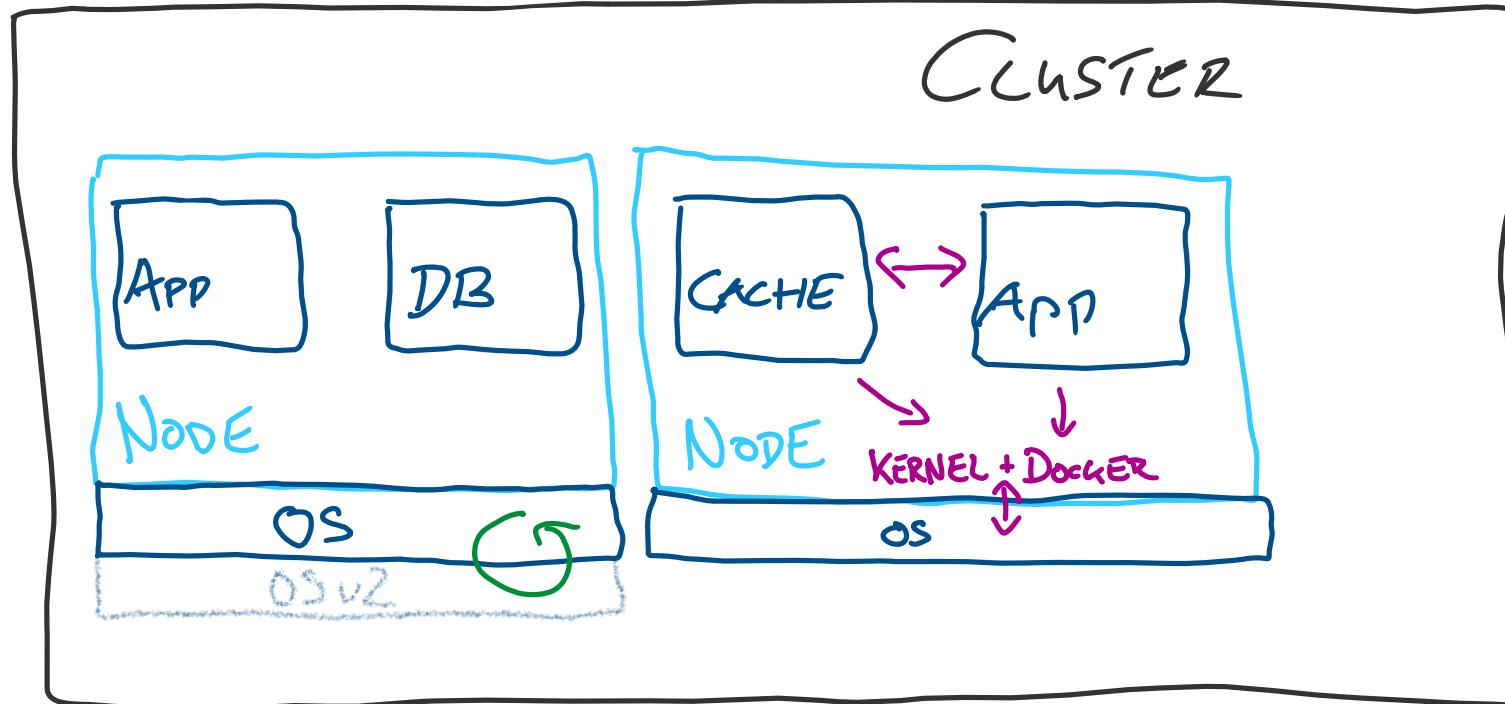


Container-optimised Linux?



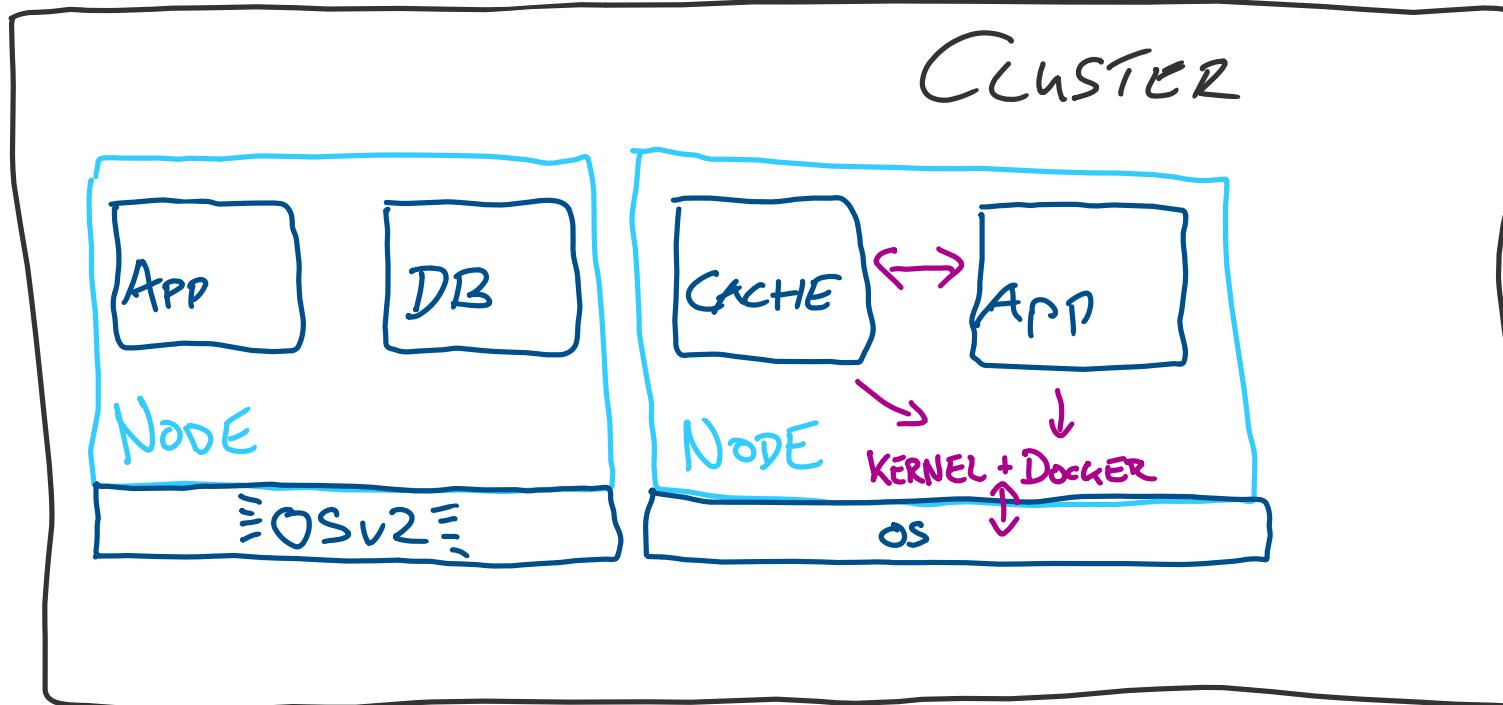


Container-optimised Linux?



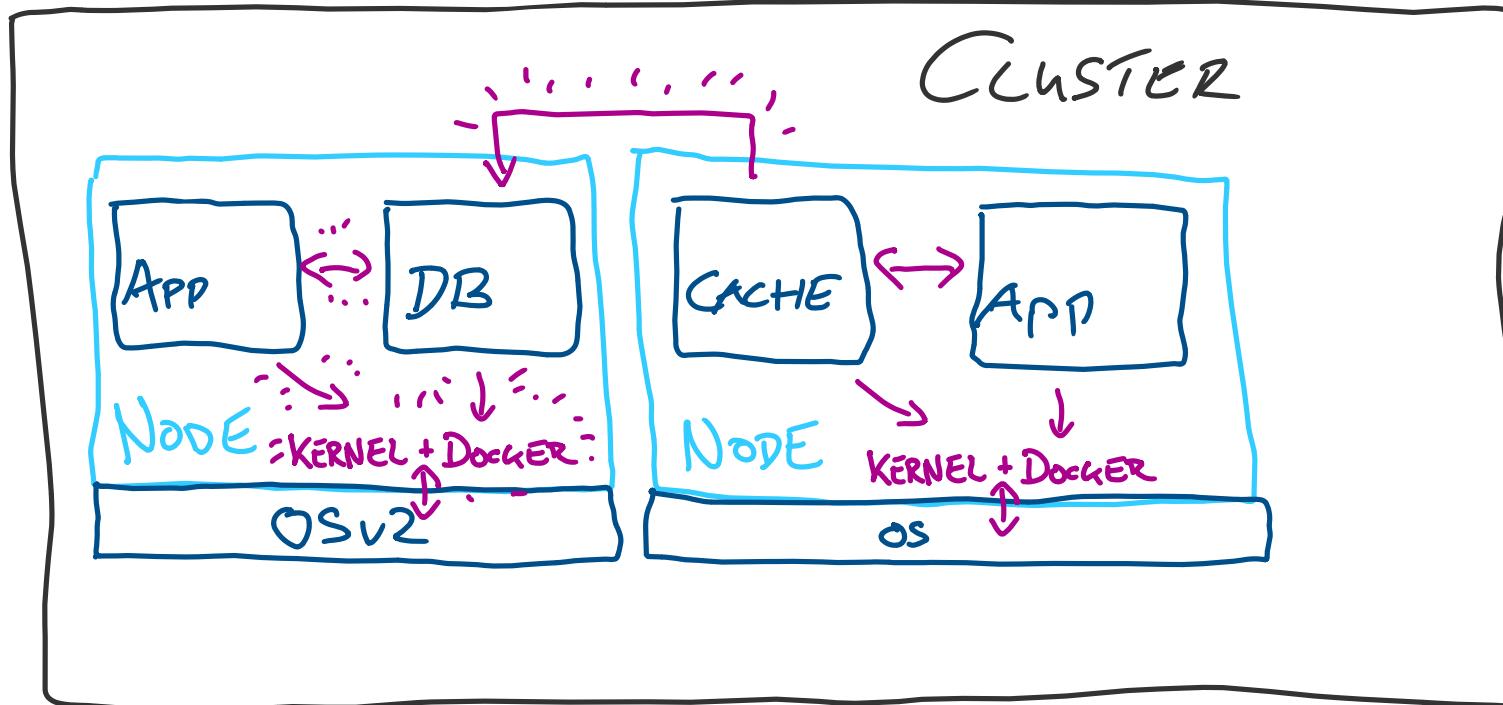


Container-optimised Linux?



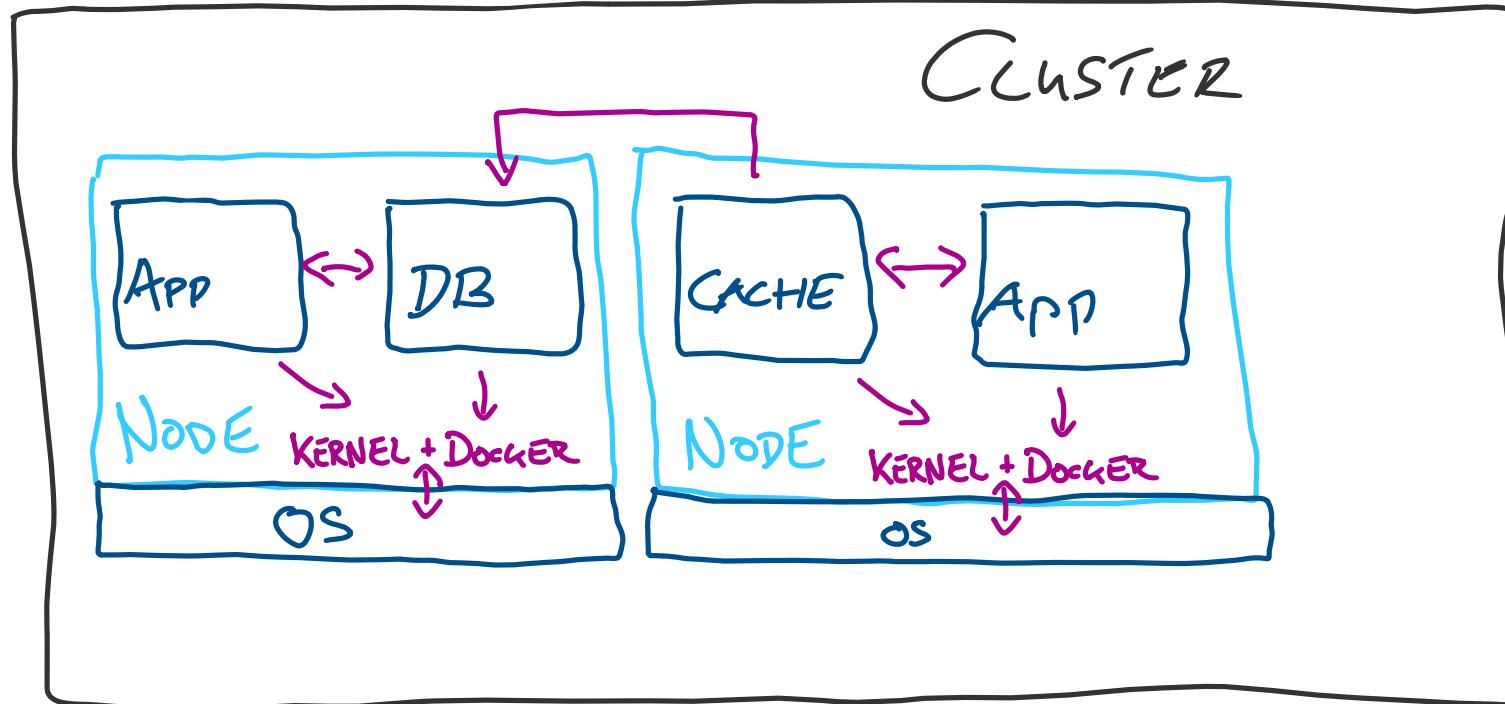


Container-optimised Linux?





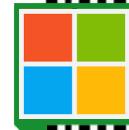
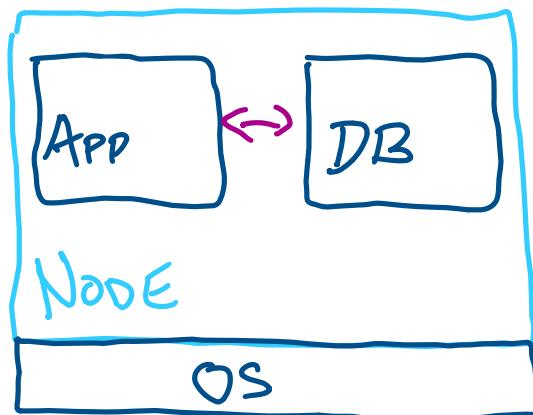
Container-optimised Linux?





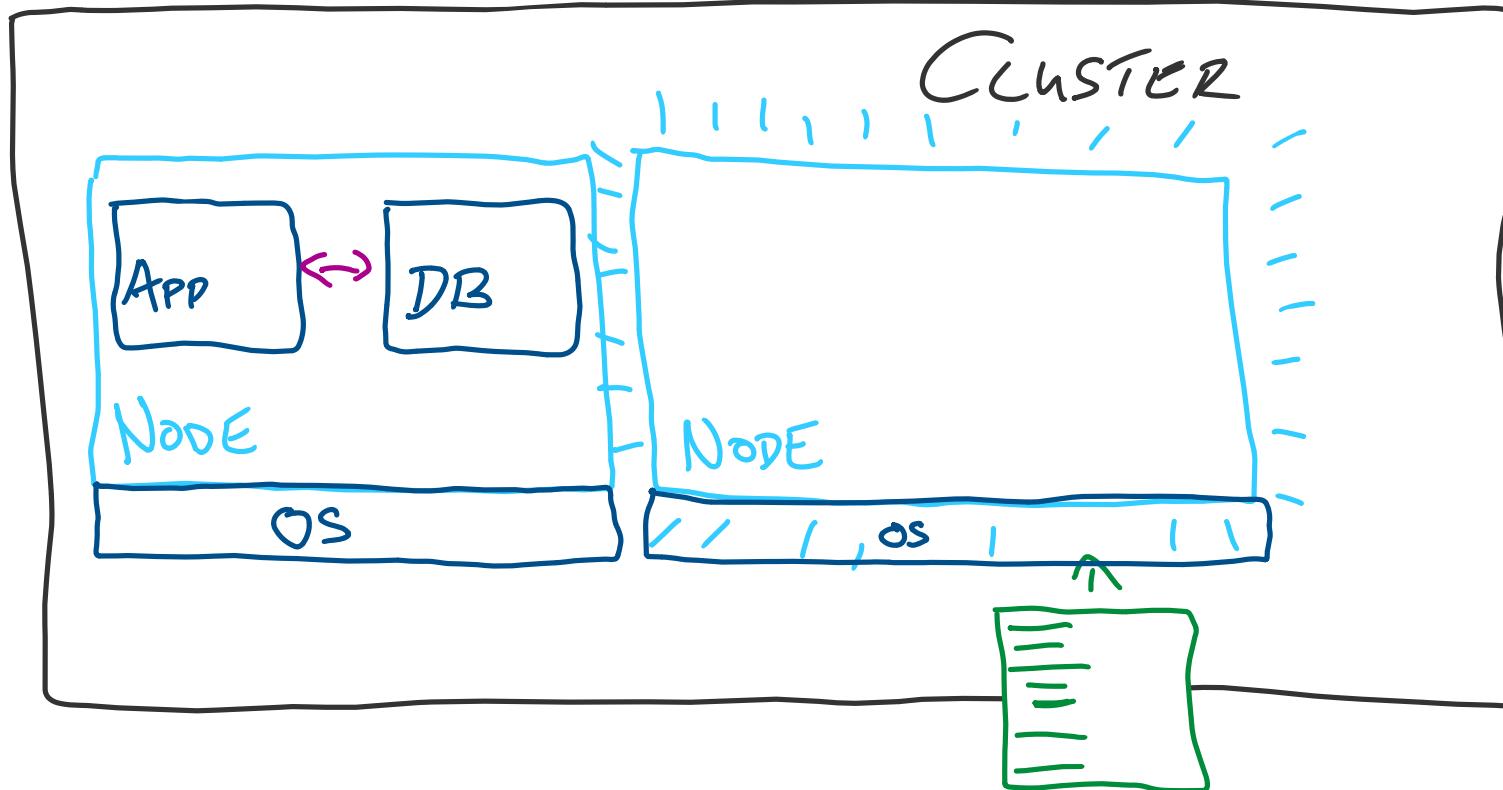
Container-optimised Linux?

Cluster



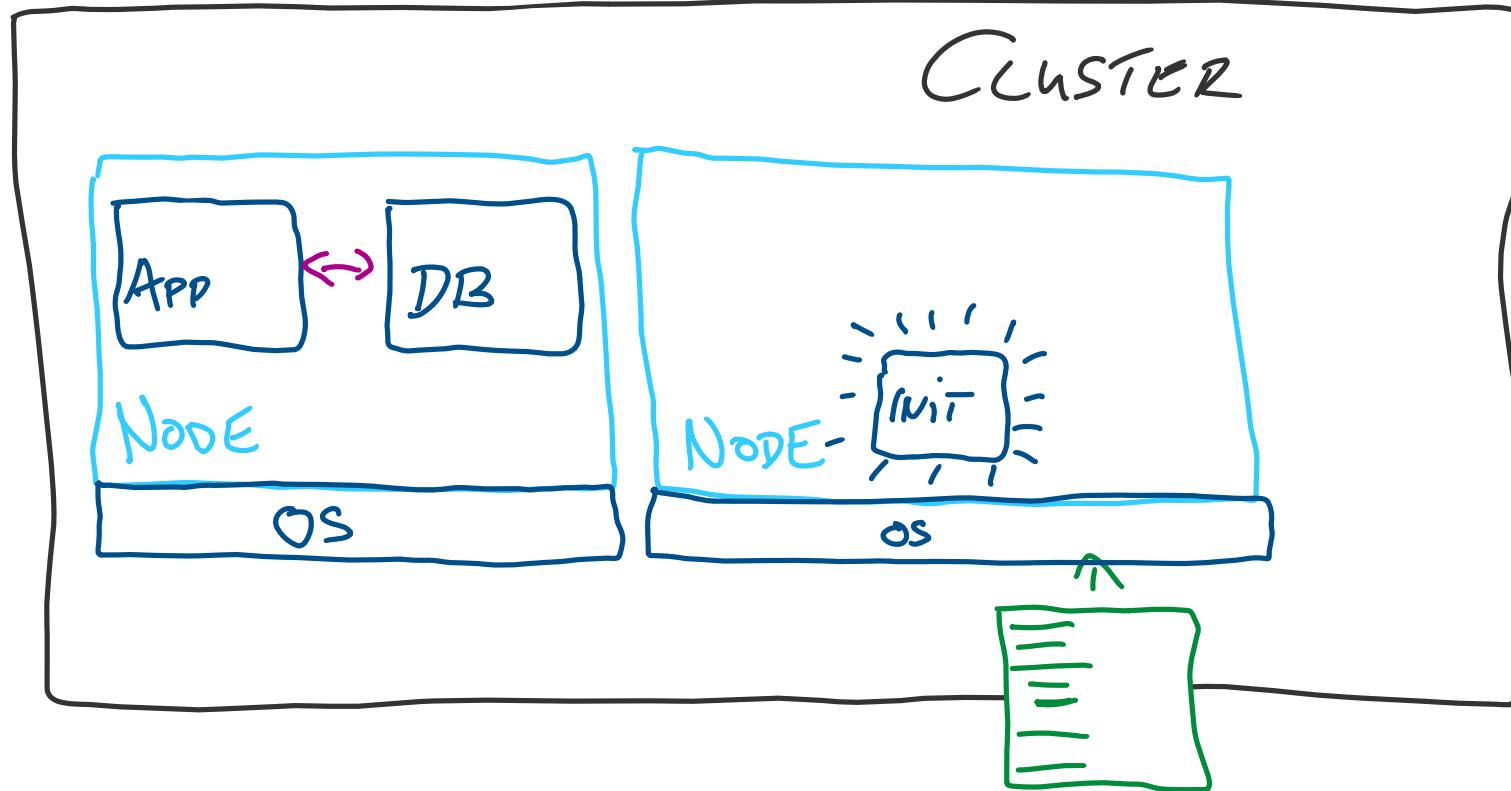


Container-optimised Linux?





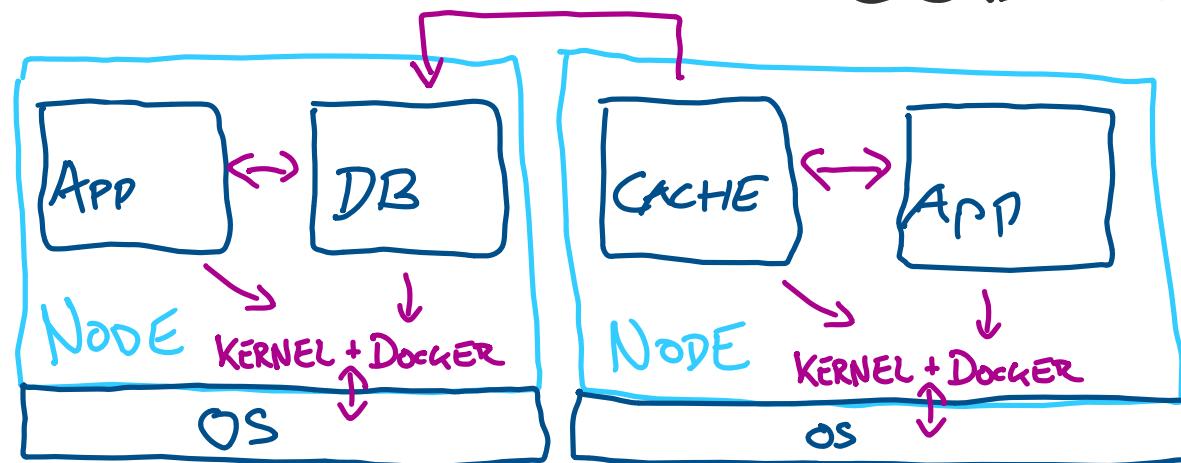
Container-optimised Linux?





Container-optimised Linux?

CLUSTER





Container-optimised Linux?

Isolated, no shared dependencies

Well-defined, limited interface to OS + "plumbing"

Declarative, automatable provisioning

Configure, apply at provisioning time. No dynamics.

Automated updates, controlled roll-out

Stage, switch, activate (or roll back)

==> Your OS as a stateless app.



Isolation



No inter-dependencies between OS and Apps

No OS <-> App shared libs. Just the Kernel ABI.

Clear and concise API/contract OS <-> Apps

No service dependencies. Just Docker/Containerd.

Updates and roll-backs with little to no side effects

Well-testable and canary friendly.



Provisioning automation



Declarative Configuration with Ignition

Applied once, at provisioning time.

Fully automatable

Thorough set of options, GO bindings, etc.

Easy to integrate with provisioning orchestration

[Terraform](#) provider, CAPI, etc.





```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]  
  
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    inline: |  
      <html><body align="center">  
      <h1>Hello KCD!</h1>  
        
      </body></html>  
- path: /srv/www/html/kcd.png  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    local: kcd.png
```

```
systemd:  
units:  
- name: kcd-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=KCD example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 \  
              -v /srv/www/html:/usr/share/caddy \  
              docker.io/caddy caddy file-server \  
              --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop nginx1  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
no_create_home: true  
groups: [ docker ]
```

```
storage:  
files:  
- path: /srv/www/html/index.html  
mode: 0644  
user:  
  name: caddy  
group:  
  name: caddy  
contents:  
  inline: |  
    <html><body align="center">  
    <h1>Hello KCD!</h1>  
      
    </body></html>  
- path: /srv/www/html/kcd.png  
mode: 0644  
user:  
  name: caddy  
group:  
  name: caddy  
contents:  
  local: kcd.png
```

```
systemd:  
units:  
- name: kcd-demo-webserver.service  
enabled: true  
contents: |  
  [Unit]  
  Description=KCD example static web server  
  After=docker.service  
  Requires=docker.service  
  [Service]  
  User=caddy  
  TimeoutStartSec=0  
  ExecStartPre=-/usr/bin/docker rm --force caddy  
  ExecStart=docker run -i -p 80:80 \  
    -v /srv/www/html:/usr/share/caddy \  
    docker.io/caddy caddy file-server \  
    --root /usr/share/caddy --access-log  
  ExecStop=/usr/bin/docker stop nginx1  
  Restart=always  
  RestartSec=5s  
  [Install]  
  WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]
```

```
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    inline: |  
      <html><body align="center">  
      <h1>Hello KCD!</h1>  
        
      </body></html>  
- path: /srv/www/html/kcd.png  
  mode: 0644  
  user:  
    name: caddy  
  group:  
    name: caddy  
  contents:  
    local: kcd.png
```

systemd:

units:

```
- name: kcd-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=KCD example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop nginx1  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
no_create_home: true  
groups: [ docker ]  
  
storage:  
files:  
- path: /srv/www/html/index.html  
mode: 0644  
user:  
  name: caddy  
group:  
  name: caddy  
contents:  
  inline: |  
    <html><body align="center">  
    <h1>Hello KCD!</h1>  
      
    </body></html>  
- path: /srv/www/html/kcd.png  
mode: 0644  
user:  
  name: caddy  
group:  
  name: caddy  
contents:  
  local: kcd.png
```

```
systemd:  
units:  
- name: kcd-demo-webserver.service  
enabled: true  
contents: |  
  [Unit]  
  Description=KCD example static web server  
  After=docker.service  
  Requires=docker.service  
  [Service]  
  User=caddy  
  TimeoutStartSec=0  
  ExecStartPre=-/usr/bin/docker rm --force caddy  
  ExecStart=docker run -i -p 80:80 \  
    -v /srv/www/html:/usr/share/caddy \  
    docker.io/caddy caddy file-server \  
    --root /usr/share/caddy --access-log  
  ExecStop=/usr/bin/docker stop nginx1  
  Restart=always  
  RestartSec=5s  
  [Install]  
  WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
  no_create_home: true  
  groups: [ docker ]
```

```
storage:  
files:  
- path: /srv/www/html/index.html  
  mode: 0644
```

```
  user:  
    name: caddy
```

```
  group:  
    name: caddy
```

```
  contents:  
    inline: |  
      <html><body align="center">  
      <h1>Hello KCD!</h1>  
        
      </body></html>
```

```
- path: /srv/www/html/kcd.png
```

```
  mode: 0644
```

```
  user:  
    name: caddy
```

```
  group:  
    name: caddy
```

```
  contents:  
    local: kcd.png
```

```
systemd:
```

```
units:
```

```
- name: kcd-demo-webserver.service  
  enabled: true  
  contents: |  
    [Unit]  
    Description=KCD example static web server  
    After=docker.service  
    Requires=docker.service  
    [Service]  
    User=caddy  
    TimeoutStartSec=0  
    ExecStartPre=-/usr/bin/docker rm --force caddy  
    ExecStart=docker run -i -p 80:80 \  
      -v /srv/www/html:/usr/share/caddy \  
      docker.io/caddy caddy file-server \  
      --root /usr/share/caddy --access-log  
    ExecStop=/usr/bin/docker stop nginx1  
    Restart=always  
    RestartSec=5s  
    [Install]  
    WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
no_create_home: true  
groups: [ docker ]
```

```
storage:  
files:  
- path: /srv/www/html/index.html  
mode: 0644  
user:  
  name: caddy  
group:  
  name: caddy  
contents:  
  inline: |  
    <html><body align="center">  
    <h1>Hello KCD!</h1>  
      
    </body></html>  
- path: /srv/www/html/kcd.png  
mode: 0644  
user:  
  name: caddy  
group:  
  name: caddy  
contents:  
  local: kcd.png
```

```
systemd:  
units:  
- name: kcd-demo-webserver.service  
enabled: true  
contents: |  
  [Unit]  
  Description=KCD example static web server  
  After=docker.service  
  Requires=docker.service  
  [Service]  
  User=caddy  
  TimeoutStartSec=0  
  ExecStartPre=-/usr/bin/docker rm --force caddy  
  ExecStart=docker run -i -p 80:80 \  
    -v /srv/www/html:/usr/share/caddy \  
    docker.io/caddy caddy file-server \  
    --root /usr/share/caddy --access-log  
  ExecStop=/usr/bin/docker stop nginx1  
  Restart=always  
  RestartSec=5s  
  [Install]  
  WantedBy=multi-user.target
```



```
passwd:  
users:  
- name: caddy  
no_create_home: true  
groups: [ docker ]
```

```
storage:  
files:  
- path: /srv/www/html/index.html  
mode: 0644
```

```
user:  
  name: caddy
```

```
group:  
  name: caddy
```

```
contents:  
  inline: |  
    <html><body align="center">  
    <h1>Hello KCD!</h1>  
      
    </body></html>
```

```
- path: /srv/www/html/kcd.png
```

```
mode: 0644
```

```
user:  
  name: caddy
```

```
group:  
  name: caddy
```

```
contents:  
  local: kcd.png
```

```
systemd:
```

```
units:
```

```
- name: kcd-demo-webserver.service  
enabled: true  
contents: |  
  [Unit]  
  Description=KCD example static web server  
  After=docker.service  
  Requires=docker.service  
  [Service]  
  User=caddy  
  TimeoutStartSec=0  
  ExecStartPre=-/usr/bin/docker rm --force caddy  
  ExecStart=docker run -i -p 80:80 \  
    -v /srv/www/html:/usr/share/caddy \  
    docker.io/caddy caddy file-server \  
    --root /usr/share/caddy --access-log  
  ExecStop=/usr/bin/docker stop nginx1  
  Restart=always  
  RestartSec=5s  
  [Install]  
  WantedBy=multi-user.target
```



Image-based OS

Stateless installations

OS is replaceable without side effects

Immutable (read-only) OS partition

No changes to OS binaries, accidentally or otherwise.

Verifiable / attestable

OS never changes, build / release signature remains valid





Updates

Atomic updates

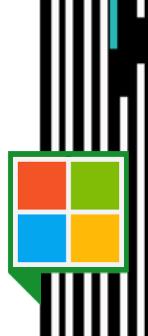
*Staged on inactive OS partition in the background
Applied at reboot*

Automated, configurable roll-back

Define required services and timeouts

Controlled roll-out

*Cluster-wide reboot coordination
Roll-back support*





Updates – release stabilisation

Alpha -> Beta -> Stable (->LTS)

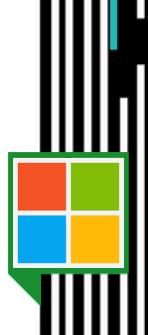
Alpha is for dev, Beta for canaries, and Stable

Extensive scenario testing for each release

Complex - Kubernetes workloads, CNI, etc.

All releases always pass full test suite

Beta canaries ensure Stable runs fine with your workloads





Scaling out

(Community) LTS

When you're so big that any maintenance is costly.

FOSS update server ([Nebraska](#))

Self-hosted updates, fine-grained roll-out, fleet mgmt.

User-centric, community-driven, fully open source

Speak up, join in, participate!





Community driven

Single vendor past

Back in the day right after our friendly fork

Freed ourselves, embracing the community

Re-orientation of project stewardship

Microsoft supports the community

And publicly commits to this

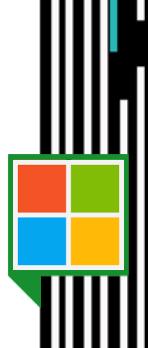


MS committed to support and to foster community



<https://azure.microsoft.com/en-us/blog/microsoft-acquires-kinvolk-to-accelerate-containeroptimized-innovation/>

“Flatcar Container Linux has a sizeable community of users on Azure, as well as other clouds, and on-premises. We know the CoreOS community has been on a winding journey over the years—we want to assure the Flatcar community that Microsoft and the Kinvolk team will continue to collaborate with the larger Flatcar community on the evolution of Flatcar Container Linux. Microsoft is committed to Flatcar Container Linux community development and will invest in working with the Flatcar community to create a growth path forward together. We’ll have our first meeting with the community within the coming weeks and invite anyone interested to attend and join the conversation.”

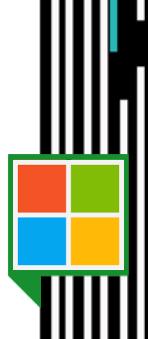


MS committed to support and to foster community



<https://azure.microsoft.com/en-us/blog/microsoft-acquires-kinvolk-to-accelerate-containeroptimized-innovation/>

“Flatcar Container Linux has a sizeable community of users on Azure, as well as other clouds, and on-premises. We know the CoreOS community has been on a winding journey over the years—we want to assure the Flatcar community that Microsoft and the Kinvolk team will continue to collaborate with the larger Flatcar community on the evolution of Flatcar Container Linux. Microsoft is committed to Flatcar Container Linux community development and will invest in working with the Flatcar community to create a growth path forward together. We’ll have our first meeting with the community within the coming weeks and invite anyone interested to attend and join the conversation.”





Work in the open

Day-to-day interactions

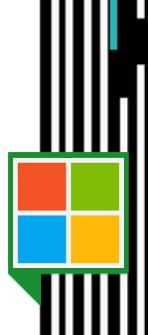
Matrix, Slack – line work and releases brawl

Planning, short and long term

Office hours, release planning, project boards

Bug smashing and docs writing days

Live streamed! Next one is 8th of July.





What's going on? - Boards

Implementation

day-to-day kanban

Release

Wrap up implementations

Roadmap

Epics with subtasks

The image displays four screenshots of GitHub boards for the Flatcar organization, illustrating the team's workflow across different stages of development and deployment:

- Upcoming / Backlog:** Shows a list of items ready for the backlog, including tasks like "Add documentation on running Flatcar with EDL" and "Update packages in coreos-stable and automation to handle them".
- Blocked:** A board for tracking blocked items, such as "Create API files, symbols and headers if they aren't already present" and "Update flatcar-configuration upstream".
- Planned / To Do:** A board for planning future work, including "Trigger the scripts CI GitHub Actions for coresos-overlay/portage-stable" and "Run legacy Intel tools as part of the GitHub Action workflow".
- In Progress:** A board for tracking current work, including "Run legacy Intel tools as part of the GitHub Action workflow" and "Run legacy Intel tools as part of the GitHub Action workflow".
- Testing / In Review:** A board for testing and reviewing code, including "Support image signing in new CI pipeline" and "Support image signing in new CI pipeline".
- Release:** A board for managing releases, showing two releases:
 - Release: 2022-07-18**: Alpha 3304.0.0 Beta 3277.1.0 Stable 3327.2.0 LTS-2022 3033.3 LTS-2021 2605.29.1
 - Release: 2022-06-20**: Alpha 3277.0.0 Beta 3277.1.1 Stable 3139.2.3 LTS-2022 3033.3 LTS-2021 2605.29.1





Contribute to other projects

Gentoo

package updates, stabilisation, bug fixes

Ignition, Butane

Extensions and bug fixes

Linux Kernel, Cilium, Falco

Bug fixes and interop



Contribute



Report bugs and issues, give feedback

We want your opinion

Bug smashing / doc writing days

Next bug smash is on the 8th of July, live streamed

Package updates / add tools to OS image

Become a Flatcar dev!





Stay in touch

Office hours

Every second Tuesday of the month, 5:30 CE[S]T

Matrix

Our day-to-day comms

Slack

Same



Demo time

We'll use the web server set-up shown earlier

(Almost. Update auto-reboot will be disabled.)

We will also update the OS.

Straightforward. Boring. Your Feierabend is safe.

The demo sources are available

<https://github.com/flatcar-linux/flatcar-demos/tree/main/kcd-berlin-2022-06-30>



Thank you 🎉



Look us up – <https://www.flatcar.org>

Join our office hours - <https://github.com/flatcar-linux/Flatcar/#monthly-office-hours-and-release-planning>

Contribute – <https://github.com/flatcar-linux/Flatcar>

Chat with us – <https://app.element.io/#/room/#flatcar:matrix.org>
– <https://kubernetes.slack.com/archives/C03GQ8B5XNJ>

